

Getting Started in Functional Programming

Introducing

Erlang



O'REILLY®

Simon St. Laurent

Introducing Erlang

If you're new to Erlang, its functional style can seem difficult, but with help from this hands-on introduction, you'll scale the learning curve and discover how enjoyable, powerful, and fun this language can be.

Author Simon St. Laurent shows you how to write simple Erlang programs by teaching you one basic skill at a time. You'll learn about pattern matching, recursion, message passing, process-oriented programming, and establishing pathways for data rather than telling it where to go. By the end of your journey, you'll understand why Erlang is ideal for concurrency and resilience.

- Get cozy with Erlang's shell, its command-line interface
- Become familiar with Erlang's basic structures by working with numbers
- Discover atoms, pattern matching, and guards: the foundations of your program structure
- Delve into the heart of Erlang processing with recursion, strings, lists, and higher-order functions
- Create processes, send messages among them, and apply pattern matching to incoming messages
- Store and manipulate structured data with Erlang Term Storage and the Mnesia database
- Learn about the Open Telecom Platform, Erlang's open source libraries and tools

Purchase the ebook edition of this O'Reilly title at oreilly.com and get free updates for the life of the edition. Our ebooks are optimized for several electronic formats, including PDF, EPUB, Mobi, and DAISY—all DRM-free.

Twitter: @oreillymedia
facebook.com/oreilly

US \$19.99

CAN \$20.99

ISBN: 978-1-449-33176-4



9 781449 331764

O'REILLY®
oreilly.com

Introducing Erlang

Simon St. Laurent

O'REILLY®

Beijing • Cambridge • Farnham • Köln • Sebastopol • Tokyo

Introducing Erlang

by Simon St. Laurent

Copyright © 2013 Simon St. Laurent. All rights reserved.

Printed in the United States of America.

Published by O'Reilly Media, Inc., 1005 Gravenstein Highway North, Sebastopol, CA 95472.

O'Reilly books may be purchased for educational, business, or sales promotional use. Online editions are also available for most titles (<http://my.safaribooksonline.com>). For more information, contact our corporate/institutional sales department: 800-998-9938 or corporate@oreilly.com.

Editors: Simon St. Laurent and Meghan Blanchette

Production Editor: Melanie Yarbrough

Cover Designer: Karen Montgomery

Interior Designer: David Futato

Illustrator: Rebecca Demarest

January 2013: First Edition

Revision History for the First Edition:

2013-01-23 First release

See <http://oreilly.com/catalog/errata.csp?isbn=9781449331764> for release details.

Nutshell Handbook, the Nutshell Handbook logo, and the O'Reilly logo are registered trademarks of O'Reilly Media, Inc. *Introducing Erlang*, the cover image of a giant red flying squirrel, and related trade dress are trademarks of O'Reilly Media, Inc.

Many of the designations used by manufacturers and sellers to distinguish their products are claimed as trademarks. Where those designations appear in this book, and O'Reilly Media, Inc., was aware of a trademark claim, the designations have been printed in caps or initial caps.

While every precaution has been taken in the preparation of this book, the publisher and author assume no responsibility for errors or omissions, or for damages resulting from the use of the information contained herein.

ISBN: 978-1-449-33176-4

[LSI]

Table of Contents

Preface.....	vii
1. Getting Comfortable.....	1
Installation	1
Firing It Up	2
First Steps: The Shell	2
Moving through Text	3
Moving through History	3
Moving through Files	4
Doing Something	4
Calling Functions	5
Numbers in Erlang	6
Working with Variables in the Shell	8
Seeing Your Bound Variables	10
Clearing Bound Variables in the Shell	10
2. Functions and Modules.....	13
Fun with fun	13
Defining Modules	15
From Module to Fun	17
Functions and Variable Scope	18
Module Directives	18
Documenting Code	20
Documenting Modules	21
Documenting Functions	22
Documenting Your Application	24
3. Atoms, Tuples, and Pattern Matching.....	27
Atoms	27
Pattern Matching with Atoms	28

Atomic Booleans	29
Guards	30
Underscoring That You Don't Care	32
Adding Structure: Tuples	34
Pattern Matching with Tuples	35
Processing Tuples	36
4. Logic and Recursion.....	39
Logic Inside of Functions	39
Evaluating Cases	39
If This, Then That	42
Variable Assignment in case and if Constructs	44
The Gentlest Side Effect: io:format	45
Simple Recursion	46
Counting Down	46
Counting Up	47
Recursing with Return Values	49
5. Communicating with Humans.....	53
Strings	53
Asking Users for Information	55
Gathering Terms	56
Gathering Characters	59
Reading Lines of Text	60
6. Lists.....	65
List Basics	65
Splitting Lists into Heads and Tails	67
Processing List Content	68
Creating Lists with Heads and Tails	70
Mixing Lists and Tuples	72
Building a List of Lists	73
7. Higher-Order Functions and List Comprehensions.....	77
Simple Higher-Order Functions	77
Creating New Lists with Higher-Order Functions	79
Reporting on a List	79
Running List Values Through a Function	80
Filtering List Values	81
Beyond List Comprehensions	81
Testing Lists	82
Splitting Lists	82

Folding Lists	83
8. Playing with Processes.....	87
The Shell Is a Process	87
Spawning Processes from Modules	89
Lightweight Processes	92
Registering a Process	92
When Processes Break	94
Processes Talking Amongst Themselves	95
Watching Your Processes	97
Breaking Things and Linking Processes	99
9. Exceptions, Errors, and Debugging.....	107
Flavors of Errors	107
Catching Runtime Errors as They Happen	108
Raising Exceptions with throw	110
Logging Progress and Failure	111
Debugging through a GUI	112
Tracing Messages	119
Watching Function Calls	122
10. Storing Structured Data.....	125
From Tuples to Records	125
Setting Up Records	126
Creating and Reading Records	127
Using Records in Functions and Modules	128
Storing Records in Erlang Term Storage	131
Creating and Populating a Table	132
Simple Queries	137
A Key Feature: Overwriting Values	138
ETS Tables and Processes	138
Next Steps	140
Storing Records in Mnesia	141
Starting up Mnesia	141
Creating Tables	142
Reading Data	145
Query List Comprehensions	146
11. Getting Started with OTP.....	151
Creating Services with gen_server	152
A Simple Supervisor	157

Packaging an Application	162
12. Next Steps Through Erlang.....	165
Moving Beyond the Shell	165
Distributed Computing	166
Processing Binary Data	166
Input and Output	166
Testing, Analyzing, and Refactoring	166
Networking and the Web	167
Data Storage	167
Extending Erlang	168
Languages Built on Erlang	168
Community	168
Sharing the Gospel of Erlang	169
A. An Erlang Parts Catalog.....	171
B. OTP Templates.....	179

Preface

Erlang has long been a mysterious dark corner of the programming universe, visited mostly by developers who need extreme reliability or scalability and people who want to stretch their brains.

Developed at Ericsson to serve on telephone switching equipment, it seemed like a strangely special-purpose language until recently, when our computer and network architectures came to look a lot more like massively parallel telephone-switching equipment. Thanks to the rise of NoSQL data stores CouchDB and Riak, you may already be using Erlang without realizing it, and Erlang is moving out into many more fields.

Exploring Erlang, if you come from pretty much any background other than functional programming, will require you to clear your mind of many techniques used in other programming languages. Forget classes, forget variables that change values—even forget the conventions of variable assignment.

Instead, you’re going to have to think about pattern matching, message passing, and establishing pathways for data rather than telling it where to go. Erlang programming can feel like making a key whose teeth set the tumblers on a lock just right for the key to pass, or playing pachinko and watching the balls fall through a maze.

Sound strange? It is—but also enjoyable, powerful, and fun.

My first explorations of Erlang confused me and excited me at the same time. I’d had some experience with what I’d called “invariant variables,” variables that can be bound to a value only once, in XSLT. That created a lot of headaches for me until I realized I was coming at the problems all wrong, and then it suddenly made sense.

Who This Book Is For

This book is mostly for people who've been programming in other languages but want to look around. Maybe you're being very practical, and Erlang's distributed model and the resulting scale and resilience advantages appeal to you. Maybe you want to see what this "functional programming" stuff is all about. Or maybe you're just going for a hike, taking your mind to a new place.

I suspect that functional programming is more approachable as a first language, before you've learned to program in other paradigms. However, getting started in Erlang—sometimes even just installing it—requires a fair amount of computing skill. If you're a complete newcomer to programming, welcome, but there will be a few challenges along the way.

Who This Book Is Not For

This book is not for people in a hurry to get things done.

If you already know Erlang, you don't likely need this book unless you're looking for a slow brush-up.

If you're already familiar with functional languages, you may find the pacing of this "Gentle Introduction" hopelessly slow. Definitely feel welcome to jump to another book that moves faster if you get bored. Come back if you find the others go too fast, and feel welcome to use this as a companion guide or reference with the other books.

What This Book Will Do For You

In *Seven Languages in Seven Weeks*, Bruce Tate suggests that "Erlang makes hard things easy and easy things hard." This book will get you through the "easy things hard" part, and show you a bit of the promised land of "hard things easy."

In practical terms, you'll learn to write simple Erlang programs. You'll understand why Erlang makes it easier to build resilient programs that can scale up and down with ease. Perhaps most importantly, you'll be able to read other Erlang resources that assume a fair amount of experience and make sense of them.

In more theoretical terms, you'll get to know functional programming. You'll learn how to design programs around message-passing and recursion, creating process-oriented programs focused more on data flow.

You'll also be better prepared to read other books and conversations about Erlang.

How This Book Works

This book tries to tell a story with Erlang. You'll probably get the most out of it if you read it in order at least the first time, though you're always welcome to come back to find whatever bits and pieces you need.

You'll start by getting Erlang installed and running, and looking around its shell. You'll spend a lot of time in the shell, so get cozy. Next, you'll start loading code into the shell to make it easier to write programs, and you'll learn how to call that code and mix it up.

You'll take a close look at numbers because they're an easy place to get familiar with Erlang's basic structures. Then you'll learn about atoms, pattern-matching, and guards—the likely foundations of your program structure. After that you'll learn about strings, lists, and the recursion at the heart of much Erlang processing. Once you've gone a few million recursions down and back, it'll be time to look at processes, a key part of Erlang that relies on the message-passing model to support concurrency and resilience.

Once you have the foundation set, you can take a closer look at debugging and data storage, and then get a quick look at a toolset that is likely at the heart of your long-term development with Erlang: the Open Telecom Platform (OTP), which is about much much more than telephones.

Some people want to learn programming languages through a dictionary. Here's a list of operators, here's a list of control structures, these are the datatypes—and then smash them together. Those lists are here, but they're in [Appendix A](#), not the main flow of the book.

The main point you should get from this book is that you can program in Erlang. If you don't get that, let me know!

Why I Wrote This Book

I'm not an Erlang expert hoping to create more Erlang experts to get a lot of work done.

I'm a developer and writer who encountered Erlang, thought it was the programming language I'd been seeking for a long time, and felt compelled to share some of that. I'm hoping that the path I followed will work for other people, probably with variations, and that a book written from a beginner's perspective (and vetted by experts) would help more people find and enjoy Erlang.

Other Resources

This book may not be the best way for you to learn Erlang. It all depends on what you want to learn and why.

If your primary interest in learning Erlang is to break out of a programming rut, you should explore Bruce Tate’s wild tour of *Seven Languages in Seven Weeks* (Pragmatic Publishers), which explores Ruby, Io, Prolog, Scala, Erlang, Clojure, and Haskell. Erlang gets only (an excellent) 37 pages, but that might be what you want.

For an online experience (now also in print from No Starch Books) with more snark and funnier illustrations, you should explore Fred Hebert’s *Learn You Some Erlang for Great Good!*, at <http://learnyousomeerlang.com/>. While much longer than Tate’s telling, it certainly moves faster and may feel more like an experienced programmer’s guide to Erlang.

The two classic general books on Erlang are the similarly-titled *Programming Erlang* (Pragmatic Publishers) by Erlang creator Joe Armstrong, and *Erlang Programming* (O’Reilly) by Francesco Cesarini and Simon Thompson. They cover a lot of similar and overlapping terrain, and both may be good places to start if this book moves too slowly or you need more reference material. *Erlang Programming* goes further into what you can do with Erlang, whereas *Programming Erlang* provides a lot of detail on setting up an Erlang programming environment.

On the more advanced side, *Erlang and OTP in Action* (Manning) by Martin Logan, Eric Merritt, and Richard Carlsson, opens with a high-speed 72-page introduction to Erlang and then spends most of its time applying the Open Telecom Platform, Erlang’s framework for building upgradeable and maintainable concurrent applications.

If you want to focus on connecting Erlang to the Web, you should definitely also explore *Building Erlang Web Applications* (O’Reilly) by Zachary Kessin.

You’ll also want to visit the main Erlang website, <http://www.erlang.org/>, for updates, downloads, documentation, and more.

Are You Sure You Want Erlang?

Though they’ve been obscure for a long time, there’s a crowd of functional languages rising into greater popularity.

Four of them in particular—Clojure, Scala, F#, and Haskell—may be more appealing than Erlang if you have specific needs.

- Clojure and Scala run on the Java Virtual Machine (JVM), making them insanely portable, and they have access to Java libraries as a result. ClojureScript does the same with JavaScript, too. (Erjang makes it possible to run Erlang on the JVM, but it’s not a core part of the language.)
- F# runs on the .NET Common Language Runtime (CLR), making it very portable in the Microsoft ecosystem, and again, has access to .NET libraries.

- Haskell doesn't run on a virtual machine, but also offers a stronger type system and a different kind of discipline (and laziness).

Personally, I got my start with these concepts in XSLT. It's a very different kind of language meant for a specific domain of document transformation, but many of the same ideas flow through it.

You don't, of course, have to decide if Erlang is your life's dream now. You can learn concepts in Erlang and apply them elsewhere if it turns out to be a better idea for your work.

Erlang Will Change You

Before you go deeper, you should know that working in Erlang may irrevocably change the way you look at programs. Its combination of functional code, process-orientation, and distributed development may seem alien at first. However, once it sinks in, Erlang can transform the way you solve problems, and potentially make it difficult to return to other languages, environments, and programming cultures.

Conventions Used in This Book

The following typographical conventions are used in this book:

Italic

Indicates new terms, URLs, email addresses, filenames, and file extensions.

Constant width

Used for program listings, as well as within paragraphs to refer to program elements such as variable or function names, statements, and keywords.

Constant width bold

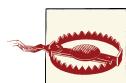
Shows commands or other text that should be typed literally by the user.

Constant width italic

Shows text that should be replaced with user-supplied values or by values determined by context.



This icon signifies a tip, suggestion, or general note.



This icon indicates a warning or caution.

A Note on Erlang Syntax

Erlang's syntax seems to be a sticking point for a lot of people. It doesn't look like the C family of languages. Punctuation is different and capitalization matters. Periods even get used as conclusions rather than connectors!

To me, Erlang syntax mostly feels natural, and I'm especially happy that it's different from the other languages I typically use. I make a lot fewer mix-ups that way.

Rather than dwell on syntax, I've chosen just to present it as it is. Comparing it to other languages doesn't seem likely to be helpful, especially when different readers may come from different programming backgrounds. Hopefully you will find Erlang syntax as pleasant as I do.

Using Code Examples

The examples in this book are meant to teach basic concepts in small bites. While you may certainly borrow code and reuse it as you see fit, you won't be able to take the code of this book and build a stupendous application instantly (unless perhaps you have an unusual fondness for calculating the speeds of falling objects). You should, however, be able to figure out the steps you need to take to build a great application.

You can download the code from the Examples link on the book's page at http://oreil.ly/introducing_erlang.

This book is here to help you get your job done. In general, you may use the code in this book in your programs and documentation. You do not need to contact us for permission unless you are reproducing a significant portion of the code. For example, writing a program that uses several chunks of code from this book does not require permission. Selling or distributing a CD-ROM of examples from O'Reilly books does require permission. Answering a question by citing this book and quoting example code does not require permission. Incorporating a significant amount of example code from this book into your product's documentation does require permission.

We appreciate, but do not require, attribution. An attribution usually includes the title, author, publisher, and ISBN. For example: "*Introducing Erlang*, by Simon St.Laurent (O'Reilly). Copyright 2013 Simon St.Laurent, 9781449331764."

If you feel your use of code examples falls outside fair use or the permission given above, feel free to contact us at permissions@oreilly.com.

Help This Book Grow

While I hope that you will enjoy reading this book and learn from it, I also hope that you can contribute to helping other readers learn Erlang here. You can help your fellow readers in a number of ways:

- If you find specific technical problems, bad explanations, or things that can be improved, please report them through the errata system at <http://oreilly.com/catalog/errata.csp?isbn=0636920025818>.
- If you like (or don't like) the book, please leave reviews. The most visible places to do so are on Amazon.com (or its international sites) and at the O'Reilly page for the book at http://oreil.ly/introducing_erlang. Detailed explanations of what worked and what didn't work for you (and the broader target audience of programmers new to Erlang) are helpful to other readers and to me.
- If you find you have much more you want to say about Erlang, please consider sharing it, whether on the Web, in a book of your own, in training classes, or in whatever form you find easiest.

I'll update the book for errata, and try to address issues raised in reviews. Even once the book is "complete," I may still add some extra pieces to it. If you purchased it as an ebook, you'll receive these updates for free at least up to the point where it's time for a whole new edition. I don't expect that new edition declaration to come quickly, however, unless the Erlang world changes substantially.

Hopefully this book will engage you enough to make you consider sharing.

Please Use It For Good

I'll let you determine what "good" means, but think about it. Please try to use Erlang's power for projects that make the world a better place, or at least not a worse place.

Safari® Books Online

 *Safari Books Online* is an on-demand digital library that delivers expert **content** in both book and video form from the world's leading authors in technology and business.

Technology professionals, software developers, web designers, and business and creative professionals use Safari Books Online as their primary resource for research, problem solving, learning, and certification training.

Safari Books Online offers a range of **product mixes** and pricing programs for **organizations**, **government agencies**, and **individuals**. Subscribers have access to thousands of books, training videos, and prepublication manuscripts in one fully searchable database

from publishers like O'Reilly Media, Prentice Hall Professional, Addison-Wesley Professional, Microsoft Press, Sams, Que, Peachpit Press, Focal Press, Cisco Press, John Wiley & Sons, Syngress, Morgan Kaufmann, IBM Redbooks, Packt, Adobe Press, FT Press, Apress, Manning, New Riders, McGraw-Hill, Jones & Bartlett, Course Technology, and dozens [more](#). For more information about Safari Books Online, please visit us [online](#).

How to Contact Us

Please address comments and questions concerning this book to the publisher:

O'Reilly Media, Inc.
1005 Gravenstein Highway North
Sebastopol, CA 95472
800-998-9938 (in the United States or Canada)
707-829-0515 (international or local)
707-829-0104 (fax)

We have a web page for this book, where we list errata, examples, and any additional information. You can access this page at http://oreil.ly/introducing_erlang.

To comment or ask technical questions about this book, send email to bookquestions@oreilly.com.

For more information about our books, courses, conferences, and news, see our website at <http://www.oreilly.com>.

Find us on Facebook: <http://facebook.com/oreilly>

Follow us on Twitter: <http://twitter.com/oreillymedia>

Watch us on YouTube: <http://www.youtube.com/oreillymedia>

Acknowledgments

Many thanks to Zachary Kessin for interesting me in Erlang in the first place, and to him and Francesco Cesarini for encouraging me to write this. Detailed feedback from Steve Vinoski and Fred Hebert has made it possible, I hope, for this book to get readers started on the right track. J. David Eisenberg and Chuck Ha helped make it especially possible for beginners to get started right, pointing out gaps and issues in my prose.

In particular, thanks to my wife Angelika for encouraging me to finish this, to my son Konrad for not throwing the printouts around too much, and to my daughter Sungiva for understanding that after I told her the story about Ned and Ernie, adventuring snakes, I needed to go back downstairs and work on this.

Getting Comfortable

Erlang has a funny learning curve for many people. It starts gently for a little while, then gets much much steeper as you realize the discipline involved, and then goes nearly vertical for a little while as you try to figure out how that discipline affects getting work done—and then it's suddenly calm and peaceful with a gentle grade for a long time as you reapply what you've learned in different contexts.

Before that climb, it's best to get comfortable in the sunny meadows at the bottom of the learning curve. Erlang's shell, its command-line interface, is a cozy place to get started and a good place to start figuring out what works and what doesn't work in Erlang. Its features will spare you headaches later, so settle in!

Installation

Erlang is officially available from <http://www.erlang.org/download.html>. For this book, I used Erlang/OTP R15B, attracted by this new feature:

Line number and filename information are now included in exception backtraces. This information will be pretty-printed in the shell and used in crash reports etc. In practice it will be much easier to find where something failed.

When you're getting started, extra information is useful, even in a language like Erlang that's designed to make debugging easier. Starting out with these features, of course, means that you can't tell newcomers that they need to appreciate how much harder it was for you getting started, but that's a minor loss.

If you're on Windows, it's easy. Download the Windows binary file, run the installer, and you're set. If you are a brave beginner tackling your first programming language, this is easily your best bet.

On Linux or Mac OS X, you may be able to download the source file and compile it. For me, on Mac OS X, I just had to unzip and untar it, and then, from the directory created by the untarring, run `./configure`, `make`, and `sudo make install`. However, that simple sequence works only if you have the right files previously installed, and can give you mysterious errors if they weren't. In particular, Apple's shift to the LLVM compiler in newer versions of XCode instead of GCC makes it less likely that GCC will be on newer Mac OS X systems, and Erlang needs GCC.

(You can also ignore the error about FOP, which Erlang uses to generate PDF documentation you can download elsewhere. Also, on newer Macs, you'll get an error at the end that wxWidgets doesn't work on 64-bit Mac OS X. For now, ignore this.)

If the compilation approach doesn't work or isn't for you, Erlang Solutions offers a number of installs at <http://www.erlang-solutions.com/section/132/download-erlang-otp>. Also, many different package managers (Debian, Ubuntu, MacPorts, brew, and so on) include Erlang. It may not be the very latest version, but having Erlang running is much better than not having Erlang running.



As CouchDB, which uses Erlang, spreads, Erlang is now part of the default installation on more systems, including Ubuntu.

Firing It Up

On Mac OS X or Linux, go to the command line and type `erl`. On Windows, go to the command line and type `werl`.

You'll see something like the following code sample, likely with a cursor next to the `1>` prompt.

```
Erlang R15B (erts-5.9) [source] [smp:2:2] [async-threads:0] [hipe] [kernel-poll:false]
Eshell V5.9 (abort with ^G)
1>
```

You're in Erlang!

First Steps: The Shell

Before moving on to the excitement of programming Erlang, it's always worth noting how to quit. The shell suggests `^G`, Ctrl-G, which will bring you to a mysterious (for now) user switch command. (Ctrl-C will bring you to a menu.) The simplest way to quit, allowing everything that might be running in the shell to exit normally, is `q()`.

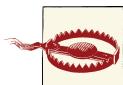
```
1> q().
ok
2> SimonMacBook:~ simonstl$
```

So what have you done here? You've issued a shell command, calling a function `q` that itself calls the `init:stop()` function built into Erlang. The period after the command tells Erlang you're done with the line. It reports back with `ok`, prints a new line number (it always does that after a period), and drops you back out to the regular command line, in this case a bash shell on your laptop.

If you had left off the period after `q()`, the results would look a little different. You'd have started a new line but the command count wouldn't update, so the line would still start with `1>`. When this happens, you can just type `.` and press Enter to finish your command.

```
1> q()
1> .
ok
2> SimonMacBook:~ simonstl$
```

Including the period at the end of the line will soon become second nature, but leaving it off can create a lot of confusion at the start.



Quitting Erlang with `q()`. turns off everything Erlang is doing, period. That's fine when you're working locally, but will become a bad idea when you're connecting to a remote shell. To quit the shell without the risk of shutting down the Erlang runtime on another system, try `Ctrl-G` and then entering `q`, followed by the Enter key.

Moving through Text

If you explore the shell, you'll find that many things work the way they do in other shells. The left and right arrow keys move you backward and forward through the line you're editing. Some of the key bindings echo the emacs text editor. `Ctrl-A` will take you to the beginning of a line, while `Ctrl-E` will take you back to the end of the line. If you get two characters in the wrong sequence, pressing `Ctrl-T` will transpose them.

Like most Unix shells, pressing the `Tab` key will make the shell try to autocomplete what you've written, though in this case it's looking for module or function names (you'll see them soon), not filenames.

Also, as you type closing parentheses or square brackets, the cursor will highlight the corresponding opening parenthesis or square bracket.

Moving through History

The up and down arrow keys run through the history, making it easy to reissue commands.

When you use the up and down arrows, the history will be broken down by newlines, not by periods, so if you left a period off in a prior command you'll need to add it again. If you want to see what's in the history, try `h()`. You can also specify how much history to keep around with `history(N)` and `results(N)`. You can tell Erlang to execute a given line again with `e(N)`, and reference a given result value with `v(N)`. Those line numbers can be useful!

Moving through Files

The Erlang shell does understand filesystems to some extent because you may need to move through them to reach the files that will become part of your program. The commands have the same names as Unix commands but are expressed as functions. The Erlang shell starts wherever you opened the shell, and you can figure out where that is with `pwd()`:

```
4> pwd().
/Users/simonstl
ok
5>
```

To change directories, use the `cd()` command, but you'll need to wrap the argument not only in parentheses but in quotes, preferably double quotes.

```
5> cd(..).
* 1: syntax error before: '..'
5> cd("..").
/Users
ok
6> cd("simonstl").
/Users/simonstl
ok
7>
```

You can look around with the `ls()` command, which will list files in the current directory if you give it no arguments, and list files in a specified directory if you give it one argument.

Doing Something

One of the easiest ways to get started playing with Erlang is to use the shell as a calculator. Unlike your typical command line, you can enter mathematical expressions and get useful results:

```
Eshell V5.9  (abort with ^G)
1> 2+2.
4
2> 27-14.
13
3> 35*42023943.
```

```
1470838005
4> 200/15.
13.33333333333334
5> 200 div 15.
13
6> 200 rem 15.
5
7> 3*(4+15).
57
```

The first three operators are addition(+), subtraction(-), and multiplication(*), which work the same way whether you're working with integer values or floating points. The fourth, /, supports division where you expect a floating point (a number with a decimal part) result. If you want an integer result (and have integer arguments), use the `div` operator instead, with `rem` to get the remainder, as shown on lines 5 and 6. Parentheses let you modify the order in which operators are processed, as shown on line 7. (The normal order of operations is listed in [Appendix A](#).)

Erlang will accept integers in place of floats, but floats are not always welcome where integers are used. If you need to convert a floating point number to an integer, you can use the `round()` built-in function:

```
8> round(200/15).
13
```

The `round()` function drops the decimal part of the number. If the decimal part was greater than or equal to .5, it increases the integer part by 1, rounding up. If you'd rather just drop the decimal part completely, use the `trunc()` function, which effectively always rounds down.

You can also refer to a previous result by its line number using `v()`. For example:

```
9> 4*v(8).
52
```

The result on line 8 was 13, and 4×13 is 52.



If you're feeling adventurous, you can use negative numbers to reference prior results. `v(-1)` is the previous result, `v(-2)` is the result before that, and so on.

Calling Functions

If you want to do more powerful calculations, Erlang's math module offers pretty much the classic set of functions supported by a scientific calculator. They return floating point values. The constant `pi` is available as a function, `Math:pi()`. Trigonometric,

logarithmic, exponential, square root, and (except on Windows) even the Gauss error functions are readily available. (The trigonometric functions take their arguments in radians, not degrees, so be ready to convert if necessary.) Using these functions is a little verbose because of the need to prefix them with `math`:, but it's still reasonably sane.

For example, to get the sine of zero radians, you'd write:

```
1> math:sin(0).  
0.0
```

Note that it's `0.0`, not just `0`, indicating that the number is floating point.

To calculate the cosine of pi and 2π radians, you'd write:

```
2> math:cos(math:pi()).  
-1.0  
3> math:cos(2*math:pi()).  
1.0
```

To calculate 2 taken to the 16th power, you'd use:

```
4> math:pow(2,16).  
65536.0
```

The full set of mathematical functions supported by Erlang's `math` module is listed in [Appendix A](#).

Numbers in Erlang

Erlang recognizes two kinds of numbers: integers and floating-point numbers (often called floats). It's easy to think of integers as "whole numbers," with no decimal part, and floats as "decimal numbers," with a decimal point and some value (even if it's `0`) to the right of the decimal. `1` is an integer, `1.0` is a floating-point number.

However, it's a little trickier than that. Erlang stores integers and floats in a very different way. Erlang lets you store massive numbers as integers, but whether they're big or small, they are always precise. You don't need to worry about their values being off by just a little.

Floating-point numbers, on the other hand, cover a wide range of numbers but with limited precision. Erlang uses the 64-bit IEEE 754-1985 "double precision" representation. This means that it keeps track of about 15 decimal digits plus an exponent. It can also represent some large numbers—powers up to positive or negative 308 are available—but because it tracks only a limited number of digits, results will vary a little more than may seem convenient, especially when you want to do comparisons.

```
1> 3487598347598347598437583475893475843749245.0.  
3.4875983475983474e42
```

```
2> 2343243.345435893850234543339545.  
2343243.3454358936  
3> 0.0000000000000000000000000000000023432432432432234232324.  
2.3432432432432235e-30
```

As you can see, some digits get left behind, and the overall magnitude of the number represented with an exponent.

When you enter floating point numbers, you must always also have at least one number to the left of the decimal point, even if it's zero. Otherwise Erlang reports a syntax error—it doesn't understand what you're doing.

```
4> .0000000000000000000000000000000023432432432432234232324.  
* 1: syntax error before: 23432432432432234232324
```

You can also write floats using the digits plus exponent notation:

```
7> 2.923e127.  
2.923e127  
8> 7.6345435e-231.  
7.6345435e-231
```

Floats' lack of precision can cause anomalous results. For example, the sine of zero is zero, and the sine of pi is also zero. However, if you calculate this in Erlang, you won't quite get to zero with the float approximation Erlang provides for pi:

```
1> math:sin(0).  
0.0  
2> math:sin(math:pi()).  
1.2246467991473532e-16
```

If Erlang's representation of pi went further, and its calculations of pi went further, the result for line 2 would be closer to zero.

If you need to keep track of money, integers are going to be a better bet. Use the smallest available unit—cents for US dollars, for instance—and remember that those cents are 1/100 of a dollar. (Financial transactions can go to much smaller fractions, but you'll still want to represent them as integers with a known multiplier.) For more complex calculations, though, you'll want to use floats, and just be aware that results will be imprecise.

If you need to do calculations on integers using a base other than 10, you can use *Base#Value* notation. For example, if you wanted to specify the binary value of 1010111, you could write:

```
3> 2#1010111.  
87
```

Erlang reports back with the base 10 value of the number. Similarly, you can specify hexadecimal numbers by using 16 instead of 2:

4> 16#cafe.

51966

Erlang lets you use either upper- or lower-case for hexdecimal numbers - 16#CAFE and 16#CaFe also produce 51966. You aren't limited to the traditional binary (base 2), octal (base 8), and hexadecimal (base 16) choices. If you want to work in base 18, or any base up to 36, you can:

5> 18#gaffe.

1743080



Why might you use base 36? It's an extremely easy way to create keys that look like a combination of letters and numbers, but resolve neatly to numbers. The 6-digit codes airlines use to identify tickets, like G6ZV1N, are easily treated as base 36. (However, they usually leave out some digits and letters that are easily confused, such as -0 and 0, and 1 and l.)

To make any of these numbers negative just put a minus sign (-) in front of them. This works with normal integers, *Base#Value* notation, and floats:

6> -1234.

-1234

7> -16#cafe.

-51966

8> -2.045234324e6.

-2045234.324

Working with Variables in the Shell

The v() function lets you refer to the results of previous expressions, but it's not exactly convenient to keep track of result numbers, and the v() function works only in the shell. It isn't a general-purpose mechanism. A more reasonable solution stores values with textual names, creating variables.

Erlang variable names begin with a capital letter or an underscore. Normal variables start with a capital letter, whereas underscores start "don't care" variables. For now, stick with normal variables. You assign a value to a variable using a syntax that should be familiar from algebra or other programming languages, here with N as the variable:

1> N=1.

1

To see the value of a variable, just type its name.

2> N.

1

To see Erlang protest at your rude behavior, try assigning the variable a new value:

```
3> N=2.  
** exception error: no match of right hand side value 2  
4> N=N+1.  
** exception error: no match of right hand side value 2
```

What's happening here? Erlang expects the righthand side of an expression, after the `=`, to match the lefthand side. It's willing to make that happen if a variable on the left side isn't bound yet, as was the case with `N=1` in the first line. However, once the variable `N` is set to 1, Erlang interprets `N=2` as `1=2`, which it won't accept. `N=N+1` also evaluates to `1=2`, and doesn't work. Erlang's *single assignment* model, where each variable can be assigned a value only once in a given context, imposes discipline whose value you will see in later chapters.

Erlang expressions work like algebra, where `N` never equals `N+1`. It just can't happen that way. However, once you've set `N` to 1, it's fine to try expressions that also come to one:

```
5> N=2-1.  
1  
6> N=15 div (3*5).  
1
```

This will get much more important when you start to take advantage of Erlang's pattern matching capabilities. You can also write the following:

```
7> 1=N.  
1
```

Erlang won't attempt to bind any variables when they appear on the right side of the equals sign, and this just effectively asks Erlang to compare 1 to 1. Try it with 2, however, and Erlang complains that there isn't a match; 2 does not equal 1:

```
8> 2=N.  
** exception error: no match of right hand side value 1
```

You can also use bound variables in calculations, for example to create new bound variables, here's one called `Number`:

```
9> Number=N*4+N.  
5  
10> 6*Number.  
30
```

When you assign a value to a variable, you should make sure that all the calculations are on the right side of the equals sign. Even though I know that `M` should be 6 when `2*M = 3*4`, Erlang doesn't:

```
11> 2*M=3*4.  
* 1: illegal pattern
```

The shell will remember your variables until you quit or tell it to forget them. Code in Erlang functions doesn't forget, until the functions stop running.

Seeing Your Bound Variables

After a while poking around the shell using it as a calculator (try it!), you may find you've forgotten what variables you've already bound. If you need a reminder, the `b()` shell command can help:

```
11> b().  
N = 1  
Number = 5  
ok
```

Clearing Bound Variables in the Shell

In the shell, and only in the shell, you can clear all variable bindings and you can clear specific variable bindings. This may prove useful after an egregious typo or to reset your console for new calculations, but it isn't an option you'll have in regular code.

To clear a specific variable, removing its binding and letting you set a new value, use the `f()` function, giving the variable name as an argument:

```
12> f(N).  
ok  
13> b().  
Number = 5  
ok  
14> N=2.  
2
```

To clear all the bound variables in the shell, just call `f()` with no arguments.

```
15> b().  
N = 2  
Number = 5  
ok  
16> f().  
ok  
17> b().  
ok
```

They all disappeared.

Before moving on to the next chapter, which will introduce modules and functions, spend some time playing in the Erlang shell. The experience, even at this simple level, will help you move forward. Use variables, and see what happens with large integers. Erlang supports large numbers very well. Try mixing numbers with decimal values (floats) and integers in calculations, and see what happens. Nothing should be difficult yet, though I suspect the idea of variables that don't change values gives you a hint of what's to come.

Functions and Modules

Like most programming languages, Erlang lets you define functions to help you represent repeated calculations. While Erlang functions can become complicated, they start out reasonably simple.

Fun with fun

You can create functions in the Erlang shell using the appropriately named `fun`. For example, to create a function that calculates the velocity of a falling object based on the distance it drops in meters, you could create the following:

```
1> FallVelocity = fun(Distance) -> math:sqrt(2 * 9.8 * Distance) end.  
#Fun<erl_eval.6.111823515>
```

You can read that as a pattern match that binds the variable `FallVelocity` to a function that takes an argument of `Distance`. The function returns (I like to read the `->` as yields) the square root of 2 times a gravitational constant for Earth of 9.8 m/s, times `Distance` (in meters). Then the function comes to an `end`, and a period closes the statement.



If you want to include multiple statements in a `fun`, separate them with commas, like `FallVelocity = fun(Distance) -> X = (2 * 9.8 * Distance), math:sqrt(X) end.`

The return value in the shell, `#Fun<erl_eval.6.111823515>`, isn't especially meaningful by itself, but it tells you that you've created a function and didn't just get an error. If you want a slightly more detailed sign that Erlang understood you, you can use the `b()` shell command to see what it thinks:

```
2> b().  
FallVelocity =  
    fun(Distance) ->  
        math:sqrt(2 * 9.8 * Distance)  
    end  
ok
```

Conveniently, binding the function to the variable `FallVelocity` lets you use that variable to calculate the velocity of objects falling to Earth:

```
3> FallVelocity(20).  
19.79898987322333  
4> FallVelocity(200).  
62.609903369994115  
5> FallVelocity(2000).  
197.9898987322333
```

If you want those meters per second in miles per hour, just create another function. You can copy-and-paste the earlier results into it (as I did here), or pick shorter numbers:

```
6> Mps_to_mph = fun(Mps) -> 2.23693629 * Mps end.  
#Fun<erl_eval.6.111823515>  
7> Mps_to_mph(19.79898987322333).  
44.289078952755766  
8> Mps_to_mph(62.609903369994115).  
140.05436496173314  
9> Mps_to_mph(197.9898987322333).  
442.89078952755773
```

I think I'll stay away from 2000 meter drops. Prefer the fall speed in kilometers per hour?

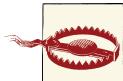
```
10> Mps_to_kph = fun(Mps) -> 3.6 * Mps end.  
#Fun<erl_eval.6.111823515>  
11> Mps_to_kph(19.79898987322333).  
71.27636354360399  
12> Mps_to_kph(62.609903369994115).  
225.3956521319788  
13> Mps_to_kph(197.9898987322333).  
712.76363543604
```

You can also go straight to your preferred measurement by nesting the following calls:

```
14> Mps_to_kph(FallVelocity(2000)).  
712.76363543604
```

However you represent it, that's really fast, though air resistance will slow those down a lot in reality.

This is handy for repeated calculations, but you probably don't want to push this kind of function use too far in the shell, as flushing your variables or quitting the shell session makes your functions vanish.



If you get an error that looks like `** exception error: no function clause matching erl_eval:'-inside-an-interpreted-fun-'(value),` check your capitalization. It may take a while to get used to capitalizing all your variables, including arguments in functions.

Defining Modules

Most Erlang programs, except things like the preceding simple calculations, define their functions in compiled modules rather than in the shell. Modules are a more formal place to put programs, and they give you the ability to store, encapsulate, share, and manage your code more effectively.

Each module should go in its own file, with an extension of `.erl`. You should use `name_of_module.erl`, where `name_of_module` is the name you specify inside of the module file. **Example 2-1**, which you can find in the examples archive at `ch02/ex1-drop`, shows what a module, `drop.erl`, containing the functions previously defined might look like.

Example 2-1. Module for calculating and converting fall velocities

```
-module(drop).
-export([fall_velocity/1, mps_to_mph/1, mps_to_kph/1]).

fall_velocity(Distance) -> math:sqrt(2 * 9.8 * Distance).

mps_to_mph(Mps) -> 2.23693629 * Mps.

mps_to_kph(Mps) -> 3.6 * Mps.
```

There are two key kinds of information in this module. At the top, the `-module` and `-export` directives tell the compiler key things about the module—its name and which functions it should make visible to other code that uses this module. The `-export` directive gives a list of functions that should be made visible—not just their names, but their *arity*, the number of arguments they take. Erlang considers functions with the same name but different arity to be different functions.

All of the code in a module must be contained in functions.

Below the directives is a set of expressions defining functions, which look similar to the fun declarations used earlier but not quite the same. The function names start with lowercase, not uppercase, and the syntax is slightly different. `fun` and `end` don't appear, and the function name is immediately followed by parentheses containing a set of arguments.



If you get errors like "drop.erl:2: bad function arity drop.erl:6: syntax error before: Fall_velocity", it's probably because you didn't convert the names from your fun/s so they start with a lowercase letter.

How do you make this actually do something?

It's time to start compiling Erlang code. The shell will let you compile modules and then use them immediately. The `c()` function lets you compile code. You need to be in the same directory as the file, whether you started the Erlang shell from there or navigated there with the commands shown in the previous chapter. You don't need to (and shouldn't) include the `.erl` file extension in the name you pass to `c()`, though you can specify directory paths.

```
1> ls().
drop.erl
ok
2> c(drop).
{ok,drop}
3> ls().
drop.beam    drop.erl
ok
```

Line 1 checks to see if the `drop.erl` source file is there, and you see the directory listing. Line 2 actually compiles it, and line 3 shows that a new file, `drop.beam`, is now available. Now that you have `drop.beam`, you can call functions from the module. You need to prefix those calls with `drop`, as shown in lines 4 and 5 of the following code.

```
4> drop:fall_velocity(20).
19.79898987322333
5> drop:mps_to_mph(drop:fall_velocity(20)).
44.289078952755766
```

It works the same as its predecessors, but now you can quit the shell, return, and still use the compiled functions.

```
6> q().
ok
$ erl
Erlang R15B (erts-5.9) [source] [smp:8:8] [async-threads:0] [hipe] [kernel-poll:false]

Eshell V5.9 (abort with ^G)
1> drop:mps_to_mph(drop:fall_velocity(20)).
44.289078952755766
```

Most Erlang programming (beyond tinkering in the shell) is creating functions in modules and connecting them into larger programs.

Erlang Compilation and the Runtime System

When you write Erlang in the shell, it has to interpret every command, whether or not you've written it before. When you tell Erlang to compile a file, it converts your text into something it can process without having to re-interpret all the text, tremendously improving efficiency when you run the code.

That "something it can process," in Erlang's case, is a BEAM file. It contains code that the BEAM processor, a key piece of the Erlang Runtime System (ERTS) can run. BEAM is Bodgan's Erlang Abstract Machine, a virtual machine that interprets optimized BEAM code. This may sound slightly less efficient than the traditional compilation to machine code that runs directly on the computer, but it resembles other virtual machines. (Oracle's Java Virtual Machine (JVM) and Microsoft's .NET Framework are the two most common virtual machines.)

Having its own virtual machine and runtime system lets Erlang optimize some key things, making it easier to build applications that scale reliably. Its process scheduler simplifies distributing work across multiple processors in a single computer. You don't have to think about how many processors your application might get to use—you just write independent processes, and Erlang spreads them out. Erlang also manages input and output in its own way, avoiding connection styles that block other processing. The virtual machine also uses a garbage collection strategy that fits its style of processing, allowing for briefer pauses in program execution. (Garbage collection releases memory that processes needed at one point but are no longer using.)

When you create and deliver Erlang programs, you will be distributing them as a set of compiled BEAM files. You don't need to compile each one from the shell as we're doing here, though. `erlc` will let you compile Erlang files directly and combine that compilation into `make` tasks and similar things, whereas `escript` can compile or interpret and run Erlang code from outside of the Erlang shell.

From Module to Fun

If you like the style of code that `fun` allowed but also want your code stored more reliably in modules where it's easier to debug, you can get the best of both worlds by using the `fun` keyword to refer to a function you've already defined. To do that, you *don't* use parentheses after `fun`, and give the module name, function name, and arity.

```
1> F_v = fun drop:fall_velocity/1.  
#Fun<drop.fall_velocity.1>  
2> F_v(20).  
19.79898987322333
```

You can also do this within code in a module, and if you’re referring to code in the same module, you can leave off the module name preface. (In this case, that would mean leaving off `drop:` and just using `fall_velocity/1`.)

Functions and Variable Scope

Erlang lets you bind a variable only once, but you might call a function many times over the course of a program. Doesn’t that mean the same variable will be bound many times?

Yes, it will be bound many times but always in separate contexts. Erlang doesn’t consider multiple calls to the same function to be the same thing. It starts with a fresh set of unassigned variables every time you call that function.

Similarly, Erlang doesn’t worry if you use the same variable name in different functions or function clauses. They aren’t going to be called in the same context at the same time, so there isn’t a collision.

The place you need to avoid re-assigning values to an already bound variable is within a given path through a given function. As long as you don’t try to reuse a variable in a given context, you shouldn’t have to worry.

Module Directives

By default, modules have very thick walls, and everything inside of them is considered private. Everything going in or out of the module needs a pass to do so, and you grant those passes through module directives (sometimes called module attributes).

The example above showed two module directives—`-module` and `-export`. The `-module` directive sets the name for the module, which outside code will need to know in order to call the functions. The `-export` directive specifies which functions that outside code can reach.

The `drop` module currently mixes two different kinds of functions. The `fall_velocity/1` function fits the name of the module, `drop`, very well, providing a calculation based on the height from which an object falls. The `mps_to_mph/1` and `mps_to_kph/1` functions, however, aren’t about dropping. They are generic measurement conversion functions that are useful in other contexts and really belong in their own module. [Example 2-2](#) and [Example 2-3](#), both in `ch02/ex2-combined`, show how this might be improved.

Example 2-2. Module for calculating fall velocities

```
-module(drop).
-export([fall_velocity/1]).

fall_velocity(Distance) -> math:sqrt(2 * 9.8 * Distance).
```

Example 2-3. Module for converting fall velocities

```
-module(convert).  
-export([mps_to_mph/1, mps_to_kph/1]).  
  
mps_to_mph(Mps) -> 2.23693629 * Mps.  
  
mps_to_kph(Mps) -> 3.6 * Mps.
```

Next, you can compile them, and then the separated functions are available for use:

```
Eshell V5.9 (abort with ^G)  
1> c(drop).  
{ok,drop}  
2> c(convert).  
{ok,convert}  
3> ls().  
convert.beam convert.erl drop.beam drop.erl  
  
ok  
4> convert:mps_to_mph(drop:fall_velocity(20)).  
44.289078952755766
```

That reads more neatly, but how might this code work if a third module needed to call those functions? Modules that call code from other modules need to specify that explicitly. **Example 2-4**, in *ch02/ex3-combined*, shows a module that uses functions from both the drop and convert modules.

Example 2-4. Module for combining drop and convert logic

```
-module(combined).  
-export([height_to_mph/1]).  
  
height_to_mph(Meters) -> convert:mps_to_mph(drop:fall_velocity(Meters)).
```

That looks much like the way you called it from the Erlang shell, but if you have a lot of calls to external modules, that can get verbose quickly. The `-import` directive, shown in **Example 2-5**, lets you simplify your code, though it comes with a possible risk of confusing other people who think the imported functions must be defined within this module. (You can find this in *ch02/ex4-combined*.)

Example 2-5. Module for combining drop and convert logic using import

```
-module(combined).  
-export([height_to_mph/1]).  
-import(drop, [fall_velocity/1]).  
-import(convert, [mps_to_mph/1]).  
  
height_to_mph(Meters) -> mps_to_mph(fall_velocity(Meters)).
```

For now, it's probably best to know about the `-import` directive so you can read other people's code, but not to use it unless you just can't resist. It can make it harder to figure where bugs are coming from, which may cost you more time than the extra typing.

Erlang includes one other directive that's similarly convenient but not best practice to use: `-compile(export_all)`. That directive tears down the module wall, making all functions available for external calls. In a module where everything is supposed to be public, that might save you typing out all the functions and all the arities of your module. However, it also means anyone can call anything in your code, exposing a lot more surface area for misunderstandings and complex debugging. If you just can't resist, it's available, but try to resist.



You can also make up your own user directives. `-author(Name)` and `-date(Date)` are commonly used. If you make up your own directives, they can have only one argument. If you spend enough time in Erlang, you'll also encounter the following: `-behaviour(Behaviour)`, `-record(Name, Fields)`, and `-vsn(Version)`.

Documenting Code

Your programs can run perfectly well without documentation. Your projects, however, will have a much harder time.

While programmers like to think they write code that anyone can look at and sort out, the painful reality is that code even a little more complicated than that shown in the previous examples can prove mystifying to other developers. If you step away from code for a while, the understanding you developed while programming it may have faded, and even your own code can seem incomprehensible.

The simplest way to add more explicit explanations to your code is to insert comments. You can start a comment with %, and it runs to the end of the line. Some comments take up an entire line, while others are short snippets at the end of a line. [Example 2-6](#) shows both varieties of comments.

Example 2-6. Comments in action

```
-module(combined).  
-export([height_to_mph/1]).    % there will be more soon!  
  
%%% combines logic from other modules into a convenience function.  
height_to_mph(Meters) -> convert:mps_to_mph(drop:fall_velocity(Meters)).
```

The Erlang compiler will ignore all text between the % sign and the end of the line, but humans exploring the code will be able to read them.

Why are there multiple percent signs at the start of the line? The Erlang Emacs mode and many other Erlang tools expect the number of percent signs to indicate levels of indentation. Three percent signs (%%%) means that the comment will be formatted flush left, two percent signs (%%) means the comment is indented with surrounding code, and a single percent sign (%) is used for comments on the end of a line.

Informal comments are useful, but developers have a habit of including comments that help them keep track of what they're doing while they're writing the code. Those comments may or may not be what other developers need to understand the code, or even what you need when you return to the code after a long time away. More formal comment structures may be more work than you want to take on in the heat of a programming session, but they also force you to ask who might be looking at your code in the future and what they might want to know.

Erlang includes a documentation system called *EDoc*, which converts comments placed in the code into navigable HTML documentation. It relies on specially formatted comments, a directive, and occasionally an extra file to provide structured information about your modules and application.

Documenting Modules

The modules in this chapter are very simple so far, but there is enough there to start documenting, as shown in the files at *ch02/ex5-docs*. Example 2-7 presents the *drop* module with more information about who created it and why.

Example 2-7. Documented module for calculating fall velocities

```
%% @author Simon St.Laurent <simonstl@simonstl.com> [http://simonstl.com]
%% @doc Functions calculating velocities achieved by objects
%% dropped in a vacuum.
%% @reference from <a href= "http://shop.oreilly.com/product/0636920025818.do" >Introducing Erlang</a>,
%% O'Reilly Media, Inc., 2012.
%% @copyright 2012 by Simon St.Laurent
%% @version 0.1

-module(drop).
-export([fall_velocity/1]).

fall_velocity(Distance) -> math:sqrt(2 * 9.8 * Distance).
```

Erlang can build the files for you using the EDoc `file` function:

```
Eshell V5.9  (abort with ^G)
1> edoc:files(["drop.erl"], [{dir, "doc"}]).
ok
```

You'll now have a collection of files in the *doc* subdirectory. If you open *drop.html* in a browser, you'll see something like Figure 2-1.

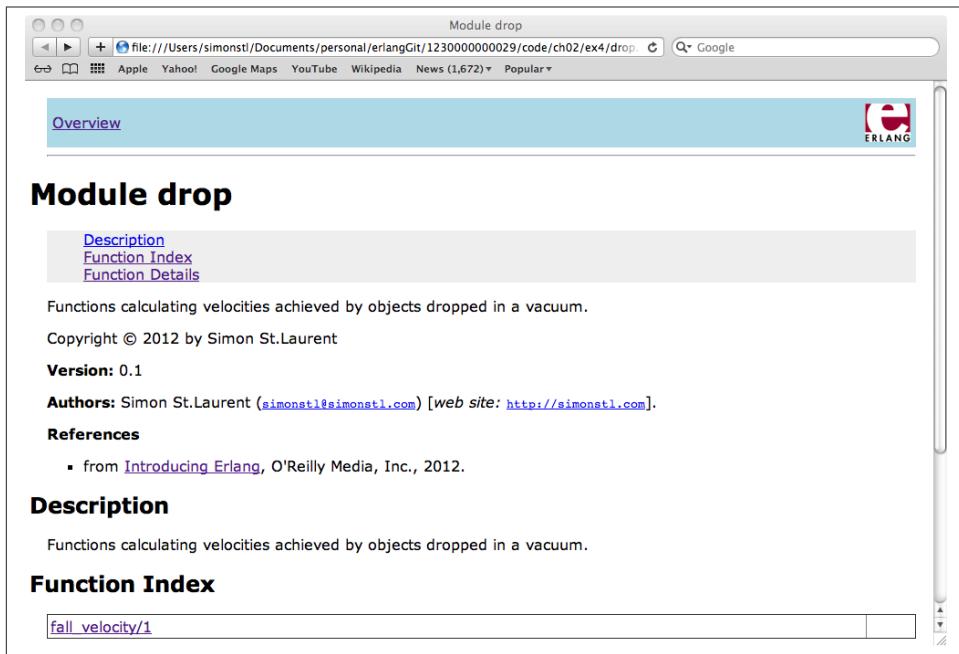


Figure 2-1. Module documentation generated from the drop.erl file

All of that metadata is great, and it can be gratifying to see your name “in lights.” However, unless you have a complex story to tell about your module as a whole, it’s likely that the core of the documentation will appear at the function level.

Documenting Functions

The drop module contains one function: `fall_velocity/1`. You probably know that it takes a distance in meters and returns a velocity in meters per second for an object dropped in a vacuum on Earth, but the code doesn’t actually say that. [Example 2-8](#) shows how to fix that with EDoc comments and the `@doc` tag.

Example 2-8. Documented function for calculating fall velocities

```
%<code>%@doc Calculates the velocity of an object falling on Earth</code>
%<code>% as if it were in a vacuum (no air resistance). The distance is</code>
%<code>% the height from which the object falls, specified in meters,</code>
%<code>% and the function returns a velocity in meters per second.</code>
```

```
fall_velocity(Distance) -> math:sqrt(2 * 9.8 * Distance).
```

Figure 2-2 shows the result, which is considerably more helpful than the previous blank space around the function. It neatly takes the first sentence of the information following `@doc` and put it in the index, using the whole description for the Function Details section. You can also use XHTML markup in the `@doc` section.

The screenshot shows a web browser window with the title "Module drop". The address bar indicates the page is at "file:///Users/simonstl/Documents/personal/erlangGit/1230000000029/code/ch02/ex4/drop.erl". The page content is as follows:

Version: 0.1

Authors: Simon St.Laurent (simonstl@simonstl.com) [**web site:** <http://simonstl.com>].

References

- from [Introducing Erlang](#), O'Reilly Media, Inc., 2012.

Description

Functions calculating velocities achieved by objects dropped in a vacuum.

Function Index

[fall_velocity/1](#) Calculates the velocity of an object falling on Earth as if it was in a vacuum (no air resistance).

Function Details

fall_velocity/1

```
fall_velocity(Distance) -> any()
```

Calculates the velocity of an object falling on Earth as if it was in a vacuum (no air resistance). The distance is the height from which the object falls, specified in meters, and the function returns a velocity in meters per second.

[Overview](#) 

Generated by EDoc, Jun 17 2012, 16:45:32.

Figure 2-2. Function documentation generated from the drop.erl file

That's a major improvement, but what if a user specifies “twenty” meters instead of 20 meters? Because Erlang doesn't worry much about types, the Erlang code doesn't say that the value for `Distance` has to be a number or the function will return an error.

You can add a directive, `-spec`, to add that information. It's a little strange, as in some ways it feels like a duplicate of the method declaration. In this case, it's simple, as shown in [Example 2-9](#).

Example 2-9. Documented function for calculating fall velocities

```
%% @doc Calculates the velocity of an object falling on Earth  
%% as if it was in a vacuum (no air resistance). The distance is  
%% the height from which the object falls, specified in meters,  
%% and the function returns a velocity in meters per second.
```

```
-spec(fall_velocity(number()) -> number()).
```

```
fall_velocity(Distance) -> math:sqrt(2 * 9.8 * Distance).
```

EDoc will combine the types specified in the `-spec` directive with the parameter names in the actual function declaration to produce the documentation shown in [Figure 2-3](#).

Module drop

file:///Users/simonstl/Documents/personal/erlangGit/123000000029/code/ch02/ex4/drop

Function Index

`fall_velocity/1` Calculates the velocity of an object falling on Earth as if it was in a vacuum (no air resistance).

Function Details

fall_velocity/1

`fall_velocity(Distance::number()) -> number()`

Calculates the velocity of an object falling on Earth as if it was in a vacuum (no air resistance). The distance is the height from which the object falls, specified in meters, and the function returns a velocity in meters per second.

Overview

Generated by EDoc, Jun 17 2012, 16:54:27.

ERLANG

Figure 2-3. EDoc documentation with type information

This chapter has really demonstrated only the `number()` type, which combines `integer()` and `float()`. [Appendix A](#) includes a full list of types.

Documenting Your Application

Sometimes you want information like the author and copyright data to appear in every module, often when it varies from module to module. Other times that becomes clutter, and it's easier to put it into one place where it applies to all of your modules.

You can create an `overview.edoc` file in your project's `doc` directory. Its content looks much like the markup used in the modules, but because it isn't mixed with code, you don't need to preface every line with `%%`. The `overview.edoc` file for this project might look like [Example 2-10](#).

Example 2-10. Documented module for calculating fall velocities

```
@author Simon St.Laurent <simonstl@simonstl.com> [http://simonstl.com]
@doc Functions for calculating and converting velocities.
@reference from <a href= "http://shop.oreilly.com/product/0636920025818.do"
>Introducing Erlang</a>, O'Reilly Media, Inc., 2012.
@copyright 2012 by Simon St.Laurent
@version 0.1
```

Now, if you re-generate documentation and click on the Overview link, you'll see something like [Figure 2-4](#).

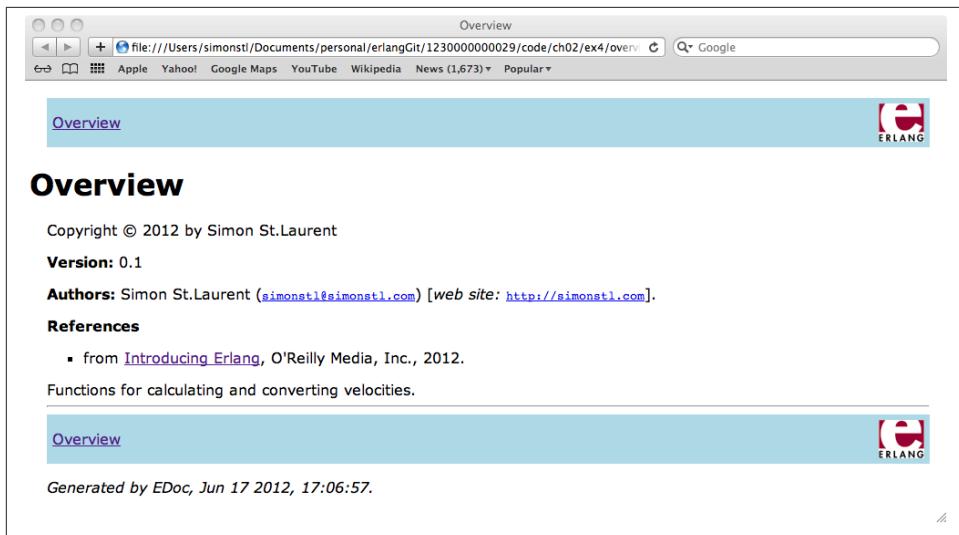


Figure 2-4. EDoc documentation with type information

If you create similar documentation in each of the Erlang files and run `edoc:files(["drop.erl", "convert.erl", "combined.erl"]).` in the Erlang shell, EDoc will build a neat if somewhat plain set of frame-based documentation for your application, as shown in Figure 2-5.

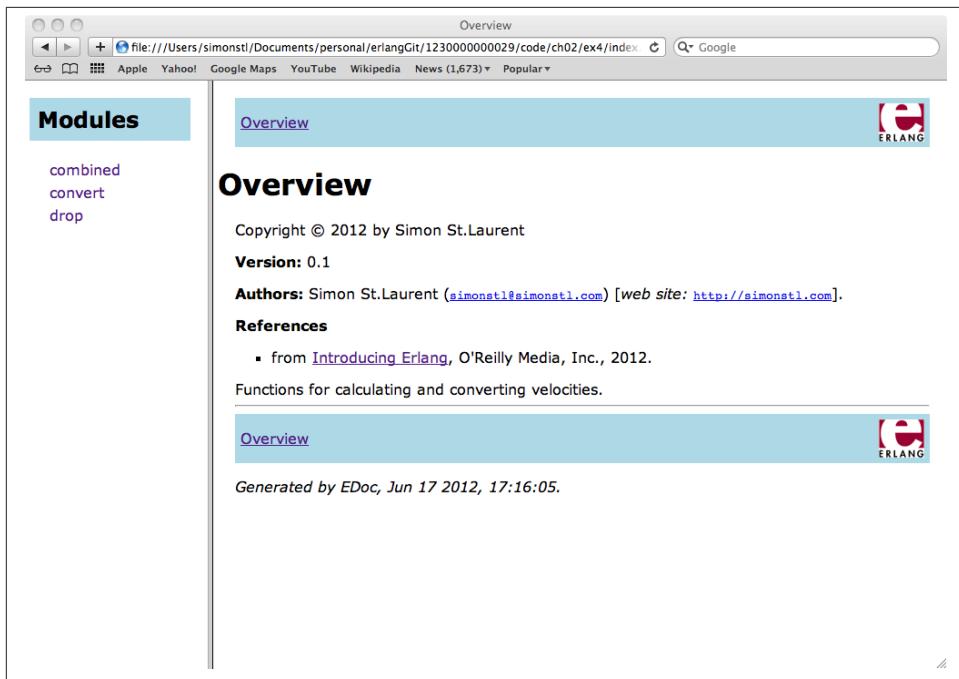


Figure 2-5. The opening to the complete set of module documentation

This is just an introduction to EDoc. For more, see Chapter 18 of *Erlang Programming*, where you can learn about fun things like the @todo tag.

Atoms, Tuples, and Pattern Matching

Erlang programs are at heart a set of message requests and tools for processing them. Erlang provides tools that simplify the efficient handling of those messages, letting you create code that is readable (to programmers at least) while still running efficiently when you need speed.

Atoms

Atoms are a key structure in Erlang. Technically they're just another type of data, but it's hard to overstate their impact on Erlang programming style.

Usually, atoms are bits of text that start with a lowercase letter, like `ok` or `earth`. They can also contain (though not start with) underscores (`_`) and at symbols (`@`), like `this_is_a_short_sentence` or `me@home`. If you want more freedom to start with uppercase letters or use spaces, you can put them in single quotes, like `'Today is a good day'`. Generally, the one word lowercase form is easier to read.

Atoms have a value—it's the same as their text. (Remember, the period after `hello` isn't part of the atom—it ends the expression.)

```
1> hello.  
hello
```

That's not very exciting in itself. What makes atoms exciting is the way that they can combine with other types and Erlang's pattern matching techniques to build simple but powerful logical structures.

Pattern Matching with Atoms

Erlang used pattern matching to make the examples in [Chapter 2](#) work, but it was very simple. The name of the function was the one key piece that varied, and as long as you provided a numeric argument Erlang knew what you meant. Erlang's pattern matching offers much more sophisticated possibilities, however, allowing you to match on arguments as well as on function names.

For example, suppose you want to calculate the velocity of falling objects not just on Earth, where the gravitational constant is 9.8 meters per second squared, but on Earth's moon, where it is 1.6 meters per second squared, and on Mars, where it is 3.71 meters per second squared. [Example 3-1](#), which you can find in *ch03/ex1-atoms*, shows one way to build code that supports this.

Example 3-1. Pattern matching on atoms as well as function names

```
-module(drop).
-export([fall_velocity/2]).  
  
fall_velocity(earth, Distance) -> math:sqrt(2 * 9.8 * Distance);  
fall_velocity(moon, Distance) -> math:sqrt(2 * 1.6 * Distance);  
fall_velocity(mars, Distance) -> math:sqrt(2 * 3.71 * Distance).
```

It looks like the `fall_velocity` function gets defined three times here, and it certainly provides three processing paths for the same function. However, because those definitions are separated with semicolons, they are treated as choices—selected by pattern-matching—rather than duplicate definitions. As in English, these pieces are called *clauses*.



If you use periods instead of semicolons, you'll get errors like `drop.erl: 5: function fall_velocity/2 already defined.`

Once you have this, you can calculate velocities for objects falling a given distance on Earth, the Earth's moon, and Mars:

```
1> c(drop).
{ok,drop}
2> drop:fall_velocity(earth,20).
19.79898987322333
3> drop:fall_velocity(moon,20).
8.0
4> drop:fall_velocity(mars,20).
12.181953866272849
```

You'll quickly find that atoms are a critical component for writing readable Erlang code.

Atomic Booleans

Two of Erlang's atoms have special properties: `true` and `false`, representing the boolean logic values of the same names. Erlang will return these atoms if you ask it to compare something:

```
1> 3<2.  
false  
2> 3>2.  
true  
3> 10 == 10.  
true
```

Erlang also has special operators that work on these atoms (and on comparisons that resolve to these atoms):

```
1> true and true.  
true  
2> true and false.  
false  
3> true or false.  
true  
4> false or false.  
false  
5> true xor false.  
true  
6> true xor true.  
false  
7> not true.  
false
```

The `and`, `or`, and `xor` operators both take two arguments. For `and`, the result is true if and only if the two arguments are true. For `or`, the result is true if at least one of the arguments is true. For `xor`, exclusive or, the result is true if one but not both arguments is true. In all other cases they return false. If you're comparing expressions more complicated than `true` and `false`, it's wise to put them in parentheses.



There are two additional operators for situations where you don't want or need to evaluate all of the arguments. The `andalso` operator behaves like `and` but evaluates the second argument only if the first one is true. The `orelse` operator evaluates the second argument only if the first one is false.

The `not` operator is simpler, taking just one argument. It turns `true` into `false` and `false` into `true`. Unlike the other boolean operators, which go between their arguments, `not` goes before its single argument.

If you try to use these operators with any other atoms, you'll get a bad argument exception.



There are other atoms that often have an accepted meaning, like `ok` and `error`, but those are more conventions than a formal part of the language.

Guards

The `fall_velocity` calculations work fairly well, but there's still one glitch: if the function gets a negative value for distance, the square root (`sqrt`) function in the calculation will be unhappy:

```
5> drop:fall_velocity(earth,-20).
** exception error: bad argument in an arithmetic expression
   in function  math:sqrt/1
      called as math:sqrt(-392.0)
   in call from drop:fall_velocity/2 (drop.erl, line 4)
```

Since you can't dig a hole 20 meters down, release an object, and marvel as it accelerates to the surface, this isn't a terrible result. However, it might be more elegant to at least produce a different kind of error.

In Erlang, you can specify which data a given function will accept with *guards*. Guards, indicated by the `when` keyword, let you fine-tune the pattern matching based on the content of arguments, not just their shape. Guards have to stay simple, can use only a very few built-in functions, and are limited by a requirement that they evaluate only data without any side effects, but they can still transform your code.



You can find a list of functions that can safely be used in guards in [Appendix A](#).

Guards evaluate their expressions to `true` or `false`, as previously described, and the first one with a `true` result wins. That means that you can write `when true` for a guard that always gets called if it is reached, or block out some code you don't want to call (for now) with `when false`.

In this simple case, you can keep negative numbers away from the square root function by adding guards to the `fall_velocity` clauses, as shown in [Example 3-2](#), which you can find at `ch03/ex2-guards`.

Example 3-2. Adding guards to the function clauses

```
-module(drop).  
-export([fall_velocity/2]).  
  
fall_velocity(earth, Distance) when Distance >= 0 -> math:sqrt(2 * 9.8 * Distance);  
fall_velocity(moon, Distance) when Distance >= 0 -> math:sqrt(2 * 1.6 * Distance);  
fall_velocity(mars, Distance) when Distance >= 0 -> math:sqrt(2 * 3.71 * Distance).
```



In Erlang, greater-than-or-equal-to is written `>=`, and less-than-or-equal-to is written `=<`. Don't make them look like arrows.

The `when` expression describes a condition or set of conditions in the function head. In this case, the condition is simple: the `Distance` must be greater than or equal to zero. If you compile that code and ask for a negative velocity, the result is different:

```
5> drop:fall_velocity(earth,-20).  
** exception error: no function clause matching  
drop:fall_velocity(earth,-20) (drop.erl, line 12)
```

Because of the guard, Erlang doesn't find a function clause that works with a negative argument. The error message may not seem like a major improvement, but as you add layers of code, "not handled" may be a more appealing response than "broke my formula."

A clearer, though still simple, use of guards might be code that returns an absolute value. Yes, Erlang has a built-in function, `abs/1`, for this, but [Example 3-3](#) makes clear how this works.

Example 3-3. Calculating absolute value with guards

```
-module(mathdemo).  
-export([absolute_value/1]).  
  
absolute_value(Number) when Number < 0 -> -Number;  
absolute_value(Number) when Number == 0 -> 0;  
absolute_value(Number) when Number > 0 -> Number.
```

When `mathdemo:absolute_value` is called with a negative (less than zero) argument, Erlang calls the first clause, which returns the negation of that negative argument,

making it positive. When the argument equals (==) zero, Erlang calls the second clause, returning 0. Finally, when the argument is positive, Erlang calls the third clause, just returning the number. (The first two clauses have processed everything that isn't positive, so the guard on the last clause is unnecessary and will go away in [Example 3-4](#).)

```
1> c(mathdemo).
{ok,mathdemo}
2> mathdemo:absolute_value(-20).
20
3> mathdemo:absolute_value(0).
0
4> mathdemo:absolute_value(20).
20
```

This may seem like an unwieldy way to calculate. Don't worry—Erlang has simpler logic switches you can use inside of functions. However, guards are critically important to choosing among function clauses, which will be especially useful as you start to work with recursion in [Chapter 4](#).

Erlang runs through the function clauses in the order you list them, and stops at the first one that matches. If you find your information is heading to the wrong clause, you may want to re-order your clauses or fine-tune your guard conditions.

Also, when your guard clause is testing for just one value, you can easily switch to using pattern-matching instead of a guard. This `absolute_value` function in [Example 3-4](#) does the same thing as the one in [Example 3-3](#).

Example 3-4. Calculating absolute value with guards

```
absolute_value(Number) when Number < 0 -> -Number;
absolute_value(0) -> 0;
absolute_value(Number) -> Number.
```

In this case, it's up to you whether you prefer the simpler form or preserving a parallel approach.



You can also have multiple comparisons in a single guard. If you separate them with semicolons it works like an OR statement, succeeding if any of the comparisons succeeds. If you separate them with commas, it works like an AND statement, and they all have to succeed for the guard to succeed.

Underscoring That You Don't Care

Guards let you specify more precise handling of incoming arguments. Sometimes you may actually want handling that is less precise, though. Not every argument is essential to every operation, especially when you start passing around complex data structures.

You could create variables for arguments and then never use them, but you'll get warnings from the compiler (which suspects you must have made a mistake) and you may confuse other people using your code who are surprised to find your code cares about only half of the arguments they sent.

You might, for example, decide that you're not concerned with what planemo (for planetary mass object, including planets, dwarf planets, and moons) a user of your velocity function specifies and you're just going to use Earth's value for gravity. Then, you might write something like [Example 3-5](#), from *ch03/ex3-underscore*.

Example 3-5. Declaring a variable and then ignoring it

```
-module(drop).  
-export([fall_velocity/2]).  
  
fall_velocity(Planemo, Distance) -> math:sqrt(2 * 9.8 * Distance).
```

This will compile, but you'll get a warning, and if you try to use it for, say, Mars, you'll get the wrong answer for Mars.

```
1> c(drop).  
drop.erl:5: Warning: variable 'Planemo' is unused  
{ok,drop}  
2> drop:fall_velocity(mars, 20).  
19.79898987322333
```

On Mars, that should be more like 12 than 19, so the compiler was right to scold you.

Other times, though, you really only care about some of the arguments. In these cases, you can use a simple underscore (_). The underscore accomplishes two things: it tells the compiler not to bother you, and it tells anyone reading your code that you're not going to be using that argument. In fact, Erlang won't let you. You can try to assign values to the underscore, but Erlang won't give them back to you. It considers the underscore permanently unbound:

```
3> _ = 20.  
20  
4> _.  
* 1: variable '_' is unbound
```

If you really wanted your code to be earth-centric and ignore any suggestions of other planemos, you could instead write something like [Example 3-6](#).

Example 3-6. Deliberately ignoring an argument with an underscore

```
-module(drop2).  
-export([fall_velocity/2]).  
  
fall_velocity(_, Distance) -> math:sqrt(2 * 9.8 * Distance).
```

This time there will be no compiler warning, and anyone who looks at the code will know that first argument is useless.

```
5> c(drop2).  
{ok,drop2}  
6> drop2:fall_velocity(you_dont_care, 20).  
19.79898987322333
```

You can use underscore multiple times to ignore multiple arguments. It matches anything for the pattern match, and never binds, so there's never a conflict.



You can also start variables with underscores—like `_Planemo`—and the compiler won't warn if you never use those variables. However, those variables do get bound, and you can reference them later in your code if you change your mind. However, if you use the same variable name more than once in a set of arguments, even if the variable name starts with an underscore, you'll get an error from the compiler for trying to bind twice to the same name.

Adding Structure: Tuples

Erlang's tuples let you combine multiple items into a single composite data type. This makes it easier to pass messages between components, letting you create your own complex data types as you need. Tuples can contain any kind of Erlang data, including numbers, atoms, other tuples, and the lists and strings you'll encounter in later chapters.

Tuples themselves are simple, a group of items surrounded by curly braces:

```
1> {earth, 20}.  
{earth, 20}
```

Tuples might contain 1 item, or they might contain 100. Two to five seem typical (and useful, and readable). Often (but not always) an atom at the beginning of the tuple indicates what it's really for, providing an informal identifier of the complex information structure stored in the tuple.

Erlang includes rarely used built-in functions that give you access to the contents of a tuple on an item by item basis. You can retrieve the values of items with the `element` function, set values in a new tuple with the `setelement` function, and find out how many items are in a tuple with the `tuple_size` function.

```
2> Tuple = {earth, 20}.  
{earth,20}  
3> element(2, Tuple).  
20
```

```
4> NewTuple = setelement(2, Tuple, 40).  
{earth,40}  
5> tuple_size(NewTuple).  
2
```

If you can stick with pattern matching tuples, however, you'll likely create more readable code.

Pattern Matching with Tuples

Tuples make it easy to package multiple arguments into a single container, and let the receiving function decide what to do with them. Pattern matching on tuples looks much like pattern matching on atoms, except that there is, of course, a pair of curly braces around each set of arguments, as [Example 3-7](#), which you'll find in *ch03/ex4-tuples*, demonstrates.

Example 3-7. Encapsulating arguments in a tuple

```
-module(drop).  
-export([fall_velocity/1]).  
  
fall_velocity({earth, Distance}) -> math:sqrt(2 * 9.8 * Distance);  
  
fall_velocity({moon, Distance}) -> math:sqrt(2 * 1.6 * Distance);  
  
fall_velocity({mars, Distance}) -> math:sqrt(2 * 3.71 * Distance).
```

The arity changes: this version is `fall_velocity/1` instead of `fall_velocity/2` because the tuple counts as only one argument. The tuple version works much like the atom version but requires the extra curly braces when you call the function as well.

```
1> c(drop).  
{ok,drop}  
2> drop:fall_velocity({earth,20}).  
19.79898987322333  
3> drop:fall_velocity({moon,20}).  
8.0  
4> drop:fall_velocity({mars,20}).  
12.181953866272849
```

Why would you use this form when it requires a bit of extra typing? Using tuples opens more possibilities. Other code could package different things into tuples—more arguments, different atoms, even functions created with `fun()`. Passing a single tuple rather than a pile of arguments gives Erlang much of its flexibility, especially when you get to passing messages between different processes.

Processing Tuples

There are many ways to process tuples, not just the simple pattern matching shown in [Example 3-7](#). If you receive the tuple as a single variable, you can do many different things with it. A simple place to start is using the tuple as a pass through to a private version of the function. That part of [Example 3-8](#) may look familiar, as it's the same as the `fall_velocity/2` function in [Example 3-2](#). (You can find this at `ch03/ex5-tuplesMore.`)

Example 3-8. Encapsulating arguments in a tuple and passing them to a private function

```
-module(drop).
-export([fall_velocity/1]).

fall_velocity({Planemo, Distance}) -> fall_velocity(Planemo, Distance).

fall_velocity(earth, Distance) when Distance >= 0 -> math:sqrt(2 * 9.8 * Distance);
fall_velocity(moon, Distance) when Distance >= 0 -> math:sqrt(2 * 1.6 * Distance);
fall_velocity(mars, Distance) when Distance >= 0 -> math:sqrt(2 * 3.71 * Distance).
```

The `-export` directive makes *only* `fall_velocity/1`, the tuple version, public. The `fall_velocity/2` function is available within the module, however. It's not especially necessary here, but this “make one version public, keep another version with different arity private” is common in situations where you want to make a function accessible but don't necessarily want its inner workings directly available.

If you call this function—the tuple version, so curly braces are necessary—`fall_velocity/1` calls the private `fall_velocity/2`, which returns the proper value to `fall_velocity/1`, which will return it to you. The results should look familiar.

```
1> c(drop).
{ok,drop}
2> drop:fall_velocity({earth,20}).
19.798987322333
3> drop:fall_velocity({moon,20}).
8.0
4> drop:fall_velocity({mars,20}).
12.181953866272849
```

There are a few different ways to extract the data from the tuple. You could reference the components of the tuple by number using the built-in function `element`, which takes a numeric position and a tuple as its arguments. The first component of a tuple can be reached at position 1, the second at position 2, and so on.

```
fall_velocity(Where) -> fall_velocity(element(1,Where) , element(2,Where)).
```

You could also break things up a bit and do pattern matching after getting the variable:

```
fall_velocity(Where) ->
    {Planemo, Distance} = Where,
    fall_velocity(Planemo, Distance).
```

This function has more than one line. Note that actions are separated with commas, and that only the last line ends with a period. The result of that last line will be the value the function returns.

The pattern matching is a little different. The function accepted a tuple as its argument and assigned it to the variable `Where`. (If `Where` is not a tuple, the function will fail with an error.) Extracting the contents of that tuple, since we know its structure, can be done with a pattern match inside the function. The `Planemo` and `Distance` variables will be bound to the values contained in the `Where` tuple, and can then be used in the call to `fall_velocity/2`.

Logic and Recursion

So far, Erlang seems logical but fairly simple. Pattern matching controls the flow through a program, and requests that match a form return certain responses. While this is enough to get many things done, there are times when you'll want more powerful options, especially as you start working with larger and more complicated data structures.

Logic Inside of Functions

Pattern matching and guards are powerful tools, but there are times when it's much easier to do some comparisons inside of a function clause instead of creating new functions. Erlang's designers agreed, and created two constructs for evaluating conditions inside of functions: the `case` expression and the less frequently used `if` expression.

The `case` construct lets you use pattern matching and guards inside of a function clause. It reads most clearly when a single value (or set of values) needs to be compared with multiple possibilities. The `if` construct evaluates only a series of guards, without pattern matching. The `if` construct tends to produce more readable code in situations where the multiple possibilities are specified by combinations of different values.

Both constructs return a value your code can capture.

Evaluating Cases

The `case` construct lets you perform pattern-matching inside of your function clause. If you found the multiple function clauses of [Example 3-2](#) hard to read, you might prefer to create a version that looks like [Example 4-1](#), which you can find in *ch04/ex1-case*.

Example 4-1. Moving pattern matching inside the function

```
-module(drop).  
-export([fall_velocity/2]).
```

```

fall_velocity(Planemo, Distance) when Distance >= 0  ->
case Planemo of
    earth -> math:sqrt(2 * 9.8 * Distance);
    moon -> math:sqrt(2 * 1.6 * Distance);
    mars -> math:sqrt(2 * 3.71 * Distance) % no closing period!
end.

```

The `case` construct will compare the atom in `Planemo` to the values listed, going down the list in order. It won't process beyond the match it finds. The `case` construct will return the result of different calculations based on which atom is used, and because the `case` construct returns the last value in the function clause, the function will return that value as well.



You can use the underscore (`_`) for your pattern match if you want a choice that matches "everything else." However, you should always put that last—nothing that comes after that will ever be evaluated.

The results should look familiar:

```

1> c(drop).
{ok,drop}
2> drop:fall_velocity(earth,20).
19.79898987322333
3> drop:fall_velocity(moon,20).
8.0
4> drop:fall_velocity(mars,20).
12.181953866272849
5> drop:fall_velocity(mars,-20).
** exception error: no function clause matching
   drop:fall_velocity(mars,-20) (drop.erl, line 5)

```

The `case` construct switches among planemos, while the guard clause on the function definition keeps out negative distances, producing (rightly) the error on line 5. This way the guard needs to appear only once.

You can also use the return value from the `case` construct to reduce duplicate code and make the logic of your program clearer. In this case, the only difference between the calculations for `earth`, `moon`, and `mars` is a gravitational constant. [Example 4-2](#), which you can find in `ch04/ex2-case`, shows how to make the `case` construct return the gravitational constant for use in a single calculation at the end.

Example 4-2. Using the return value of the case construct to clean up the function

```

-module(drop).
-export([fall_velocity/2]).

```

```

fall_velocity(Planemo, Distance) when Distance >= 0 ->
    Gravity = case Planemo of
        earth -> 9.8;
        moon -> 1.6;
        mars -> 3.71
    end, % note comma - function isn't done yet
    math:sqrt(2 * Gravity * Distance).

```

This time, the `Gravity` variable is set to the return value of the `case` construct. Note the comma after the `end`. This function isn't done yet! Commas let you separate constructs inside of function declarations. The now more readable formula `math:sqrt(2 * Gravity * Distance)`. is the last line of the function, and the value it produces will be the return value.

You can also use guards with a `case` statement, as shown, perhaps less than elegantly, in [Example 4-3](#), which is in `ch04/ex3-case`. This might make more sense if there were different planemos with different rules about distances.

Example 4-3. Moving guards into the case statement

```

-module(drop).
-export([fall_velocity/2]).

fall_velocity(Planemo, Distance) ->
    Gravity = case Planemo of
        earth when Distance >= 0 -> 9.8;
        moon when Distance >= 0 -> 1.6;
        mars when Distance >= 0 -> 3.71
    end, % note comma - function isn't done yet
    math:sqrt(2 * Gravity * Distance).

```

This produces similar results, except that the error message at the end changes from `no function clause matching drop:fall_velocity(mars,-20)` to `no case clause matching mars in function drop:fall_velocity/2:`

```

1> c(drop).
{ok,drop}
2> drop:fall_velocity(mars,20).
12.181953866272849
3> drop:fall_velocity(mars, -20).
** exception error: no case clause matching mars
   in function drop:fall_velocity/2 (drop.erl, line 6)

```

The error is correct, in that the `case` construct is trying to match `mars`, but misleading because the problem isn't with `mars` but rather with the guard that's checking the `Distance` variable. If Erlang tells you that your `case` doesn't match but a match is obviously right there in front of you, check your guard statements.

If This, Then That

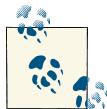
The `if` construct is broadly similar to the `case` statement, but without the pattern matching. This allows you to write a catch-all clause—a guard matching `true` at the end if you would like, and often makes it easier to express logic based on broader comparisons than simple matching.

Suppose, for example, that the precision of the `fall_velocity` function is too much. Instead of an actual speed you'd like to describe the speed produced by dropping from a tower of a given height. You can add an `if` construct that does that to the earlier code from [Example 4-2](#), as shown in [Example 4-4](#), in `ch04/ex4-if`.

Example 4-4. Adding an if construct to convert numbers into atoms

```
-module(drop).
-export([fall_velocity/2]).  
  
fall_velocity(Planemo, Distance) when Distance >= 0  ->
    Gravity = case Planemo of
        earth -> 9.8;
        moon -> 1.6;
        mars -> 3.71
    end,
    Velocity = math:sqrt(2 * Gravity * Distance),  
  
    if
        Velocity == 0 -> 'stable';
        Velocity < 5 -> 'slow';
        Velocity >= 5, Velocity < 10 -> 'moving';
        Velocity >= 10, Velocity < 20 -> 'fast';
        Velocity >= 20 -> 'speedy'
    end.
```

This time, the `if` construct returns a value (an atom describing the velocity) based on the many guards it includes. Because that value is the last thing returned within the function, that becomes the return value of the function.



The commas in the `if` behave like the `and` operator.

The results are a little different from past trials:

```
1> c(drop).
{ok,drop}
2> drop:fall_velocity(earth,20).
```

```

fast
3> drop:fall_velocity(moon,20).
moving
4> drop:fall_velocity(mars,20).
fast
5> drop:fall_velocity(earth,30).
speedy

```

If you want to capture the value produced by the `if` construct and use it for something else, you can. [Example 4-5](#), in `ch04/ex5-if`, sends a warning to standard output (in this case the Erlang shell) if you drop an object too fast.

Example 4-5. Sending an extra warning if the velocity is too high

```

-module(drop).
-export([fall_velocity/2]).

fall_velocity(Planemo, Distance) when Distance >= 0  ->
    Gravity = case Planemo of
        earth -> 9.8;
        moon -> 1.6;
        mars -> 3.71
    end,
    Velocity = math:sqrt(2 * Gravity * Distance),
    Description = if
        Velocity == 0 -> 'stable';
        Velocity < 5 -> 'slow';
        (Velocity >= 5) and (Velocity < 10) -> 'moving';
        (Velocity >= 10) and (Velocity < 20) -> 'fast';
        Velocity >= 20 -> 'speedy'
    end,
    if
        (Velocity > 40) -> io:format("Look out below!~n");
        true -> true
    end,
    Description.

```

The new (second) `if` clause checks the `Velocity` variable to see if it's above 40. If it is, it calls `io:format`, which creates a side effect: a message on the screen. However, every `if` must find some true statement or it will report an error in those cases when nothing matches. Here, you could add an explicit case matching when the `Velocity` is less than or equal to 40. In many cases, however, it won't matter. The `true -> true` line is a catch-all that returns true no matter what reaches it. After the `if` concludes, the single line `Description.` returns the contents of the `Description` variable from the function.



The catchall approach works in cases where you only want to test for a subset of cases among a complicated set of possibilities. In cases as simple as this example, however, it's probably cleaner to create a more explicit test.

The function produces an extra result—the message—when the distance is large enough (and the planemo's gravity strong enough) to produce a velocity faster than 40 meters per second:

```
1> c(drop).
{ok,drop}
2> drop:fall_velocity(earth,10).
fast
3> drop:fall_velocity(earth,200).
Look out below!
speedy
```

Variable Assignment in case and if Constructs

Every possible path created in a `case` or `if` statement has the opportunity to bind values to variables. This is usually a wonderful thing, but could let you create unstable programs by assigning different variables in different clauses. This might look something like [Example 4-6](#), which you can find in `ch04/ex6-broken`.

Example 4-6. A badly broken if construct

```
-module(broken).
-export([bad_if/1]).

bad_if(Test_val) ->

    if
        Test_val < 0 -> X = 1;
        Test_val >= 0 -> Y = 2
    end,

    X+Y.
```

In theory, after the `case` or `if` is over, the program might crash because of unbound variables. However, Erlang won't let you get that far:

```
1> c(broken).
broken.erl:11: variable 'X' unsafe in 'if' (line 6)
broken.erl:11: variable 'Y' unsafe in 'if' (line 6)
error
```

The compilation errors turn up where your program actually uses the variables. The Erlang compiler double-checks to make sure that the variables it's about to put to use are properly defined. It won't let you compile something this broken.

You *can* bind variables in an `if` or `case` construct. You have to define all of the variables in every single clause, however. If you're defining only one variable, it's also much cleaner to bind the return value of the `if` or `case` clause to a variable instead of defining that variable in every clause.

The Gentlest Side Effect: `io:format`

Up until [Example 4-5](#), all of the Erlang examples you've seen focused on a single path through a group of functions. You put an argument or arguments in, and got a return value back. That approach is the cleanest way to do things: you can count on things that worked before to work again because there's no opportunity to muck up the system with leftovers of past processing.

[Example 4-5](#) stepped outside of that model, creating a side effect that will linger after the function is complete. The side effect is just a message that appears in the shell (or in standard output when you start running Erlang outside of the shell). Applications that share information with multiple users or keep information around for longer than a brief processing cycle will need stronger side effects, like storing information in databases.

Erlang best practice suggests using side effects **only** when you really need to. An application that presents an interface to a database, for example, really will need to read and write that database. An application that interacts with users will need to put information on the screen (or other interface) so that users can figure out what they're expected to do.

Side effects are also extremely useful for tracing logic when you are first starting out. The simplest way to see what a program is doing, before you've learned how to use Erlang's built-in tracing and debugging tools for processes, is to have the program report its status at points you consider interesting. This is not a feature you want to leave in shipping code, but when you're getting started, it can give you an easily understandable window into your code's behavior.

The `io:format` function lets you send information to the console, or, when you're eventually running code outside of the console, to other places. For now, you'll just use it to send messages from the program to the console. [Example 4-5](#) showed the simplest way to use `io:format`, just printing a message it takes in double quotes:

```
io:format("Look out below!~n") ;
```

The `~n` represents a *newline*, telling the console to start any new messages it sends at the beginning of the next line. It makes your results look a bit neater.

The more typical way to use `io:format` includes two arguments: a double-quoted formatting string, and a list of values that can be included in the string. In this case (which you can see in `ch04/ex7/format`), it might look like the following:

```
io:format("Look out below! ~w is too high.~n", [Distance]) ;
```

or:

```
io:format("Look out below! ~w is too high on ~w.~n", [Distance, Planemo]) ;
```

`io:format/2` offers many formatting options beyond `~w` and `~n`. You'll encounter them as they become necessary, but if you're impatient, there's a list in [Appendix A](#). You may also want to explore the section on error logging in [Chapter 9](#), if you find yourself using `io:format` for tasks that might be helped by more sophisticated logging tools.



Erlang flatly prohibits operations that could cause side effects in guard expressions. If side effects were allowed in guards, then any time a guard expression was evaluated—whether it returned true or false—the side effect would happen. `io:format` wouldn't likely do anything terrible, but these rules mean that it too is blocked from use in guard expressions.

Simple Recursion

Because variables can't change values, the main tool you'll use to repeat actions is recursion: having a function call itself until it's (hopefully) reached a conclusion. This can sound complicated, but it doesn't have to be.

There are two basic kinds of useful recursion. In some situations, you can count on the recursion to reach a natural end. The process runs out of items to work on, or reaches a natural limit. In other situations, there is no natural end, and you need to keep track of the result so the process will end. If you can master these two basic forms, you'll be able to create many more complex variations.



There is a third form, in which the recursive calls never reach an end. This is called an *infinite loop*, and is an error you'll want to avoid.

Counting Down

The simplest model of recursion with a natural limit is a countdown, like the one used for rockets. You start with a large number, and count down to zero. When you reach zero, you're done (and the rocket takes off, if there is one).

To implement this in Erlang, you'll pass a starting number to an Erlang function. If the number is greater than zero, it will then announce the number and call itself with the number minus one as the argument. If the number is zero (or less), it will announce `blastoff!` and end. [Example 4-7](#), found in `ch04/ex8-countdown`, shows one way to do this.

Example 4-7. Counting down

```
-module(count).
-export([countdown/1]).  
  
countdown(From) when From > 0 ->
    io:format("~w!~n", [From]),
    countdown(From-1);  
  
countdown(From) ->
    io:format("blastoff!~n").
```

The last clause could have a guard—when `From =< 0`—but it would be useful only to make clear when the blastoff happens to human readers. Unnecessary guard clauses may lead to weird errors, so brevity is probably the best option here, though you’ll get a warning that `From` is unused in the final clause. Here’s a test run:

```
1> c(count).
count.erl:9: Warning: variable 'From' is unused
{ok,count}
2> count:countdown(2).
2!
1!
blastoff!
ok
```

The first time through, Erlang chose the first clause of `countdown(From)`, passing it a value of 2. That clause printed 2, plus an exclamation point and a newline, and then it called the `countdown` function again, passing it a value of 1. That triggered the first clause again. It printed 1, plus an exclamation point and a newline, and then it called the `countdown` function again—this time passing it a value of 0.

The value of 0 triggered the second clause, which printed `blastoff!` and ended. After running three values through the same set of code, the function comes to a neat conclusion.



You could also implement this conclusion with an `if` statement inside a single `countdown(From)` function clause. This is unusual in Erlang. I find guards more readable in these cases, but you may see things differently.

Counting Up

Counting up is trickier because there’s no natural endpoint, so you can’t model your code on [Example 4-7](#). Erlang’s single-assignment approach to variables rules out some approaches, but there’s another way to make this work, using an *accumulator*. An

accumulator is an extra argument that keeps track of the current result of past work, passing it back into a recursive function. (You can have more than one accumulator argument if you need, though one is often sufficient.) [Example 4-8](#), which you can find in *ch04/ex9-countup*, shows how to add a `countup` function to the `count` module, which lets Erlang count up to a number.

Example 4-8. Counting up

```
-module(count).
-export([countdown/1, countup/1]).  
  
countup(Limit) ->
    countup(1, Limit).  
  
countup(Count, Limit) when Count <= Limit ->
    io:format("~w!~n", [Count]),
    countup(Count+1, Limit);  
  
countup(Count, Limit) ->
    io:format("Finished.~n").  
  
...
```

It produces results such as the following:

```
1> c(count).
{ok,count}
2> count:countup(2).
1!
2!
Finished.
ok
```

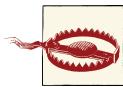
The `export` directive makes the `countup/1` function visible (as well as the earlier `countdown/1`, which you'll find in the sample code).

The `countup/2` function, which does most of the work, remains private, not exported. This isn't mandatory; you might make it public if you wanted to support counting between arbitrary values, but it's common Erlang practice. Keeping the recursive internal functions private makes it less likely that someone will misuse them for purposes they're not well-suited to. In this case, it doesn't matter at all, but it can make a big difference in other more complex situations, especially when data is modified.

When you call `countup/1`, it calls `countup/2` with an argument of 1 (for the current count) and the `Limit` value you provided for the upper limit.

If the current count is less than or equal to the upper limit, the first clause of the `countup/2` function reports the current `Count` value with `io:format`. Then it calls itself again, increasing the `Count` by one but leaving the `Limit` alone.

If the current count is greater than the upper limit, it fails the guard on the first clause, so the second clause kicks in, reports “Finished.”, and is done.



The guards here are sufficient to avoid infinite loops. You can enter zero, negative numbers, or decimals as arguments to `countup/1` and it will terminate neatly. You can get into serious trouble, however, if your termination test relies on `==` or `=:=` for a more exact comparison rather than `>=` or `<=` for a rough comparison.

Recurse with Return Values

The counting examples are simple—they demonstrate how recursion works, but just discard the return values. There are return values—the `io:format` calls return the atom `ok`—but they aren’t of much use. More typically, a recursive function call will make use of the return value.

A classic recursive call calculates factorials. A factorial is the product of all positive integers equal to or less than the argument. The factorial of 1 is 1; 1 by itself yields 1. The factorial of 2 is 2; 2×1 yields 2. It starts to get interesting at 3, where $3 \times 2 \times 1$ is six. At 4, $4 \times 3 \times 2 \times 1$ is 24, and the results get larger rapidly with larger arguments.

There was a pattern to that, though. You can calculate any factorial by multiplying the integer by the factorial of one less. That makes it a perfect case for using recursion, using the results of smaller integers to calculate the larger ones. This approach is similar to the countdown logic, but instead of just counting, the program collects calculated results. That could look like [Example 4-9](#), which you’ll find in `ch04/ex10-factorial-down`.

Example 4-9. A factorial written with the counting down approach

```
-module(fact).
-export([factorial/1]).

factorial(N) when N > 1 ->
    N * factorial(N-1);

factorial(N) when N <= 1 ->
    1.
```

The first clause of `factorial` uses the pattern previously described. The first clause, used for numbers above one, returns a value that is the number `N` times the factorial of the next integer down. The second clause returns the value 1 when it reaches 1. Using `<=` in that comparison, rather than `==`, gives the function more resilience against non-integer or negative arguments, though the answers it returns aren’t quite right: factorials really only work for integers of 1 or higher. The results are as previously suggested:

```
1> c(fact).
{ok,fact}
```

```

2> fact:factorial(1).
1
3> fact:factorial(3).
6
4> fact:factorial(4).
24
5> fact:factorial(40).
815915283247897734345611269596115894272000000000

```

This works, but it may not be clear why it works. Yes, the function counts down and collects the values, but if you want to see the mechanism, you need to add some `io:format` calls into the code, as shown in [Example 4-10](#). (You can find this at *ch04/ex10-factorial-down-instrumented*.)

Example 4-10. Looking into the factorial recursion calls

```

-module(fact).
-export([factorial/1]).

factorial(N) when N > 1 ->
    io:format("Calling from ~w.~n", [N]),
    Result = N * factorial(N-1),
    io:format("~w yields ~w.~n", [N, Result]),
    Result;

factorial(N) when N == 1 ->
    io:format("Calling from 1.~n"),
    io:format("1 yields 1.~n"),
    1.

```

There's a bit more overhead here. To present the result of the recursive call and still return that value to the next recursive call requires storing it in a variable, here called `Result`. The `io:format` call makes visible which value produced the result. Then, because the last value expression in a function clause is the return value, `Result` appears again. The second clause for 1 is similar, except that it can report simply that `1 yields 1.` because it always will.

When you compile this and run it, you'll see something such as the following:

```

7> fact:factorial(4).
Calling from 4.
Calling from 3.
Calling from 2.
Calling from 1.
1 yields 1.
2 yields 2.
3 yields 6.
4 yields 24.
24

```

Although the calls count down the values, as the logic would suggest, the messages about results don't appear until the countdown is complete, and then they all appear in order, counting up.

The reason this happens is that the function calls don't return values until the countdown is complete. Until then, the Erlang runtime builds a stack of frames corresponding to the function calls. You can think of the frames as paused versions of the function logic, waiting for an answer to come back. Once the call with an argument of 1 returns a simple value, not calling any further, Erlang can unwind those frames and calculate the Result. That unwinding presents the results—"X yields Y"—in the order that the frames unwind.

That "unwinding" also means that the code in [Example 4-9](#) and [Example 4-10](#) is not *tail recursive*. When Erlang encounters code that ends with a simple recursive call, it can optimize the handling to avoid keeping that stack of calls around. This probably doesn't matter for a one-time calculation, but it makes a huge difference when you write code that will stay running for a long time.

You can achieve tail recursion for factorials by applying the counting up approach to factorials. You'll get the same results (at least for integer values), but the calculations will work a little differently, as shown in [Example 4-11](#), at *ch04/ex12-factorial-up*.

Example 4-11. A factorial written with the counting up approach

```
-module(fact).
-export([factorial/1]).  
  
factorial(N) ->
    factorial(1, N, 1).  
  
factorial(Current, N, Result) when Current ==< N ->
    NewResult = Result*Current,
    io:format("~w yields ~w!~n", [Current, NewResult]),
    factorial(Current+1, N, NewResult);  
  
factorial(Current, N, Result) ->
    io:format("Finished.~n"),
    Result.
```

As in the counting up example, the main function call, here `factorial/1`, calls a private function, `factorial/3`. In this case, there are two accumulators. `Current` stores the current position in the count, whereas `Result` is the answer from the previous multiplication. When the value of `Current` climbs past the limiting value `N`, the first guard fails, the second clause is invoked, and the function is finished and returns the `Result`. (You'll get a compilation warning because the final clause doesn't use the accumulator variables `Current` or `N`. You can ignore it.)

Because `factorial/3`'s last call in the recursive section is to itself, without any complications to track, it is tail recursive. Erlang can minimize the amount of information it has to keep around while the calls all happen.

The calculation produces the same results, but does the math in a different order:

```
9> fact:factorial(4).  
1 yields 1!  
2 yields 2!  
3 yields 6!  
4 yields 24!  
Finished.  
24
```

Although the code is tracking more values, the Erlang runtime has less to do. When it finally hits the final result, there's no further calculation needed. That result is the result, and it passes back through to the original call. This also makes it easier to structure the `io:format` calls. If you remove them or comment them out, the rest of the code stays the same.

Communicating with Humans

Erlang's origins in telecom switching have left it with a fairly minimal set of tools for communicating with people, but there's enough there to do worthwhile things. You've already used some of it (`io:format/1-2`) but there are more pieces you'll want to learn to handle communications with people and sometimes with other applications. At the very least, this chapter will let you build more convenient interfaces for testing your code than calling functions from the Erlang shell.



If you're feeling completely excited about the recursion you learned in [Chapter 4](#), you may want to jump ahead to [Chapter 6](#), where that recursion will once again be front and center.

Strings

Atoms are great for sending messages within a program, even messages that the programmer can remember, but they're not really designed for communicating outside of the context of Erlang processes. If you need to be assembling sentences or even presenting information, you'll want something more flexible. Strings, sequences of characters, are the structure you need. You've already used strings a little bit, as the double-quoted arguments to `io:format` in [Chapter 4](#):

```
io:format("Look out below!~n") ;
```

The double-quoted content (`Look out below!~n`) is a string. A string is a sequence of characters. If you want to include a double-quote within the string, you can escape it with a backslash, like `\"`. To include a backslash, you have to use `\\`, and [Appendix A](#) includes a complete list of escapes and other options. If you create a string in the shell, Erlang will report back the string *with* the escapes. To see what it "really" contains, use `io:format`:

```
1> X = "Quote - \" in a string. \n Backslash, too: \\ . \n".
"Quote - \" in a string. \n Backslash, too: \\ ."
2> io:format(X).
Quote - " in a string.
Backslash, too: \ .
ok
```



If you start entering a string and don't close the quotes, when you press Enter, the Erlang shell will just give you a new line with the same number. This lets you include newlines in strings, but it can be very confusing. If you think you're stuck, usually entering ". will get you out of it.

Erlang development hasn't focused heavily on text historically, but if your programs involve sharing information with humans, you'll want to get familiar with how to get information into and out of strings. This is an area where you may want to spend a fair amount of time in the shell playing with different tools.



Technically, strings don't really exist as a type in Erlang because strings are lists of characters. Thinking about strings as lists of characters, however, is useful in only a few situations, typically where you want to process a string from start to end. You'll learn about lists in [Chapter 6](#), and the code here will have to use a list built-in function, but for now you should just think about string operations rather than lists.

The simplest, usually, is concatenation, where you combine two strings into one. Erlang offers two easy ways to do this. The first uses the ++ operator:

```
1> "erl" ++ "ang".
"erlang"
2> A="ang".
"ang"
3> "erl" ++ A.
"erlang"
```

The other approach uses an explicit `string:concat/2` function:

```
4> string:concat("erl", "ang").
"erlang"
5> N="ang".
"ang"
6> string:concat("erl", N).
"erlang"
```

The ++ operator is usually more convenient because it lets you work with more than two arguments without nesting functions.



Erlang has a shortcut where you can concatenate two strings just by putting them next to each other: "erl" "ang" will end up as "erlang". However, if you try to mix variables into that, you'll get a syntax error. This shortcut is of limited value except maybe when you're cutting and pasting quoted values as you're writing your code, and doesn't work in every context.

Erlang also offers three options for comparing string equality, the `==` operator, the `=:=` (exact equality) operator, and a `string:equal/2` function. The `==` operator is generally the simplest for this, though the others produce the same results:

```
7> "erl" == "erl".
true
8> "erl" == "ang".
false
9> G ="ang".
"ang"
10> G == "ang".
true
```

Erlang doesn't offer functions for changing strings in place, as that would work badly with a model where variable contents don't change. However, it does offer a set of functions for finding content in strings and dividing or padding those strings, which together let you extract information from a string (or multiple strings) and recombine it into a new string.

If you want to do more with your strings, you should definitely explore the documentation for the `string` and `re` (regular expressions) Erlang modules. If the strings you want to work with represent file or directory names, definitely explore the `filename` module. If you need to perform Unicode encoding conversion on Erlang strings, you'll also want to explore the `unicode` module. (By default, Erlang represents characters using UTF-8 values.)



I'm working on creating a single wrapper module that assembles Erlang's tools for working with strings into one place. For more, visit <https://github.com/simonstl/erlang-simple-string>.

Asking Users for Information

Many Erlang applications run kind of like wholesalers—in the background, providing goods and services to retailers who interact directly with users. Sometimes, however, it's nice to have a direct interface to code that is a little more customized than the Erlang

console. You probably won't write many Erlang applications whose primary interface is the command line, but you may find that interface very useful when you first try out your code. (Odds are good that if you're working with Erlang, you don't mind using a command-line interface, either.)

You *can* mix input and output with your program logic, but for this kind of simple facade, it probably makes better sense to put it in a separate module. In this case, the `ask` module will work with the `drop` module from [Example 3-8](#).



Erlang's `io` functions for input have a variety of strange interactions with the Erlang shell, as discussed in the following section. You will have better luck working with them in other contexts.

Gathering Terms

The simplest way to build an interface—an interface probably just for programmers—is to create a way for users to enter Erlang terms using `io:read/1`. This lets users enter a complete Erlang term—an atom, number, or tuple, for example. An initial version of this might look like [Example 5-1](#), which you can find in `ch05/ex1-ask`.

Example 5-1. Asking the user for an Erlang term

```
-module(ask).  
-export([term/0]).  
  
term() ->  
    Input = io:read("What {planemo, distance} ? >>"),  
    Term = element(2,Input),  
    drop:fall_velocity(Term).
```

The `Input` variable will be set by the call to `io:read/1`, getting an Erlang term. If all goes well, it will contain a tuple like `{ok,{mars,20}}`, where the first value is `ok` and the second value of the tuple is the term the user entered. Extracting that value—in this case a tuple—requires a call to the `element/2` method. Finally, the code calls the `drop:fall_velocity` method with that value.



If you wanted, you could cram that all into one line as `term() -> drop:fall_velocity(element(2,io:read("What {planemo, distance} ? >>")))`, but that's both hard to read and hard to modify.

For your own use, this could be perfectly fine. A simple session might look like the following:

```

1> c(drop).
{ok,drop}
2> c(ask).
{ok,ask}
3> ask:term().
What {planemo, distance} ? >>{mars,20}.
12.181953866272849

```

If you leave off the period at the end of the term, Erlang will repeat the prompt but not show where you were, trusting you to read the line above. Also, the things you enter at an `io:read/1` prompt become part of the console's command history, and you can repeat them with the up arrow. (These issues are interactions with the Erlang shell, not issues with the function itself.)

Things can get weird quickly, however, if the user enters unexpected terms—a number instead of a tuple—or broken terms, with bad syntax.

```

4> ask:term().
What {planemo, distance} ? >>20.
** exception error: no function clause matching
     drop:fall_velocity(20) (drop.erl, line 4)
5> ask:term().
What {planemo, distance} ? >>.
** exception error: no function clause matching
     drop:fall_velocity({1,erl_parse,
     ["syntax error before: "','".'"]}) (drop.erl, line 4)

```

In both cases, passing the extracted `Term` directly to `fall_velocity/1` is a bad idea. In the first case, it's because `fall_velocity/1` expects a tuple, not a bare number. In the second case, `fall_velocity/1` has a similar problem, but it's being sent an error message, not a term it can process. [Example 5-2](#), in `ch05/ex2-ask`, shows a better way to handle these kinds of problems. It gives the user a direct error message when it encounters the wrong type of information or broken information. (It also uses pattern matching instead of `element/2`.)

Example 5-2. Asking the user for an Erlang term and handling bad results

```

-module(ask).
-export([term/0]).

term() ->
    Input = io:read("What {planemo, distance} ? >>"),
    process_term(Input).

process_term({ok, Term}) when is_tuple(Term) -> drop:fall_velocity(Term);

process_term({ok, _}) -> io:format("You must enter a tuple.~n");

process_term({error, _}) -> io:format("You must enter a tuple with correct syntax.~n").

```

This doesn't solve every possible problem. Users could still enter tuples with the wrong content, and `drop:fall_velocity` will report an error. [Chapter 9](#) will explore how to address that problem in much greater detail.

When you go to the trouble of building this kind of interface, however, it's probably not because typing `ask:term()` is shorter than typing `drop:fall_velocity`. Odds are good that you want to try a number of values and possibilities, so you want the question repeated. [Example 5-3](#), in `ch05/ex3-ask`, presents the result of a (correctly formatted) call to the user and then calls `term()` again, setting up a recursive loop. (It also offers a nice way to exit the loop.)

Example 5-3. Asking the user for an Erlang term and handling bad results

```
-module(ask).
-export([term/0]).

term() ->
    Input = io:read("What {planet, distance} ? >>"),
    process_term(Input).

process_term({ok, Term}) when is_tuple(Term) ->
    Velocity = drop:fall_velocity(Term),
    io:format("Yields ~w.~n",[Velocity]),
    term();

process_term({ok, quit}) ->
    io:format("Goodbye.~n");
    % does not call term() again

process_term({ok, _}) ->
    io:format("You must enter a tuple.~n"),
    term();

process_term({error, _}) ->
    io:format("You must enter a tuple with correct syntax.~n"),
    term().
```

When you compile the `ask` module and call `ask:term/0`, you'll see the question repeated as long as you keep entering appropriate tuples. To break out of that loop, just enter the atom `quit` followed by a period.

```
6> c(ask).
{ok,ask}
7> ask:term().
What {planet, distance} ? >>{mars,20}.
Yields 12.181953866272849.
What {planet, distance} ? >>20.
You must enter a tuple.
What {planet, distance} ? >>quit.
Goodbye.
ok
```

Gathering Characters

The `io:get_chars/2` function will let you get just a few characters from the user. This seems like it should be convenient if, for example, you have a list of options. Present the options to the user, and wait for a response. In this case, the list of planemos is the option, and they're easy to number 1 through 3, as shown in the code for [Example 5-4](#), which you can find at `ch05/ex4-ask`. That means you just need a single character response.

Example 5-4. Presenting a menu and waiting for a single character response

```
-module(ask).  
-export([chars/0]).  
  
chars() ->  
    io:format("Which planemo are you on?~n"),  
    io:format(" 1. Earth ~n"),  
    io:format(" 2. Earth's Moon~n"),  
    io:format(" 3. Mars~n"),  
    io:get_chars("Which? > ",1).
```

Most of that is presenting the menu, and you could combine all of those `io:format/1` calls into a single call if you wanted. The key piece is the `io:get_chars/2` call at the end. The first argument is a prompt, and the second is the number of characters you want returned. The function still lets users enter whatever they want until they press Enter, but it will tell you only the first however many characters you specified.

```
1> c(ask).  
{ok,ask}  
2> ask:chars().  
Which planemo are you on?  
 1. Earth  
 2. Earth's Moon  
 3. Mars  
Which? > 3  
"3"  
3>  
3>
```

The `io:get_chars` function returns the string "3", the character the user entered, after they hit Enter. However, as you can tell by the duplicated command prompt, the Enter still gets reported to the Erlang shell. This can get stranger if users enter more content than is needed:

```
4> ask:chars().  
Which planemo are you on?  
 1. Earth  
 2. Earth's Moon  
 3. Mars
```

```
Which? > 222222
"2"
5> 22222
5>
```

There may be times when `io:get_chars` is exactly what you want, but odds are good, at least when working within the shell, that you'll get cleaner results by taking in a complete line of user input and picking what you want from it.

Reading Lines of Text

Erlang offers a few different functions that pause to request information from users. The `io:get_line/1` function waits for the user to enter a complete line of text terminated by a newline. You can then process the line to extract the information you want, and nothing will be left in the buffer. [Example 5-5](#), in `ch05/ex5-ask`, shows how this could work, though extracting the information is somewhat more complicated than I would like.

Example 5-5. Collecting user responses a line at a time

```
-module(ask).
-export([line/0]).

line() ->
    Planemo = get_planemo(),
    Distance = get_distance(),
    drop:fall_velocity({Planemo, Distance}).

get_planemo() ->
    io:format("Where are you?~n"),
    io:format(" 1. Earth ~n"),
    io:format(" 2. Earth's Moon~n"),
    io:format(" 3. Mars~n"),
    Answer = io:get_line("Which? > "),

    Value = hd(Answer),
    char_to_planemo(Value).

char_to_planemo(Char) ->
    case
        [Char] == "1" -> earth;
        Char == $2 -> moon;
        Char == $1 -> mars
    end.

get_distance() ->
    Input = io:get_line("How far? (meters) > "),
    Value = string:strip(Input, right, $\n),
    {Distance, _} = string:to_integer(Value),
    Distance.
```

To clarify the code, the `line/0` function just calls three other functions. It calls `get_planemo/0` to present a menu to the user and get a reply, and it similarly calls `get_distance/0` to ask the user the distance of the fall. Then it calls `drop:fall_velocity/1` to return the velocity at which a frictionless object will hit the ground when dropped from that height at that location.

The `get_planemo/0` function is a combination of `io:format/1` calls to present information and an `io:get_line/1` call to retrieve information from the user. Unlike `io:get_chars/1`, `io:get_line/1` returns the entire value the user entered, including the newline, and leaves nothing in the buffer.

```
get_planemo() ->
    io:format("Where are you?~n"),
    io:format(" 1. Earth ~n"),
    io:format(" 2. Earth's Moon~n"),
    io:format(" 3. Mars~n"),
    Answer = io:get_line("Which? > "),

    Character = hd(Answer),
    char_to_planemo(Character).
```

The last two lines are the actual string processing. The only piece of the response that matters to this application is the first character of the string. The easy way to grab that is with the built-in function `hd/1`, which pulls the first item from a string or list.



Because strings are really lists of numbers, you could instead call `lists:nth(1, Answer)`. The first argument, 1, is the position you want to retrieve, and the second argument, `Answer`, is the list, in this case a string, from which you want to retrieve it. For this function, the first character in an Erlang string is in position 1, not 0 as in many other languages. That makes the function name `nth` make sense when it's time to retrieve the 4th, 5th, 6th, and so on values.

The `drop:fall_velocity/1` function won't know what to do with a `planemo` listed as 1, 2, or 3; it expects an atom of `earth`, `moon`, or `mars`. The `get_planemo/0` function concludes by returning the value of that conversion, performed by the `char_to_planemo/1` function:

```
char_to_planemo(Char) ->
    case
        [Char] == "1" -> earth;
        Char == $2 -> moon;
        Char == $3 -> mars
    end.
```

The `case` statement shows three different ways of testing the character. If you prefer to evaluate the character as text, you can put square brackets around it and compare it to a string, like "1" here. You can also test against Erlang's character notation, in which `$2` is the value for the character two. Finally, if you're comfortable with character values, you can compare it to those values, like 51, which corresponds to 3. The atom returned by the `case` statement will be returned to the `get_planemo/0` function, which will in turn return it to the `line/0` function for use in the calculation.

You could also rewrite that function to skip the `case` statement and just use pattern matching:

```
char_to_planemo($1) -> earth;
char_to_planemo($2) -> moon;
char_to_planemo($3) -> mars.
```



Erlang's character notation understands Unicode as well. If you try `$☃`, the Unicode Snowman, Erlang will understand that it is character 9731, hex 2603. It also understood Emoji characters from Unicode's Astral Plane, which are often difficult for simple Unicode implementations. However, many Erlang functions do not understand Unicode, and one of the major updates planned to the language for R16 is more thorough Unicode support.

Getting the distance is somewhat easier:

```
get_distance() ->
    Input = io:get_line("How far? (meters) > "),
    {Distance, _} = string:to_integer(Value),
    Distance.
```

The `Input` variable collects the user's response to the question "How far?", and the `string:to_integer/1` function extracts an integer from that response. The pattern match on the left grabs the first piece of the tuple it returns, which is the integer, while the underscore discards the rest of what it sends, which is anything else on the line. That will include the newline, but also any decimal part users enter. You could use `string:to_float/1` for more precision, but that won't accept an integer. Using `string:to_integer/1` isn't perfect, but for these purposes it's probably acceptable.



It isn't necessary for this conversion, but if you *just* want to strip newlines out of user responses, you can use `string:strip(Input, right, $\n)`, where `Input` is what just came from the user.

A sample run demonstrates that it produces the right results given the right input.

```
1> c(ask).
{ok,ask}
2> ask:line().
Where are you?
1. Earth
2. Earth's Moon
3. Mars
Which? > 1
How far? (meters) > 20
19.79898987322333
3> ask:line().
Where are you?
1. Earth
2. Earth's Moon
3. Mars
Which? > 2
How far? > 20
8.0
```

Chapter 9 will return to this code, looking at better ways to handle the errors users can provoke by entering unexpected answers.

Strings are not Erlang’s strongest suit, but it has the facilities to make pretty much anything you need work. As you read the next two chapters on lists, remember that strings are actually lists underneath, and you can use any of the list tools on strings.

Erlang is great at handling lists, long series of similar (or not) values. List processing makes it easy to see the value of recursion, and offers opportunities to get a lot of work done for very little effort.

List Basics

An Erlang list is an ordered set of elements. Generally you will process a list in order, from the first item (the *head*) to the last item, though there are times when you may want to grab a particular item from the list. Erlang also provides built-in functions for manipulating lists when you don't want to go through the entire sequence.

Erlang syntax encloses lists in square brackets and separates elements with commas. A list of numbers might look like the following:

```
[1,2,4,8,16,32]
```

The elements can be of any type, including numbers, atoms, tuples, strings, and other lists. When you're starting out, it's definitely easiest to work with lists that contain only a single type of element, rather than mixing all the possibilities, but Erlang itself has no such constraint. There is also no limit on the number of items a list can contain, though eventually you may find practical limits of memory.

You can pattern match with lists just as you can with other Erlang data structures:

```
1> [1,X,4,Y] = [1,2,4,8].  
[1,2,4,8]  
2> X.  
2  
3> Y.  
8
```



While it's possible to use lists instead of tuples, your code will make more sense if you use tuples to handle data structures containing various kinds of data in a known structure, and lists to handle data structures containing less varied data in unknown quantities. (Tuples are expected to come in a certain order and can also contain lists, so if you have a data structure that's mostly known except for an expanding part or two, including a list inside of a tuple can be a workable solution.)

Lists can contain lists, and sometimes this can produce surprising results. If, for example, you want to add a list to a list, you may end up with more levels of list than you planned:

```
4> Insert=[2,4,8].  
[2,4,8]  
5> Full = [1, Insert, 16, 32].  
[1,[2,4,8],16,32]
```

You can fix that (if you want to) with the `lists:flatten/1` function:

```
6> Neat = lists:flatten(Full).  
[1,2,4,8,16,32]
```

This also means that if you want to append lists, you need to decide whether you're creating a list of lists or a single list containing the contents of the component lists. To create a list of lists, you just put lists into lists.

```
7> A = [1,2,4].  
[1,2,4]  
8> B = [8,16,32].  
[8,16,32]  
9> ListOfLists = [A,B].  
[[1,2,4],[8,16,32]]
```

To create a single list from multiple lists, you can use the `lists:append/2` function or the equivalent `++` operator.

```
10> Combined1 = lists:append(A,B).  
[1,2,4,8,16,32]  
11> Combined2 = A ++ B.  
[1,2,4,8,16,32]
```

Both produce the same result: a combined and flattened list.



The `++` operator is right associative, which can change the order of the resulting list when you append multiple lists.

If you have a set of lists you'd like combined, you can use the `lists:append/1` function, which takes a list of lists as its argument and returns a single list containing their contents:

```
12> C = [64,128,256].  
[64,128,256]  
13> Combined4 = lists:append([A,B,C]).  
[1,2,4,8,16,32,64,128,256]
```



If you want to generate a list of sequential integers (or characters), the `lists:seq/2` function is handy. Its arguments are the start and end of the list values. For example, `lists:seq(-2,8)` produces `[-2,-1,0,1,2,3,4,5,6,7,8]`, and `lists:seq($A,$z)` produces the string (list) "ABCDEFIGHJKLMNOPQRSTUVWXYZ[\\"`\`abcdefghijklm\nopqrstuvwxyz".

Splitting Lists into Heads and Tails

Lists are a convenient way to hold piles of similar data, but their great strength in Erlang is the way they make it easy to do recursion. Lists are a natural fit for the “counting down” style of logic explored in [Chapter 4](#): you can run through a list until you run out of items. In many languages, running through a list means finding out how many items it contains and going through them sequentially. Erlang takes a different approach, letting you process the first item in a list, the *head*, while extracting the rest of the list, the *tail*, so that you can pass it to another call recursively.

To extract the head and the tail, you use pattern matching, with a special form of the list syntax on the left:

```
[Head | Tail] = [1,2,4].
```

The two variables separated by a vertical bar (`|`), or cons, for list constructor, will be bound to the head and tail of the list on the right. In the console, Erlang will just report the contents of the right side of the expression, not the fragments created by the pattern match, but if you work through a list you can see the results:

```
1> List = [1,2,4].  
[1,2,4]  
2> [H1 | T1] = List.  
[1,2,4]  
3> H1.  
1  
4> T1.  
[2,4]  
5> [H2 | T2] = Tail1.  
[2,4]  
6> H2.  
2  
7> T2.  
[4]  
8> [H3 | T3] = T2.
```

```

[4]
9> H3.
4
10> T3.
[]
11> [H4 | T4] = T3.
** exception error: no match of right hand side value []

```

Line 2 copies the initial list into two smaller pieces. H1 will contain the first item of the List, whereas T1 will contain a list that has everything *except* the first element. Line 5 repeats the process on the smaller list, breaking T1 into a H2 and a T2. This time T2 is still a list, as shown on line 7, but contains only one item. Line 8 breaks that single-item list again, putting the value into H3 and an *empty* list into T3.

What happens when you try to split an empty list, as shown on line 11? Erlang reports an error, “no match...” This fortunately does not mean that recursion on lists is doomed to produce errors. That lack of a match will naturally stop the recursive process, which is probably what you want.



Head and tail work only moving forward through a list. If order matters and you really need to go through a list backwards, you’ll need to use the `lists:reverse` function and then walk through the reversed list.

Processing List Content

The head and tail notation was built for recursive processing. Actually making that work typically follows a pattern in which a list arrives as an argument and is then passed to another (usually private) function with an accumulator argument. A simple case might perform a calculation on the contents of the list. *Example 6-1*, in `ch06/ex1-product`, shows this pattern in use, multiplying the values of a list together.

Example 6-1. Calculating the product of values in a list

```

-module(overall).
-export([product/1]).

product([]) -> 0; % in case the list is empty, return zero
product(List) -> product(List,1).

product([], Product) -> Product; % when list empty, stop, report

product([Head|Tail], Product) -> product(Tail, Product * Head).

```

In this module, the `product/1` function is the gateway, passing the list (if the list has content) plus an accumulator to `product/2`, which does the real work. If you wanted to test the arriving list to make sure it meets your expectations, it probably makes the most sense to do that work in `product/1`, and let `product/2` focus on recursive processing.



Is the product of an empty list really zero? It might make more sense for an empty list to fail and produce a crash. Erlang's "let it crash" philosophy is, as you'll see later, pretty calm about such things. In the long run, you'll have to decide which cases are better left to crash and which shouldn't.

The `product/2` function has two clauses. The first matches the empty list, and will get called at the end of the recursive process when there are no more entries to process, or if the list arrives empty. It returns its second argument, the accumulator.

The second clause does more work if the arriving list is not empty. First, the pattern match (`[Head|Tail]`) splits off the first value the list from the rest of the list. Next, it calls `product/2` again, with the remaining (if any) portion of the list and a new accumulator that is multiplied by the value of the first entry in the list. The result will be the product of the values included in the list:

```
1> c(overall).
{ok,overall}
2> overall:product([1,2,3,5]).
30
```

That went smoothly, but what happened? After `product/1` called `product/2`, it made five iterations over the list, concluding with an empty list, as shown in [Table 6-1](#).

Table 6-1. Recursive processing of a simple list in `product/2`

Arriving List	Arriving Product	Head	Tail
[1,2,3,5]	1	1	[2,3,5]
[2,3,5]	1 (1*1)	2	[3,5]
[3,5]	2 (1*2)	3	[5]
[5]	6 (2*3)	5	[]
[]	30 (6*5)	None	None

The last arriving `Product`, `30`, will be handled by the clause for the empty list and reported as the return value for `product/2`. When `product/1` receives that value, it will also report `30` as its return value and exit.



Because strings are lists, you can do strange things like enter over `all:product("funny")`. `product/1` will interpret the character values as numbers, and return `17472569400`.

Creating Lists with Heads and Tails

While there are times you want to calculate a single value from a list, much list processing involves modifying lists or converting a list into another list. Because you can't actually change a list, modifying or converting a list means creating a new list. To do that, you use the same vertical bar head/tail syntax, but on the right side of the pattern match instead of the left. You can try this out in the console, though it's more useful in a module:

```
1> X=[1|[2,3]].  
[1,2,3]
```

Erlang interprets `[1|[2,3]]` as creating a list. If the value to the right of the vertical bar is a list, it gets appended to the head as a list. In this case, the result is a neat list of numbers. There are a few other forms you should be aware of:

```
2> Y=[1,2 | [3]].  
[1,2,3]  
3> Z=[1,2 | 3].  
[1,2|3]
```

In line 2, there isn't a list wrapped around the now two items in the head, but the constructor still blends the head and the tail together seamlessly. (If you do wrap them in square brackets, the list constructor assumes that you want a list as the first item in the list, so `[[1,2] | [3]]` will produce `[[1,2],3]`.)

However, line 3 demonstrates what happens if you don't wrap the tail in square brackets —you get a list, called an *improper list*, that still contains a constructor, with a strange tail. Until you've learned your way quite thoroughly around Erlang, you probably should avoid this, as it will create runtime errors if you try to process it as a normal list. Eventually you may find rare reasons to do this, or encounter code that uses it.

More typically, you'll use list constructors to build lists inside of recursive functions. **Example 6-2**, which you can find in `ch06/ex2-drop`, starts from a set of tuples representing planemos and distances. With the help of the `drop` module from **Example 3-8**, it creates a list of velocities for the corresponding falls.

Example 6-2. Calculating a series of drop velocities, with an error

```
-module(listdrop).  
-export([falls/1]).
```

```

falls(List) -> falls(List,[]).

falls([], Results) -> Results;
falls([Head|Tail], Results) -> falls(Tail, [drop:fall_velocity(Head) | Results]).
```

Much of this is familiar from [Example 6-1](#), except that the `Results` variable gets a list instead of a number, and the last line of `falls/2` creates a list instead of a single value. If you run it, however, you'll see one minor problem:

```

1> c(drop).
{ok,drop}
2> c(listdrop).
{ok,listdrop}
3> listdrop:falls([{earth,20},{moon,20},{mars,20}]).
```

[12.181953866272849,8.0,19.79898987322333]

The resulting velocities are reversed: the Earth has more gravity than Mars, and objects should fall faster on Earth. What happened? That last key line in `falls/2` is reading a list from the beginning to the end, and creating a list from the end to the beginning. That puts the values in the wrong order. Fortunately, as [Example 6-3](#) demonstrates, this is easy to fix. You need to call `lists:reverse/1` in the clause of the `falls/2` function that handles the empty list.

Example 6-3. Calculating a series of drop velocities, with the error fixed

```

-module(listdrop).
-export([falls/1]).

falls(List) -> falls(List,[]).

falls([], Results) -> lists:reverse(Results);
falls([Head|Tail], Results) -> falls(Tail, [drop:fall_velocity(Head) | Results]).
```

Now it works:

```

4> c(listdrop).
{ok,listdrop}
5> listdrop:falls([{earth,20},{moon,20},{mars,20}]).
```

[19.79898987322333,8.0,12.181953866272849]



You could instead have put the `lists:reverse/1` call in the `falls/1` gateway function. Either way is fine, though I prefer to have `falls/2` return a finished result.

Mixing Lists and Tuples

As you get deeper into Erlang and pass around more complex data structures, you may find that you're processing lists full of tuples, or that it would be more convenient to rearrange two lists into a single list of tuples or vice-versa. The `lists` module includes easy solutions to these kinds of transformations and searches.

The simplest set of tools are the `lists:zip/2` and `lists:unzip/1` functions. They can turn two lists of the same size into a list of tuples or a list of tuples into two lists.

```
1> List1=[1,2,4,8,16].  
[1,2,4,8,16]  
2> List2=[a,b,c,d,e].  
[a,b,c,d,e]  
3> TupleList=lists:zip(List1,List2).  
[{1,a},{2,b},{4,c},{8,d},{16,e}]  
4> SeparateLists=lists:unzip(TupleList).  
{[1,2,4,8,16],[a,b,c,d,e]}
```

The two lists, `List1` and `List2`, have different contents, but the same number of items. The `lists:zip/2` function returns a list containing a tuple for each of the items in the original list. The `lists:unzip/1` function takes that list of two-component tuples, and splits it out into a tuple containing two lists.



Erlang also provides `lists:zip3/3` and `lists:unzip3/1`, which do the same combining and separating on sets of three lists or tuple values.

You will also likely encounter times when you need to process a different kind of list containing tuples, a collection of values identified by keys. Many languages include associative arrays, where access is provided through key values, but Erlang's lists are always sequential, and have no built-in concept of retrieving information with a key. However, the `lists` module provides functions that support treating a list of tuples as if it were a key/value store, such as hash table, hash tree, or associative array.

You don't need to do anything special to create a list that you'll use this way, except to make sure that the key is in a consistent location in the tuples stored in the list. Because the functions let you specify the location, you can put them anywhere in the tuple as long as you're consistent, but for this example, they'll be in the first position.

The first function to explore is `lists:keystore/4`. It takes a key value, a position, a list that is the previous state of the key/value store (defined in line 1 of the following code

sample), and a tuple. If no tuple has that key value, then the new tuple simply gets added to the list, as shown in line 2 of the following code sample. If, as in line 3, a tuple already has that key value, the function will return a list that replaces the matched tuple with the new one.

```
1> Initial=[[1,tiger], {3,bear}, {5,lion}].  
[[1,tiger],{3,bear},{5,lion}]  
2> Second=lists:keystore(7,1,Initial,{7,panther}).  
[{1,tiger},{3,bear},{5,lion},{7,panther}]  
3> Third=lists:keystore(7,1,Second,{7,leopard}).  
[{1,tiger},{3,bear},{5,lion},{7,leopard}]
```

You can also pass `lists:keystore/4` an empty list for the array, and it will just return a list containing the new tuple.

Sometimes you want to replace a value *only* if it is present, not add a new value to the list. The similar `lists:keyreplace/4` will do just that.

```
4> Fourth=lists:keyreplace(6,1,Third,{6,chipmunk}).  
[{1,tiger},{3,bear},{5,lion},{7,leopard}]
```

There was no item in the previous list with a key value of 6, so `lists:keyreplace/4` just returned a copy of the original list.



All of these functions are copying lists or creating new modified versions of a list. As you'd expect in Erlang, the original list is untouched.

If you want to get information back out of a list, the `lists:keyfind/3` argument will report the data that matches a given key:

```
5> Animal5=lists:keyfind(5,1,Third).  
{5,lion}
```

If the key isn't present, however, you'll just get a return value of `false`, instead of a tuple.

```
6> Animal6=lists:keyfind(6,1,Third).  
false
```

Building a List of Lists

While simple recursion isn't too complicated, list processing has a way of turning into lists of lists in various stages. Pascal's triangle, a classic mathematical tool, is relatively simple to create but demonstrates more intricate work with lists. It starts with a 1 at the top, and then each new row is composed of the sum of the two numbers above it:

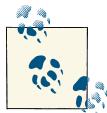
```

      1
     1   1
    1   2   1
   1   3   3   1
  1   4   6   4   1
...

```

If those numbers seem familiar, it's probably because they're the binomial coefficients that appear when you put $(x+y)$ to a power. That's just the beginning of this mathematical marvel, described in more detail at http://en.wikipedia.org/wiki/Pascal's_triangle.

This is easily calculated with Erlang in a number of ways. You can apply the list techniques already discussed in this chapter by treating each row as a list, and the triangle as a list of lists. The code will be seeded with the first row—the top 1—represented as `[0, 1, 0]`. The extra zeros make the addition much simpler.



This is not intended to be an efficient, elegant, or maximally compact implementation. At this point, a naive implementation likely explains more about lists. Once this makes sense, and you learn about list comprehensions in [Chapter 7](#), you can explore what a vastly more compact version might look like, see http://rosettacode.org/wiki/Pascal's_triangle#Erlang.

For a first step, [Example 6-4](#) calculates rows individually. This is a simple recursive process, walking over the old list and adding its contents to create a new list.

Example 6-4. Calculating a row

```

-module(pascal).
-export([add_row/1]).
add_row(Initial) -> add_row(Initial, 0, []).

add_row([], 0, Final) -> [0 | Final];

add_row([H | T], Last, New) -> add_row(T, H, [Last + H | New]).
```

The `add_row/1` function sets things up, sending the current row, a `0` to get the math started, and an empty list you can think of as “where the results go,” though it is really an accumulator. The `add_row/3` function has two clauses. The first checks to see if the list being added is empty. If it is, then it reports back the final row, adding a `0` at the front.

Most of the work gets done in the second clause of `add_row/3`. When it receives its arguments, the `[H | T]` pattern match splits the head of the list into the `H` value (a number) and the tail into `T` (a list, which may be empty if that was the last number). It also gets values for the `Last` number processed and the current `New` list being built.

It then makes a recursive call to `add_row/3`. In that new call, the tail of the old list, `T`, is the new list to process, the `H` value becomes the `Last` number processed, and the third argument, the list, opens with the actual addition being performed, which is then combined with the rest of the `New` list being built.



Because the lists in the triangle are symmetrical, there is no need to use `lists:reverse/1` to flip them. You can, of course, if you want to.

You can test this easily from the console, but remember that your test lists need to be wrapped in zeros:

```
1> c(pascal).
{ok,pascal}
2> pascal:add_row([0,1,0]).
[0,1,1,0]
3> pascal:add_row([0,1,1,0]).
[0,1,2,1,0]
4> pascal:add_row([0,1,2,1,0]).
[0,1,3,3,1,0]
```

Now that you can create a new row from an old one, you need to be able to create a set of rows from the top of the triangle, as shown in [Example 6-5](#), which you can find in `ch06/ex4-pascal`. The `add_row/3` function effectively counted down to the end of the list, but `triangle/3` will need to count up to a given number of rows. The `triangle/1` function sets things up, defining the initial row, setting the counter to 1 (because that initial row *is* the first row), and passing on the number of `Rows` to be created.

The `triangle/3` function has two clauses. The first, the stop clause, halts the recursion when enough `Rows` have been created, and reverses the list. (The individual rows may be symmetrical, but the triangle itself is not.) The second clause does the actual work of generating new rows. It gets the `Previous` row generated from the `List`, and then it passes that to the `add_row/1` function, which will return a new row. Then it calls itself with the new list, an incremented `Count`, and the `Rows` value the stop clause needs.

Example 6-5. Calculating the whole triangle with both functions

```
-module(pascal).
-export([triangle/1]).

triangle(Rows) -> triangle([[0,1,0]],1,Rows).

triangle(List, Count, Rows) when Count >= Rows -> lists:reverse(List);

triangle(List, Count, Rows) ->
  [Previous | _] = List,
  triangle([add_row(Previous) | List], Count+1, Rows).
```

```
add_row(Initial) -> add_row(Initial, 0, []).  
add_row([], 0, Final) -> [0 | Final];  
add_row([H | T], Last, New) -> add_row(T, H, [Last + H | New]).
```

Happily, this works.

```
5> c(pascal).  
{ok,pascal}  
6> pascal:triangle(4).  
[[0,1,0],[0,1,1,0],[0,1,2,1,0],[0,1,3,3,1,0]]  
7> pascal:triangle(6).  
[[0,1,0],  
 [0,1,1,0],  
 [0,1,2,1,0],  
 [0,1,3,3,1,0],  
 [0,1,4,6,4,1,0],  
 [0,1,5,10,10,5,1,0]]
```

Pascal's triangle may be a slightly neater set of lists than most you will process, but this kind of layered list processing is a very common tactic for processing and generating lists of data.

Higher-Order Functions and List Comprehensions

Higher-order functions, functions that accept other functions as arguments, are a key place where Erlang’s power really starts to shine. It’s not that you can’t do higher-order functions in other languages—you can in many—but more that Erlang treats higher-order functions as a native and natural part of the language rather than an oddity.

Simple Higher-Order Functions

Way back in [Chapter 2](#), you saw how to use a `fun` to create a function:

```
1> Fall_velocity = fun(Distance) -> math:sqrt(2 * 9.8 * Distance) end.  
#Fun<erl_eval.6.111823515>  
2> Fall_velocity(20).  
19.79898987322333  
3> Fall_velocity(200).  
62.609903369994115
```

Erlang not only lets you put functions into variables, it lets you pass functions as arguments. This means that you can create functions whose behavior you modify at the time you call it, in much more intricate ways than is normally possible with parameters. A very simple function that takes another function as an argument might look like [Example 7-1](#), which you can find in `ch07/ex1-hof`.

Example 7-1. An extremely simple higher-order function

```
-module(hof).  
-export([tripler/2]).  
  
tripler(Value, Function) -> 3 * Function(Value).
```

The argument names are generic, but fit. `tripler/2` will take a value and a function as arguments. It runs the value through the function, and multiplies that result by three. In the shell, this might look like the following:

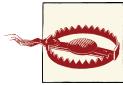
```
1> c(hof).
{ok,hof}
2> MyFunction=fun(Value)->20*Value end.
#Fun<erl_eval.6.111823515>
3> hof:tripler(6,MyFunction).
360
```

That defines another simple function taking one argument (and returning that number multiplied by 20), and stores it in the variable `MyFunction`. Then it calls the `hof:tripler/2` function with a value of six and the `MyFunction` function. In the `hof:tripler/2` function, it feeds the `Value` to the `Function`, getting back 120. Then it triples that, returning 360.

You can skip assigning the function to a variable if you want, and just include the `fun` declaration inside the `hof:tripler/2` function call:

```
4> hof:tripler(6,fun(Value)->20*Value end).
360
```

That may or may not be easier to read, depending on the functions and your expectations. This case is trivially simple, but demonstrates that it works.



While this is a powerful technique, you can outsmart yourself with it easily. (I do!) Just as with normal code, you need to make sure the number and sometimes the type of your arguments line up. The extra flexibility and power can create new problems if you aren't careful.

`fun` has a few other tricks up its sleeve that you should know. You can use a `fun` to preserve context, even context that has since vanished.

```
5> X=20.
20
6> MyFunction2=fun(Value)->X * Value end.
#Fun<erl_eval.6.82930912>
7> f(X).
ok
8> X.
* 1: variable 'X' is unbound
9> hof:tripler(6,MyFunction2).
360
```

Line 5 assigns a variable named `X` a value, and line 6 uses that variable in a `fun`. Line 7 obliterates the `X` variable, as line 8 demonstrates, but line 9 shows that `MyFunction2` still remembers that `X` was 20. Even though the value of `X` has been flushed from the shell, the `fun` preserves the value and can act upon it. (This is called a *closure*.)

You may also want to pass a function from a module, even a built-in module, to your (or any) higher-order function. That's simple, too:

```
7> hof:tripler(math:pi(), fun math:cos/1).
-3.0
```

In this case, the `hof:tripler` function receives the value `pi` and a `fun`, which is the `math:cos/1` function from the built-in `math` module. Since the cosine of `pi` is `-1`, the `tripler` returns `-3.0`.

Creating New Lists with Higher-Order Functions

Lists are one of the best and easiest places to apply higher-order functions. Applying a function to all the components of a list to create a new list, sort a list, or break a list into smaller pieces is popular work. You don't need to do much difficult work to make this happen, though: Erlang's built-in `lists` module offers a variety of higher-order functions, listed in [Appendix A](#), that take a function and list and do something with them. You can also use *list comprehensions* to do much of the same work. The `lists` module may seem easier at first, but as you'll see, list comprehensions are powerful and concise.

Reporting on a List

The simplest of these functions is `foreach/2`, which always returns the atom `ok`. That may sound strange, but `foreach/2` is a function you'll call if and only if you want to do something to the list with side effects—like present the contents of a list to the console. To do that, define a simple function that applies `io:format/2`, here stored in the variable `Print`, and a `List`, and then pass them both to `lists:foreach/2`.

```
1> Print = fun(Value) -> io:format(" ~p~n",[Value]) end.
#Fun<erl_eval.6.111823515>
2> List = [1,2,4,8,16,32].
[1,2,4,8,16,32]
3> lists:foreach(Print,List).
1
2
4
8
16
32
ok
```

The `lists:foreach/2` function walked through the list, in order, and called the function in `Print` with each item of the list as a `Value`. The `io:format/2` function inside of `Print` presented the list item, slightly indented. When it reached the end of the list, `lists:foreach/2` returned the value `ok`, which the console also displayed.



Most of the demonstrations in this chapter will be operating on that same `List` variable, containing `[1,2,4,8,16,32]`.

Running List Values Through a Function

You might also want to create a new list based on what a function does with all of the values in the original list. You can square all of the values in a list by creating a function that returns the square of its argument, and passing that to `lists:map/2`. Instead of returning `ok`, it returns a new list reflecting the work of the function it was given:

```
4> Square = fun(Value)->Value*Value end.  
#Fun<erl_eval.6.111823515>  
5> lists:map(Square,List).  
[1,4,16,64,256,1024]
```

There's another way to accomplish the same thing, with what Erlang calls a *list comprehension*.

```
6> [Square(Value) || Value <- List].  
[1,4,16,64,256,1024]
```

That produces the same resulting list, with different (and more flexible) syntax. While you saw the `[A | B]` syntax in list constructors, a list comprehension uses `[A || B]` syntax. That extra vertical bar changes the whole way this is interpreted. Instead of being a head and a tail, it's an expression—here a function—and a rule for extracting the arguments for that function from a list, called a *generator*.

In this case, the function is the `fun` you put in the `Square` variable on line 4. Its argument, `Value`, is taken on a walk through the `List`. That arrow—the `<-`—means “an element of,” or if you want be more active, “comes from.” You can read this list comprehension as “Create a list consisting of squares of a `Value`, where the `Value` comes from `List`.”

Strictly speaking, the expression on the left doesn't have to be formally declared as a function. You can get the same results with something less formal:

```
7> [Value * Value || Value <- List].  
[1,4,16,64,256,1024]
```



The multiplication operator (*) is technically a call to the */2 function, but any legal Erlang expression can be on the left of the ||.

Filtering List Values

The `lists` module offers a few different functions for filtering the content of a list based on a function you provide as a parameter. The most obvious, `lists:filter/2`, returns a list composed of the members of the original list for which the function returned `true`. For example, if you wanted to filter a list of integers down to values that could be represented as four binary digits, so numbers 0 or greater but less than 16, you could define a function and store it in `Four_bits`:

```
8> Four_bits = fun(Value)-> (Value<16) and (Value>=0) end.  
#Fun<erl_eval.6.111823515>
```

Then, if you apply it to the previously defined `List` of `[1,2,4,8,16,32]`, you'd get just the first four values:

```
9> lists:filter(Four_bits,List).  
[1,2,4,8]
```

Once again, you can create the same effect with a list comprehension. This time, you don't actually need to create a function, but instead use a guard-like construct (written *without* the `when`) on the right side of the comprehension:

```
10> [Value || Value <- List, Value<16, Value>=0].  
[1,2,4,8]
```



If you also want a list of values that didn't match, `lists:partition/2` will return a tuple that contains the matched items in its first element, and the unmatched items in its second.

Beyond List Comprehensions

List comprehensions are concise and powerful, but they lack a few key features available in other recursive processing. The only type of result they can return is a list, but there will be many times when you want to process a list and return something else, like a boolean, a tuple, or a number. List comprehensions also lack support for accumulators, and don't let you suspend processing completely when certain conditions are met.

You could write your own recursive functions to process lists, but much of the time you'll find that the `lists` module already offers a function that takes a function you define and a list and returns what you need.

Testing Lists

Sometimes you just want to know if all the values—or any of the values—in a list meet specific criteria. Are they all of a specific type, or do they have a value that meets certain criteria?

The `lists:all/2` and `lists:any/2` functions let you test a list against rules you specify in a function. If your function returns `true` for all of the list values, both of these functions will return `true`. `lists:any/2` will also return `true` if one or more values in the list results in your function returning `true`. Both will return `false` if your function consistently returns `false`.



`lists:all/2` and `lists:any/2` don't necessarily evaluate the entire list; as soon as they hit a value that provides a definitive answer, they'll stop and return that answer.

```
11> IsInt = fun(Value) -> is_integer(Value) end.  
#Fun<erl_eval.6.111823515>  
12> lists:all(IsInt, List).  
true  
13> lists:any(IsInt, List).  
true  
14> Compare = fun(Value) -> Value > 10 end.  
#Fun<erl_eval.6.111823515>  
15> lists:any(Compare, List).  
true  
16> lists:all(Compare, List).  
false
```

You can think of `lists:all/2` as an `and` function applied to lists; more precisely like `andalso` because it stops processing as soon as it encounters a false result. Similarly, `lists:any/2` is like `or`, or `orelse`, in this case stopping as soon as it finds a true result. As long as you need only to test individual values within lists, these two higher order functions can save you writing a lot of recursive code.

Splitting Lists

Filtering lists is useful, but sometimes you want to know what didn't go through the filter, and sometimes you just want to separate items.

The `lists:partition/2` function returns a tuple containing two lists. The first is the list items that met the conditions specified in the function you provided, while the

second is the items that didn't. If the `Compare` variable is defined as shown in line 14 of the previous demonstration, returning `true` when a list value is greater than 10, then you can split a list into a list of items greater than 10 and a list of items fewer than 10 easily:

```
17> lists:partition(Compare,List).  
 {[16,32],[1,2,4,8]}
```

Sometimes you'll want to split a list by starting from the beginning—the head—and stopping when a list value no longer meets a condition. The `lists:takewhile/2` and `lists:dropwhile/2` functions create a new list that contains the parts of an old list before or after encountering a boundary condition. These functions aren't filters, and to make that clear, the examples use a different list than the rest in this chapter.

```
18> Test=fun(Value) -> Value < 4 end.  
#Fun<erl_eval.6.111823515>  
19> lists:dropwhile(Test, [1,2,4,8,4,2,1]).  
[4,8,4,2,1]  
20> lists:takewhile(Test, [1,2,4,8,4,2,1]).  
[1,2]
```

Both functions run through a list from head to tail and stop when they reach a value for which the function you provide as the first argument returns `false`. The `lists:dropwhile/2` function returns what's left of the list, including the value that flunked the test. It does not, however, filter out later list entries that it might have dropped if they had appeared earlier in the list. The `lists:takewhile/2` function returns what was already processed, *not* including the value that flunked the test.

Folding Lists

Adding an accumulator to list processing lets you turn lists into much more than other lists, and opens the door to much more sophisticated processing. Erlang's `lists:foldl/3` and `lists:foldr/3` functions let you specify a function, an initial value for an accumulator, and a list. Instead of the one-argument functions you've seen so far, you need to create a two argument function, accepting the current value in the list traversal and the accumulator. The result of that function will become the new value of the accumulator.

Defining a function that works within the folding functions looks a little different, because of the two arguments:

```
21> Divide=fun(Value, Accumulator) -> Value / Accumulator end.  
#Fun<erl_eval.6.111823515>
```

This function divides its first argument—to be the `Value` coming from the list—by second, the `Accumulator` passed to it by the function doing folding.

Folding has one other key twist. You can choose whether you want the function to traverse the list from head to tail, with `lists:foldl/3`, or from tail to head, with `lists:foldr/3`. If order doesn't change the result, you should go with `lists:foldl/3`, as its implementation is tail-recursive and more efficient in most situations.

The `Divide` function is one of those cases that will produce very different results depending on the direction in which you process the list (and the initial accumulator value). In this case, folding also produces different results than you might expect in a simple division. Given the usual List of `[1,2,4,8,16,32]`, it seems like going from left to right will produce `1/2/4/8/16/32`, and going from right to left will produce `32/16/8/4/2/1`, at least if you use an initial accumulator of 1. They don't produce those results, however.

```
22> 1/2/4/8/16/32.  
3.0517578125e-5  
23> lists:foldl(Divide,1,List).  
8.0  
24> 32/16/8/4/2/1.  
0.03125  
25> lists:foldr(Divide,1,List).  
0.125
```

This code seems too simple to have a bug, so what's going on? [Table 7-1](#) walks through the calculations for `lists:foldl(Divide,1,List)`, and [Table 7-2](#) walks through `lists:foldr(Divide,1,List)` step by step.

Table 7-1. Recursive division of a list forwards with foldl/3

Value from List	Accumulator	Result of Division
1	1	1
2	1(1/1)	2
4	2(2/1)	2
8	2(4/2)	4
16	4(8/2)	4
32	4	8

Table 7-2. Recursive division of a list backwards with foldr/3

Value from List	Accumulator	Result of Division
32	1	32
16	32 (32/1)	0.5
8	0.5 (32/16)	16
4	16 (8/0.5)	0.25
2	0.25 (4/16)	8
1	8	0.125

Moving through a list step-by-step produces very different values. In this case, the simple `Divide` function's behavior changes drastically above and below the value 1, and combining that with walking through a list item by item yields results that might not be precisely what you expected.



The result of the `foldl` is the same as $32/(16/(8/(4/(2/(1/1)))))$, while the result of the `foldr` is the same as $1/(2/(4/(8/(16/(32/1)))))$. The parentheses in those perform the same restructuring as the fold, and the concluding 1 in each is where the initial accumulator value fits in.

Folding is an incredibly powerful operation. This simple if slightly weird example just used a single value, a number, as an accumulator. If you use a tuple as the accumulator, you can store all kinds of information about a list as it passes by, and even perform multiple operations. You probably won't want to try to define the functions you use for that as one-liners, but the possibilities are endless.

Playing with Processes

Erlang is a functional language, but Erlang programs are rarely structured around simple functions. Instead, Erlang's key organizational concept is the *process*, an independent component (built from functions) that sends and receives messages. Programs are deployed as sets of processes that communicate with each other. This approach makes it much easier to distribute work across multiple processors or computers, and also makes it possible to do things like upgrade programs in place without shutting down the whole system.

Taking advantage of those features, though, means learning how to create (and end) processes, how to send messages among them, and how to apply the power of pattern matching to incoming messages.

The Shell Is a Process

You've been working within a single process throughout this book so far, the Erlang shell. None of the previous examples sent or received messages, of course, but the shell is an easy place to send and (for test purposes, at least) receive messages.

The first thing to explore is the *process identifier*, often called a *pid*. The easiest pid to get is your own, so in the shell you can just try the `self()` function:

```
1> self().  
<0.36.0>
```

`<0.36.0>`, is the shell's representation of a *tuple*, a set of three integers that provide the unique identifier for this process. You may get a different set of numbers when you try it. This group of numbers is guaranteed to be unique within this run of Erlang, not

permanently the same in future use. Erlang uses pids internally, but while you can read them in the shell, you can't type pids directly into the shell or into functions. Erlang much prefers that you treat pids as abstractions, though if you really want to address a process by its pid numbers, you can use the `pid/3` shell function to do so.



Pids can even identify processes running on multiple computers within a cluster. You'll need to do more work to set up a cluster, but you won't have to throw away code you wrote with pids and processes built on them when you get there.

Every process gets its own pid, and those pids function like addresses for mailboxes. Your programs will send messages from one process to another by sending them to a pid. When that process gets time to check its mailbox, it will be able to retrieve and process the messages there.

Erlang, however, will *never* report that a message send failed, even if the pid doesn't point to a real process. It also won't report that a message was ignored by a process. You need to make sure your processes are assembled correctly.

The syntax for sending a message is pretty simple: a function or variable containing the pid, plus the send operator (!) and the message.

```
2> self() ! test1.  
test1  
3> Pid=self().  
<0.36.0>  
4> Pid ! test2.  
test2
```

Line 2 sent a message to the shell containing the atom `test1`. Line 3 assigned the pid for the shell, retrieved with the `self()` function, to a variable named `Pid`, and then line 4 used that `Pid` variable to send a message containing the atom `test2`. (The `!` always returns the message, which is why it appears right after the sends in lines 2 and 4.)

Where did those messages go? What happened to them? Right now, they're just waiting in the shell's mailbox, doing nothing.

There's a shell function—`flush()`—that you can use to see what's in the mailbox, though it also removes those messages from the mailbox. The first time you use it, you'll get a report of what's in the mailbox, but the second time, the messages are gone, already read.

```
5> flush().  
Shell got test1  
Shell got test2  
ok  
6> flush().  
ok
```

The proper way to read the mailbox, which gives you a chance to do something with the messages, is the `receive...` `end` construct. You can test this out in the shell. The first of the following tests just reports what the message was, whereas the second expects a number and doubles it.

```
7> self() ! test1.  
test1  
8> receive X -> X end.  
test1  
9> self() ! 23.  
23  
10> receive Y->2*Y end.  
46
```

So far, so good. However, if you screw up—if there isn’t a message waiting, or if you provide a pattern match that doesn’t work—the shell will just sit there, hung. Actually, it’s waiting for something to arrive in the mailbox, but you’ll be stuck. The easiest way out of that is to hit Ctrl-G, and then type q. You’ll have to restart Erlang. (X and Y become bound variables, so don’t try to reuse them, either.)

Spawning Processes from Modules

While sending messages to the shell is an easy way to see what’s happening, it’s not especially useful. Processes at their heart are just functions, and you know how to build functions in modules. The `receive...end` statement is structured like a `case...end` statement, so it’s easy to get started.

Example 8-1, which is in *ch08/ex1-simple*, shows a simple—excessively simple—module containing a function that reports messages it receives.

Example 8-1. An overly simple process definition

```
-module(bounce).  
-export([report/0]).  
  
report() ->  
    receive  
        X -> io:format("Received ~p~n",[X])  
    end.
```

When the `report/0` function receives a message, it will report that it received it. Setting this up means compiling it and then using the `spawn/3` function, which turns the function into a free-standing process. The arguments are the module name, the function name, and a list of arguments for the function. (Even if you don’t have any arguments, you need to include an empty list in square brackets.) The `spawn/3` function will return the Pid, which you should capture in a variable, here `Pid`:

```
1> c(bounce).
{ok,bounce}
2> Pid=spawn(bounce,report,[[]]).
<0.38.0>
```

Once you have the process spawned, you can send a message to that pid, and it will report that it received it:

```
3> Pid ! 23.
Received 23
23
ok
```

However, there's one small problem. The `report` process exited—it went through the `receive` clause only once, and when it was done, was done. If you try to send it a message, you'll get back the message, and nothing will report an error, but you also won't get any notification that the message was received because nothing is listening any longer.

```
4> Pid ! 23.
23
```

To create a process that keeps processing messages, you need to add a recursive call, as shown in the `receive` statement in [Example 8-2](#), in *ch08/ex2-recursion*.

Example 8-2. A function that creates a stable process

```
-module(bounce).
-export([report/0]).

report() ->
    receive
        X -> io:format("Received ~p~n",[X]),
        report()
    end.
```

That extra call to `report()` means that after the function shows the message that arrived, it will run again, ready for the next message. If you recompile the `bounce` module and spawn it to a new `Pid2` variable, you can send it multiple messages, as shown here.

```
5> c(bounce).
{ok,bounce}
6> Pid2=spawn(bounce,report,[[]]).
<0.47.0>
7> Pid2 ! 23.
Received 23
23
ok8> Pid2 ! message.
Received message
message
```

You can also pass an accumulator from call to call if you want, for a simple example, to keep track of how many messages have been received by this process. [Example 8-3](#) shows the addition of an argument, in this case just an integer that gets incremented with each call. You can find it in *ch08/ex3-counter*.

Example 8-3. A function that adds a counter to its message reporting

```
-module(bounce).
-export([report/1]).  
  
report(Count) ->
receive
    X -> io:format("Received #~p: ~p~n",[Count,X]),
        report(Count+1)
end.
```

The results are pretty predictable, but remember that you need to include an initial value in the arguments list in the `spawn/3` call.

```
1> c(bounce).
{ok,bounce}
2> Pid2=spawn(bounce,report,[1]).
<0.38.0>
3> Pid2 ! test.
Received #1: test
test
ok4> Pid2 ! test2.
Received #2: test2
test2
ok5> Pid2 ! another.
Received #3: another
```

Whatever you do in your recursive call, keeping it simple (and preferably tail-recursive) is best, as these can get called many, many times in the life of a process.



If you want to create impatient processes that stop after waiting a given amount of time for a message, you should investigate the `after` construct of the `receive` clause.

You can write this function in a slightly different way that may make what's happening clearer and easier to generalize. [Example 8-4](#), in *ch08/ex4-state*, shows how to use the return value of the `receive` clause, here the `Count` plus one, to pass state from one iteration to the next.

Example 8-4. Using the return value of the receive clause as state for the next iteration

```
-module(bounce).
-export([report/1]).
```

```

report(Count) ->
    NewCount = receive
        X -> io:format("Received #~p: ~p~n", [Count, X]),
        Count + 1
    end,
    report (NewCount).

```

In this model, all (though just one here) of the `receive` clauses return a value that gets passed to the next iteration of the function. If you use this approach, you can think of the return value of the `receive` clause as the state to be preserved between function calls. That state can be much more intricate than a counter—it might be a tuple, for instance, that includes references to important resources or work in progress.

Lightweight Processes

If you've worked in other programming languages, you may be getting worried. Threads and process spawning are notoriously complex and often slow in other contexts, but Erlang expects applications to be a group of easily spawned processes? That run recursively?

Yes, absolutely. Erlang was written specifically to support that model, and its processes are more lightweight than pretty much any of its competitors. The Erlang scheduler gets processes started and distributes processing time among them, as well as splitting them out across multiple processors.

It is certainly possible to write processes that perform badly and to structure applications so that they wait a long time before doing anything. You don't, though, have to worry about those problems happening just because you're using multiple processes.

Registering a Process

Much of the time, pids are all you need to find and contact a process. However, you will likely create some processes that need to be more findable. Erlang provides a process registration system that is extremely simple: you specify an atom and a pid, and then any process that wants to reach that registered process can just use the atom to find it. This makes it easier, for example, to add a new process to a system and have it connect with previously existing processes.

To register a process, just use the `register/2` built-in function. The first argument is an atom, effectively the name you're assigning the process, and the second argument is the pid of the process. Once you have it registered, you can just send it messages, using the atom instead of a pid:

```

1> Pid1=spawn(bounce,report,[1]).
<0.33.0>
2> register(bounce,Pid1).

```

```
true
3> bounce ! hello.
Received #1: hello
hello
ok4> bounce ! "Really?".
Received #2: "Really?"
"Really?"
ok
```

If you attempt to call a process that doesn't exist (or one that has crashed), you'll get a bad arguments error:

```
6> zingo ! test.
** exception error: bad argument
   in operator  !/2
      called as zingo ! test
```

If you attempt to register a process to a name that is already in use, you'll also get an error, but if a process has exited (or crashed), the name is effectively no longer in use and you can re-register it.

You can also use `whereis/1` to retrieve the pid for a registered process (or `undefined`, if there is no process registered with that atom), and `unregister/1` to take a process out of the registration list without killing it.

```
5> GetBounce = whereis(bounce).
<0.33.0>
6> unregister(bounce).
true
7> TestBounce = whereis(bounce).
undefined
8> GetBounce ! "Still there?".
Received #3: "Still there?"
"Still there?"
ok
```



If you want to see which processes are registered, you can use the `regs()` shell command.

If you've worked in other programming languages and learned the gospel of “no global variables,” you may be wondering why Erlang permits a systemwide list of processes like this. Most of the rest of this book, after all, has been about isolating change and minimizing shared context.

If you think of registered processes as more like services than functions, however, it may make more sense. A registered process is effectively a service published to the entire system, something usable from multiple contexts. Used sparingly, registered processes create reliable entry points for your programs, something that can be very valuable as your code grows in size and complexity.

When Processes Break

Processes are fragile. If there's an error, the function stops and the process goes away. **Example 8-5**, in *ch08/ex5-division*, shows a `report/0` function that can break if it gets input that isn't a number.

Example 8-5. A fragile function

```
-module(bounce).
-export([report/0]).

report() ->
    receive
        X -> io:format("Divided to ~p~n",[X/2]),
        report()
    end.
```

If you compile and run this (deliberately) error-inviting code, you'll find that it works well so long as you only send it numbers. Send anything else, and you'll see an ERROR REPORT in the shell, and no more responses from that pid. It died.

```
1> c(bounce).
{ok,bounce}
2> Pid3=spawn(bounce,report,[]).
<0.38.0>
3> Pid3 ! 38.
Divided to 19.0
38
ok4> Pid3 ! 27.56.
Divided to 13.78
27.56
ok5> Pid3 ! seven.

=ERROR REPORT==== 24-Aug-2012::20:59:43 ===
Error in process <0.38.0> with exit value: {badarith,[{bounce,report,0,[{file,
"bounce.erl"},{line,6}]}]}

seven
6> Pid3 ! 14.
```

14

As you get deeper into Erlang’s process model, you’ll find that “let it crash” is not an unusual design decision in Erlang, though being able to tolerate such things and continue requires some extra work. [Chapter 9](#) will also show you how to find and deal with errors of various kinds.

Processes Talking Amongst Themselves

Sending messages to Erlang processes is easy, but it’s hard for them to report back responses if you don’t leave information about where they can find you again. Sending a message without including the sender’s pid is kind of like leaving a phone message without including your own number: it might trigger action, but the recipient might not get back to you.

To establish process to process communications without registering lots of processes, you need to include pids in the messages. Passing the pid requires adding an argument to the message. It’s easy to get started with a test that calls back the shell. [Example 8-6](#), in *ch08/ex6-talking*, builds on the `drop` module from [Example 3-2](#), adding a `drop/0` function that receives messages and removing the `fall_velocity/2` function from the `export`.

Example 8-6. A process that sends a message back to the process that called it

```
-module(drop).
-export([drop/0]).  
  
drop() ->
receive
    {From, Planemo, Distance} ->
        From ! {Planemo, Distance, fall_velocity(Planemo, Distance)},
        drop()
end.  
  
fall_velocity(earth, Distance) when Distance >= 0  -> math:sqrt(2 * 9.8 * Distance);
fall_velocity(moon, Distance) when Distance >= 0  -> math:sqrt(2 * 1.6 * Distance);
fall_velocity(mars, Distance) when Distance >= 0  -> math:sqrt(2 * 3.71 * Distance).
```

To get started, it’s easy to test this from the shell:

```
1> c(drop).
{ok,drop}
2> Pid1=spawn(drop,drop,[]).
<0.38.0>
3> Pid1 ! {self(), moon, 20}.
{<0.31.0>,moon,20}
4> flush().
Shell got {moon,20,8.0}
ok
```

Example 8-7, which you'll find in *ch08/ex7-talkingProcs*, shows a process that calls that process, just to demonstrate that this can work with more than just the shell.

Example 8-7. Calling a process from a process, and reporting the results

```
-module(mph_drop).  
-export([mph_drop/0]).  
  
mph_drop() ->  
    Drop=spawn(drop:drop,[]),  
    convert(Drop).  
  
convert(Drop) ->  
    receive  
        {Planemo, Distance} ->  
            Drop ! {self(), Planemo, Distance},  
            convert(Drop);  
        {Planemo, Distance, Velocity} ->  
            MphVelocity= 2.23693629 * Velocity,  
            io:format("On ~p, a fall of ~p meters yields a velocity of ~p mph.\n",  
[Planemo, Distance, MphVelocity]),  
            convert(Drop)  
    end.
```

The `mph_drop/1` function spawns a `drop:drop/0` process when it is first set up, using the same module you saw in [Example 8-6](#), and stores the pid in `Drop`. Then it calls `convert/1`, which will listen for messages recursively.



If you don't separate the initialization from the recursive listener, your code will work, but will spawn new `drop:drop/0` processes every time it processes a message instead of using the same one repeatedly.

The `receive` clause relies on the call from the shell (or another process) including only two arguments, while the `drop:drop/0` process sends back a result with three. (As your code grows more complex, you will likely want to use more explicit flags about the kind of information contained in a message.) When the `receive` clause gets a message with two arguments, it sends a message to `Drop`, identifying itself as the sender and passing on the arguments. When the `Drop` returns a message with the result, the `receive` clause reports on the result, converting the velocity to miles per hour. (Yes, it leaves the distance metric, but makes the velocity more intelligible to Americans.)

Using this from the shell looks like the following:

```
1> c(drop).  
{ok,drop}  
2> c(mph_drop).  
{ok,mph_drop}
```

```

3> Pid1=spawn(mph_drop,mph_drop,[]).
<0.43.0>
4> Pid1 ! {earth,20}.
On earth, a fall of 20 meters yields a velocity of 44.289078952755766 mph.
{earth,20}
5> Pid1 ! {mars,20}.
On mars, a fall of 20 meters yields a velocity of 27.250254686571544 mph.
{mars,20}

```

This simple example might look like it behaves as a more complex version of a function call, but there is a critical difference. In the shell, with nothing else running, the result will come back quickly—so quickly that it reports before the shell puts up the message—but this was a series of asynchronous calls. Nothing held and waited specifically for a returned message.

The shell sent a message to `Pid1`, the process identifier for `mph_drop:convert/1`. That process sent a message to `Drop`, the process identifier for `drop:drop/0`, which `mph_drop:mph_drop:0` set up when it was spawned. That process returned another message to `mph_drop:convert/1`, which reported to standard output, in this case the shell. Those messages passed and were processed rapidly, but in a system with thousands or millions of messages in motion, those passages might have been separated by many messages, and come in later.

Watching Your Processes

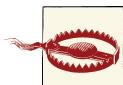
Erlang provides a simple but powerful tool for keeping track of your processes and seeing what's happening. Pman, the process manager, offers a minimal GUI that lets you look into the current state of your processes and see what's happening. Depending on how you installed Erlang, you may be able to start it from a toolbar, but you can always start it from the shell:

```

6> pman:start().
<0.48.0>

```

You'll see something like [Figure 8-1](#) appear. It's a long list of processes, more than you probably wanted to know about. To see a saner, simpler version, check "Hide System Processes," and you'll see something similar to [Figure 8-2](#).



The documentation for Pman warns that "The Pman application has been superseded by the Observer application. Pman will be removed in R16." Things will be changing soon. However, Observer in R15 doesn't run on 64-bit Macs (because of dependencies on wxWidgets). This chapter will get an update when R16 releases.

Pman: Overview on nonode@nohost

File	View	Trace	Nodes	Help	
Pid	Current Function	Name	Msgs	Reds	Size
<0.53.0>	io:wait_io_mon_reply/2		0	14	239
<0.52.0>	gstk:worker_init/1		0	1	235
<0.51.0>	gstk_port_handler:idle/1		0	6695	6767
<0.50.0>	gstk:loop/1		0	45296	4183
<0.49.0>	gs_frontend:loop/1	gs_frontend	0	1608	991
<0.48.0>	pman_process:pinfo/2		0	13088	1675
<0.44.0>	drop:drop/0		0	9	237
<0.43.0>	mph_drop:convert/1		0	37	235
<0.31.0>	shell:eval_loop/3		0	10855	2591
<0.27.0>	gen_server:loop/6	kernel_safe_sup	0	58	242
<0.26.0>	gen_server:loop/6		0	268	242
<0.25.0>	shell:get_command1/5		0	5113	2600
<0.24.0>	group:get_line1/4		0	2508	1009

Figure 8-1. The Process Manager at startup

Figure 8-2. The Process Manager after hiding (most) system processes

The list of processes is useful, but Pman also lets you look inside of process activity. If you double-click on a process, say `mph_drop:convert/1`, you'll get some basic information about the process, as shown in [Figure 8-3](#).

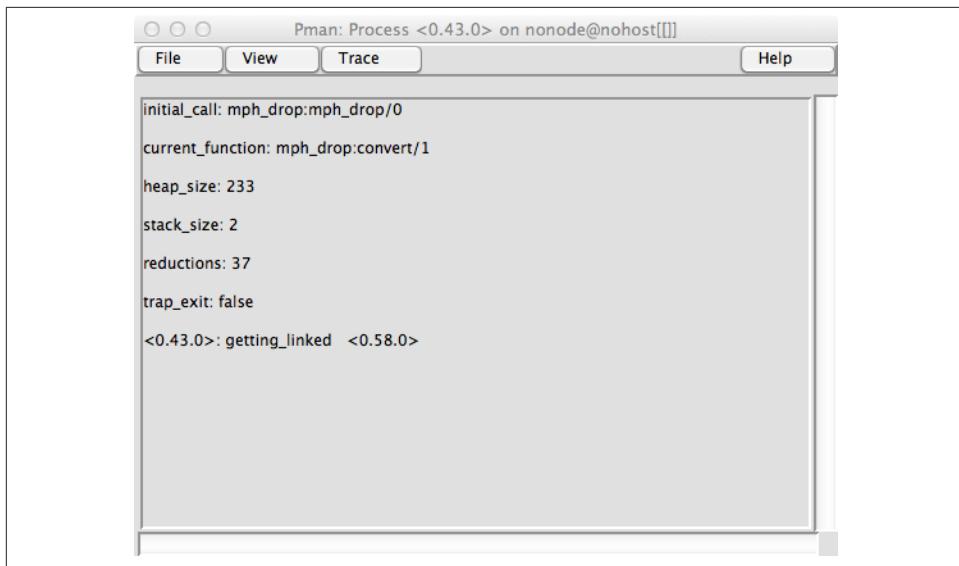


Figure 8-3. A closer look at mph_drop

Once you've opened the window, make the process do something:

```
7> Pid1 ! {mars,20}.
On mars, a fall of 20 meters yields a velocity of 27.250254686571544 mph.
{mars,20}
```

The Pman window for that process will update to show messages and calls, as shown in [Figure 8-4](#). rec means a message was received, whereas ! indicates a message sent and call indicates a function call. (io:format makes a lot of calls, which is the bottom part of that message.)

Pman is generally the easiest place to turn when you're having difficulty figuring out what is happening among your processes.

Breaking Things and Linking Processes

When you send a message, you'll always get back the message as the return value. This doesn't mean that everything went well and the message was received and processed correctly, however. If you send a message that doesn't match a pattern at the receiving process, nothing will happen (for now at least), with the message landing in the mailbox but not triggering activity. Sending a message that gets through the pattern matching but creates an error will halt the process where the error occurred, possibly even a few messages and processes down the line.

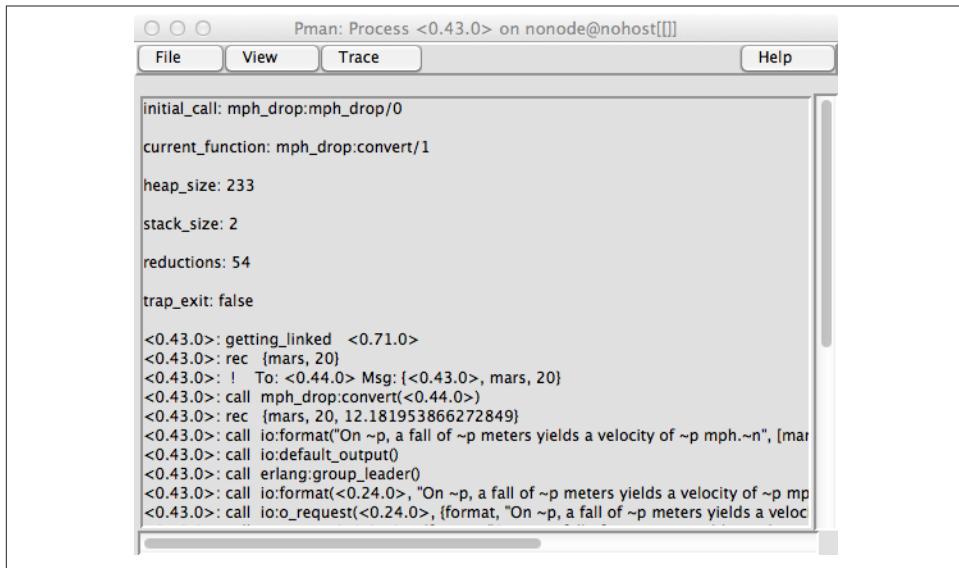


Figure 8-4. Tracing calls when you send mph_drop a message



Messages that don't match a pattern in the `receive` clause don't vanish; they just linger in the mailbox without being processed. It is possible to update a process with a new version of the code that retrieves those messages.

Because processes are fragile, you often want your code to know when another process has failed. In this case, if bad inputs halt the `drop:drop/0`, it doesn't make much sense to leave the `mph_drop:convert/1` process hanging around. You can see how this works through the shell and Pman. First, start up Pman and spawn `mph_drop:mph_drop/0`.

```
1> pman:start().  
<0.33.0>  
2> Pid1=spawn(mph_drop,mph_drop,[]).  
<0.40.0>
```

You'll see something like Figure 8-5 in Pman. Then, feed your process some bad data, an atom (`zoids`) instead of a number for the `Distance`, and Pman will look more like Figure 8-6.

```
3> Pid1 ! {moon,zoids}.  
{moon,zoids}  
4>
```

```
=ERROR REPORT===== 29-Sep-2012::13:41:18 ===
Error in process <0.41.0> with exit value: {badarith,[{drop,fall_velocity,2},
[{file,"drop.erl"},{line,12}]],{drop,drop,0,[{file,"drop.erl"},{line,7}]}}}
```

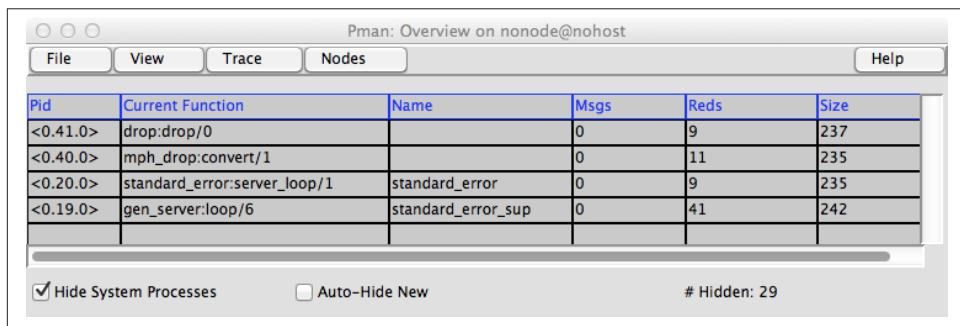


Figure 8-5. A healthy set of processes

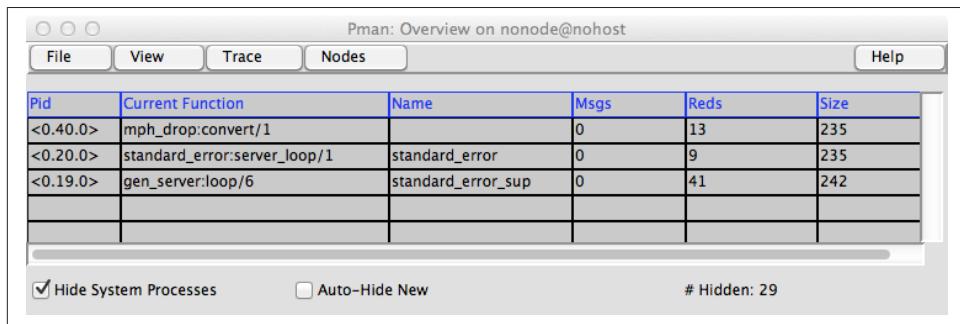


Figure 8-6. Only the drop:drop/0 process is gone

Since the remaining `mph_drop:convert/1` process is now useless, it would be better for it to halt when `drop:drop/0` fails. Erlang lets you specify that dependency with a link. The easy way to do that while avoiding potential race conditions is to use `spawn_link/3` instead of just `spawn/3`. This is shown in [Example 8-8](#), which you can find in *ch08/ex8-linking*.

[Example 8-8.](#) Calling a linked process from a process so failures propagate

```
-module(mph_drop).
-export([mph_drop/0]).

mph_drop() ->
    Drop=spawn_link(drop,drop,[]),
    convert(Drop).
```

```

convert(Drop) ->
receive
    {Planemo, Distance} ->
        Drop ! {self(), Planemo, Distance},
        convert(Drop);
    {Planemo, Distance, Velocity} ->
        MphVelocity= 2.23693629 * Velocity,
        io:format("On ~p, a fall of ~p meters yields a velocity of ~p mph.\n",[Plane
mo, Distance, MphVelocity]),
        convert(Drop)
end.

```

Now, if you recompile and test this out with Pman, you'll see that both processes vanish when `drop:drop/0` fails, as shown in [Figure 8-7](#).

```

1> c(mph_drop).
{ok,mph_drop}
2> pman:start().
<0.38.0>
3> Pid1=spawn(mph_drop,mph_drop,[ ]).
<0.45.0>
4> Pid1 ! {moon,zoids}.

=ERROR REPORT==== 29-Sep-2012::13:38:10 ===
Error in process <0.46.0> with exit value: {badarith,[{drop,fall_velocity,2,
[{file,"drop.erl"},{line,12}]},{drop,drop,0,[{file,"drop.erl"},{line,7}]}]}

{moon,zoids}

```

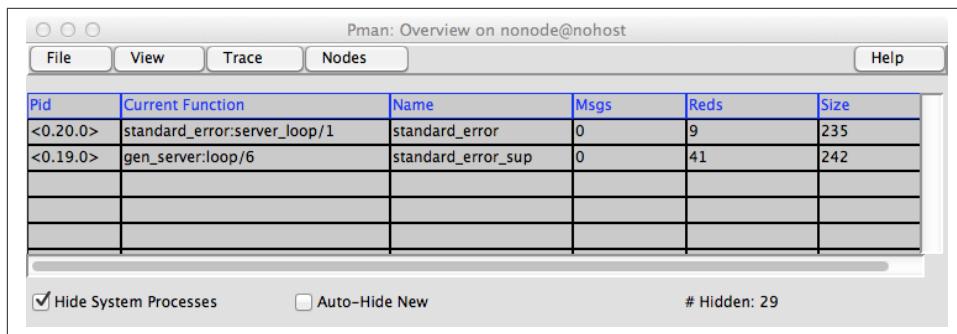


Figure 8-7. Both processes now depart when there is an error



Links are bidirectional. If you kill the `mph_drop:mph_drop/0` process—with, for example, `exit(Pid1,kill)`.—the `drop:drop/1` process will also vanish. (`kill` is the harshest reason for an exit, and isn't trappable because sometimes you really need to halt a process.)

That kind of failure may not be what you have in mind when you think of linking processes. It's the default behavior for linked Erlang processes, and makes sense in many contexts, but you can also have a process trap exits. When an Erlang process fails, it sends an explanation to other processes that are linked to it in the form of a tuple. The tuple contains the atom EXIT, the Pid of the failed process, and the error as a complex tuple. If your process is set to trap exits, through a call to `process_flag(trap_exit, true)`, these errors reports arrive as messages, rather than just killing your process.

Example 8-9, in *ch08/ex9-trapping*, shows how the initial `mph_drop/0` method changes to include this call to set the process flag, and adds another entry to the `receive` clause which will listen for exits and report them more neatly.

Example 8-9. Trapping a failure, reporting an error, and exiting

```
-module(mph_drop).
-export([mph_drop/0]).

mph_drop() ->
    process_flag(trap_exit, true),
    Drop=spawn_link(drop,drop,[]),
    convert(Drop).

convert(Drop) ->
    receive
        {Planemo, Distance} ->
            Drop ! {self(), Planemo, Distance},
            convert(Drop);
        {'EXIT', Pid, Reason} ->
            io:format("FAILURE: ~p died because of ~p.~n",[Pid, Reason]);
        {Planemo, Distance, Velocity} ->
            MphVelocity= 2.23693629 * Velocity,
            io:format("On ~p, a fall of ~p meters yields a velocity of ~p mph.~n",[Plane
mo, Distance, MphVelocity]),
            convert(Drop)
    end.
```

If you run this and feed it bad data, the `convert/1` method will report an error message (mostly duplicating the shell) before exiting neatly.

```
1> c(mph_drop).
{ok,mph_drop}
2> Pid1=spawn(mph_drop,mph_drop,[[]]).
<0.45.0>
3> Pid1 ! {moon,20}.
On moon, a fall of 20 meters yields a velocity of 17.89549032 mph.
{moon,20}
4> Pid1 ! {moon,zoids}.
FAILURE: <0.46.0> died because of {badarith,
                                         [{drop,fall_velocity,2,
                                           [{file,"drop.erl"},{line,12}]}],
                                         {drop,drop,0},
```

```

[[{file,"drop.erl"},{line,7}]]]}.

=ERROR REPORT==== 29-Sep-2012::14:37:47 ===
Error in process <0.46.0> with exit value: {badarith,[{drop,fall_velocity,2,
[{file,"drop.erl"},{line,12}]}},{drop,drop,0,[{file,"drop.erl"},{line,7}]}]}

{moon,zoids}

```

A more robust alternative would set up a new `Drop` variable, spawning a new process. That version, shown in [Example 8-10](#), which you can find at `ch08/ex10-resilient`, is much tougher. Its `receive` clause sweeps away failure, soldiering on with a new copy (New `Drop`) of the drop calculator if needed.

Example 8-10. Trapping a failure, reporting an error, and setting up a new process

```

-module(mph_drop).
-export([mph_drop/0]).

mph_drop() ->
    process_flag(trap_exit, true),
    Drop=spawn_link(drop,drop,[]),
    convert(Drop).

convert(Drop) ->
    receive
        {Planemo, Distance} ->
            Drop ! {self(), Planemo, Distance},
            convert(Drop);
        {'EXIT', _Pid, _Reason} ->
            NewDrop=spawn_link(drop,drop,[]),
            convert(NewDrop);
        {Planemo, Distance, Velocity} ->
            MphVelocity= 2.3293629 * Velocity,
            io:format("On ~p, a fall of ~p meters yields a velocity of ~p mph.\n",[Plane
mo, Distance, MphVelocity]),
            convert(Drop)
    end.

```

If you compile and run [Example 8-10](#), you'll see [Figure 8-8](#) when you first start Pman (and check the Hide System Processes box). If you feed it bad data, as shown on line 6 in the following code sample, you'll still get the error message from the shell, but the process will work just fine. As you'll see in Pman, as shown in [Figure 8-9](#), it started up a new process to handle the `drop:drop/0` calculations, and as line 8 shows, it works like its predecessor.

```

1> c(mph_drop).
{ok,mph_drop}
2> Pid1=spawn(mph_drop,mph_drop,[]).
<0.43.0>
3> pman:start().
<0.46.0>

```

```

4> Pid1 ! {moon,20}.
On moon, a fall of 20 meters yields a velocity of 17.89549032 mph.
{moon,20}
5> Pid1 ! {mars,20}.
On mars, a fall of 20 meters yields a velocity of 27.250254686571544 mph.
{mars,20}
6> Pid1 ! {mars,zoids}.
{mars,zoids}
7>
=ERROR REPORT==== 29-Sep-2012::14:53:13 ===
Error in process <0.44.0> with exit value: {badarith,[{drop,fall_velocity,2,
[{file,"drop.erl"},{line,13}]],{drop,drop,0,[{file,"drop.erl"},{line,7}]}]}}

7> Pid1 ! {moon,20}.
On moon, a fall of 20 meters yields a velocity of 17.89549032 mph.
{moon,20}

```

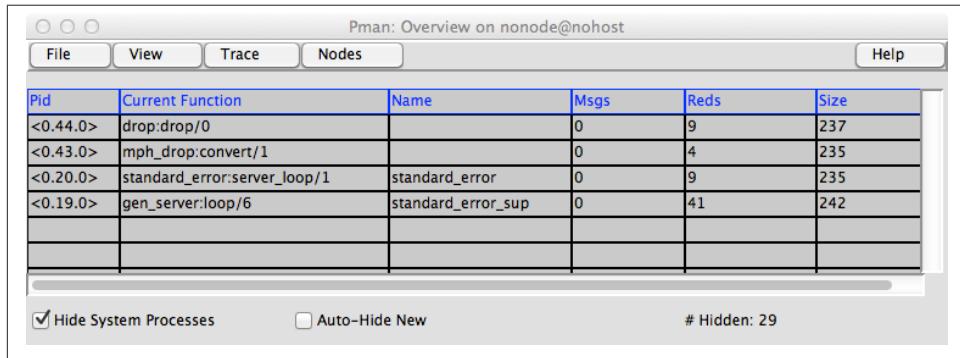


Figure 8-8. Processes before an error. Note the Pid on the top line.

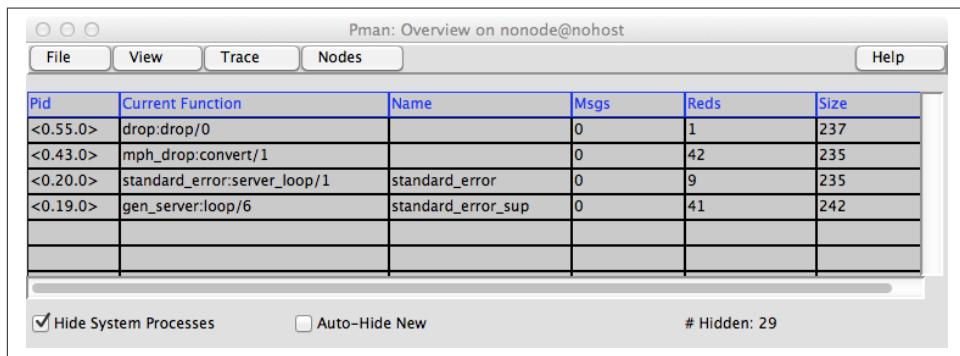


Figure 8-9. Processes after an error. Note the top line Pid change.

Erlang offers many more process management options. You can remove a link with `unlink/1`, or establish a connection for just watching a process with `erlang:monitor/2`. If you want to terminate a process, you can use `exit/1` within that process, or `exit/2` to specify a process and reason from another process.

Building applications that can tolerate failure and restore their functionality is at the core of robust Erlang programming. Developing in that style is probably a larger leap for most programmers than Erlang's shift to functional programming, but it's where the true power of Erlang becomes obvious.

Exceptions, Errors, and Debugging

“Let it crash” is a brilliant insight, but one whose application you probably want to control. While it’s possible to write code that constantly breaks and recovers, it can be easier to write and maintain code that explicitly handles failure where it happens. However you choose to deal with errors, you’ll definitely want to be able to track them down in your application.

Flavors of Errors

As you’ve already seen, some kinds of errors will keep Erlang from compiling your code, and the compiler will also give you warnings about potential issues, like variables that are declared but never used. Two other kinds of errors are common: runtime errors, which turn up when code is operating and can actually halt a function or process, and logic errors, which may not kill your program but can cause deeper headaches.

Logic errors are often the trickiest to diagnose, requiring careful thought and perhaps some time with the debugger, log files, or a test suite. Simple mathematical errors can take a lot of work to untangle. Sometimes issues are related to timing, when the sequence of operations isn’t what you expect. In severe cases, race conditions can create deadlocks and halting, but more mild cases can produce bad results and confusion.

Runtime errors can also be annoying, but they are much more manageable. In some ways you can see handling runtime errors as part of the logic of your program, though you don’t want to get carried away. In Erlang, unlike many other programs, handling errors as errors may offer only minor advantages over letting an error kill a process and then dealing with the problem at the process level, as [Example 8-10](#) showed.

Catching Runtime Errors as They Happen

If you want to catch runtime errors close to where they took place, the `try...catch` construct lets you wrap suspect code and handle problems (if any) that code creates. It makes it clear to both the compiler and the programmer that something unusual is happening, and lets you deal with any unfortunate consequences of that work.

For a simple example, look back to [Example 3-1](#), which calculated fall velocity without considering the possibility that it would be handed a negative distance. The `math:sqrt/1` function will produce a `badarith` error if it has a negative argument. [Example 4-2](#) kept that problem from occurring by applying guards, but if you want to do more than block, you can take a more direct approach with `try` and `catch`, as shown in [Example 9-1](#). (You can find it in `ch09/ex1-tryCatch`.)

Example 9-1. Using try and catch to handle a possible error

```
-module(drop).
-export([fall_velocity/2]).  
  
fall_velocity(Planemo, Distance) ->
    Gravity = case Planemo of
        earth -> 9.8;
        moon -> 1.6;
        mars -> 3.71
    end,  
  
try math:sqrt(2 * Gravity * Distance) of
    Result -> Result
  catch
    error:Error -> {error, Error}
  end.
```

The calculation itself is now wrapped in a `try`. If the calculation succeeds, the pattern match following the `of` will be used. In this case, the calculation just produces one value, so matching the variable `Result` will put that value in `Result`, which then becomes the return value.

You can leave out the `of` clause entirely if this is all you're doing; the result of the expression in the `try` will become the returned value. You probably won't see `of` very frequently in code. This `try...catch` construct produces exactly the same results as the one in [Example 9-1](#).

```
try math:sqrt(2 * Gravity * Distance)
  catch
    error:Error -> {error, Error}
  end.
```

If the calculation fails, in this case because of a negative argument, the pattern match in the `catch` clause comes into play. In this case, the atom `error` will match the class or exception type of the error (which can be `error`, `throw`, or `exit`), and the variable `Error` will collect the details of the error. It then returns a tuple, opening with the atom `error` and the contents of the `Error` variable, which will explain the type of error.

You can try the following on the command line:

```
1> c(drop).
{ok,drop}
2> drop:fall_velocity(earth,20).
19.79898987322333
3> drop:fall_velocity(earth,-20).
{error,badarith}
```

When the calculation is successful, you'll just get the result. When it fails, the tuple tells you the kind of error that caused the problem. It's not a complete solution, but it's a foundation on which you can build.

You can have multiple statements in the `try` (much as you can in a `case`), separated by commas as usual. At least when you're getting started, it's easiest to keep the code you are trying simple so you can see where failures happened. However, if you wanted to watch for requests that provided an atom that didn't match the planemos in the `case`, you could put it all into the `try`:

```
fall_velocity(Planemo, Distance) ->
try
    Gravity = case Planemo of
        earth -> 9.8;
        moon -> 1.6;
        mars -> 3.71
    end,
    math:sqrt(2 * Gravity * Distance)
of
    Result -> Result
catch
    error:Error -> {error, Error}
end.
```

If you try an unsupported planemo, you'll now see the code catch the problem, at least once you recompile the code to use the new version:

```
4> drop:fall_velocity(jupiter,20).
** exception error: no case clause matching jupiter
   in function  drop:fall_velocity/2 (drop.erl, line 5)
5> c(drop).
{ok,drop}
6> drop:fall_velocity(jupiter,20).
{error,{case_clause,jupiter}}
```

The `case_clause` error indicates that a `case` failed to match, and the second component of that tuple, `jupiter`, tells you the item that didn't match.

You can also have multiple pattern matches in the `catch`. If your patterns don't match the error in the `catch` clause, it gets reported as a runtime error, as if the `try` hadn't wrapped it.

If the code that might fail can create a mess, you may want to include an `after` clause after the `catch` clause and before the closing `end`. The code in an `after` clause is guaranteed to run whether the attempted code succeeds or fails, and can be a good place to address any side effects of the code. It doesn't affect the return value of the clause.



Erlang also includes an older `catch` construct that doesn't use `try`. You may find this in someone else's code, but probably shouldn't include it in any new code that you write. It is less sophisticated and less readable, though [Chapter 11](#) shows one use for it in the shell.

Raising Exceptions with `throw`

You may want to create your own errors, or at least report results in a way that the `try...` `catch` mechanism can work with. The `throw/1` function lets you create exceptions that can then be caught (or left to kill a process and possibly reported to the shell). It often takes a tuple as an argument, letting you provide more detail about the exception, but you can use whatever you think appropriate. You definitely want to handle exceptions close to where you want to raise them, however.

Using it in the shell provides a simple example of what it does:

```
1> throw(my_exception).
** exception throw: my_exception
```

You can pattern match for thrown exceptions in a `catch` clause by using `throw` instead of `error`:

```
try some:function(argument)
  catch
    error:Error -> {found, Error};
    throw:Exception -> {caught, Exception}
  end;
```

You probably should save `throw` for cases where you can't come up with a better approach for signaling within your code, and be sure to use it only where you know you have nearby code that will catch it. Relying on other people and distant code to understand your invented exceptions may stretch their patience.



The preceding example used `found` and `caught` to distinguish between the different kinds of exceptions, but most code will likely just use `error` for both.

Logging Progress and Failure

The `io:format/2` function is useful for simple communications with the shell, but as your programs grow (and especially as they become distributed processes), hurling text toward standard output is less likely to get you the information you need. Erlang offers a set of functions for more formal logging. They *can* hook into more sophisticated logging systems, but it's easy to get started with them as a way to structure messages from your application.

Three functions in the `error_logger` module give you three levels of reporting:

`info_msg`

For logging ordinary news that doesn't require intervention.

`warning_msg`

For news that's worse. Someone should do something eventually.

`error_msg`

Something just plain broke, and needs to be looked at.

Like `io:format`, there are two versions of each function. The simpler version takes just a string, while the more sophisticated version takes a string and a list of arguments that get added to that string. Both use the same formatting structure as `io:format`, so you can pretty much replace any `io:format` calls you'd been using for debugging directly. All of these return `ok`.

As you can see, these calls produce reports that are visually distinctive, though warnings and errors get the same ERROR REPORT treatment:

```
1> error_logger:info_msg("The value is ~p. ~n",[360]).  
ok  
  
=INFO REPORT==== 12-Oct-2012::08:00:41 ===  
The value is 360.  
2> error_logger:warning_msg("Connection lost; will retry.").  
  
=ERROR REPORT==== 12-Oct-2012::08:01:33 ===  
Connection lost; will retry.  
ok  
Connection lost; will retry.ok  
3> error_logger:error_msg("Unable to read database.~n").
```

```
=ERROR REPORT==== 12-Oct-2012::08:03:45 ===  
Unable to read database.  
ok
```

The more verbose form produces only a mild improvement over `io:format`, so why would you use it? Because Erlang has much much more lurking under the surface. By default, when Erlang starts up, it sets up the `error_logger` module to report to the shell. However, if you turn on SASL—the Erlang System Architecture Support Libraries, not the authentication layer—you’ll be able to connect these notices to a much more sophisticated system for logging distributed processes. (If you just want to write your errors to disk, you should explore the `error_logger:logfile/1` function.)



It's possible to break the logger with bad format strings, so if you want more reliable logging, you may want to check into the more spartan `_report` versions of these functions.

While logging information is useful, it's not unusual to write code with subtle errors where you're not positive what to log where. You could litter the code with reporting, or you could switch to a different set of tools, Erlang's debugging facilities.

Debugging through a GUI

Erlang's graphical debugger is the friendliest place to start, requiring only a minor change in how you compile code to get started. This demonstration will use the same code shown in [Example 9-1](#), but you need to compile it with the `debug_info` flag, and start the debugger with `debugger:start()`. You'll see a window like the one shown in [Figure 9-1](#).

```
1> c(drop, [debug_info]).  
{ok,drop}  
2> debugger:start().  
  
=ERROR REPORT==== 16-Oct-2012::05:30:03 ===  
WX Failed loading "wxe_driver"@"/usr/local/lib/erlang/lib/wx-0.99.2/priv"  
{ok,<0.40.0>}
```



The ERROR REPORT is because of issues with my 64-bit Mac, and hopefully they will go away in a future release. The debugger is somewhat unstable here but mostly still works, and you can take that report as a good sign that the `error_logger` module is used throughout Erlang.

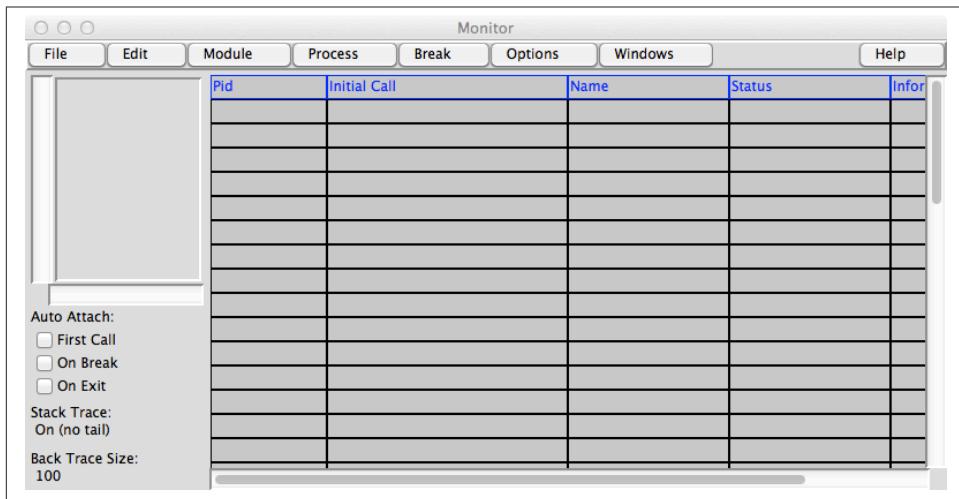


Figure 9-1. The debugger window when first opened

When it first opens, the debugger window is pretty empty looking. You need to tell it what you want to watch, by choosing Interpret... from the Module menu (which looks like a button in the top row). As shown in [Figure 9-2](#), you should see the drop module (you may need to navigate to it if you didn't start in the same directory). If you select the drop module and click Choose, drop will appear in the lefthand pane of the Monitor window, as shown in [Figure 9-3](#). You can then click Done to close the window and get back to the Monitor.

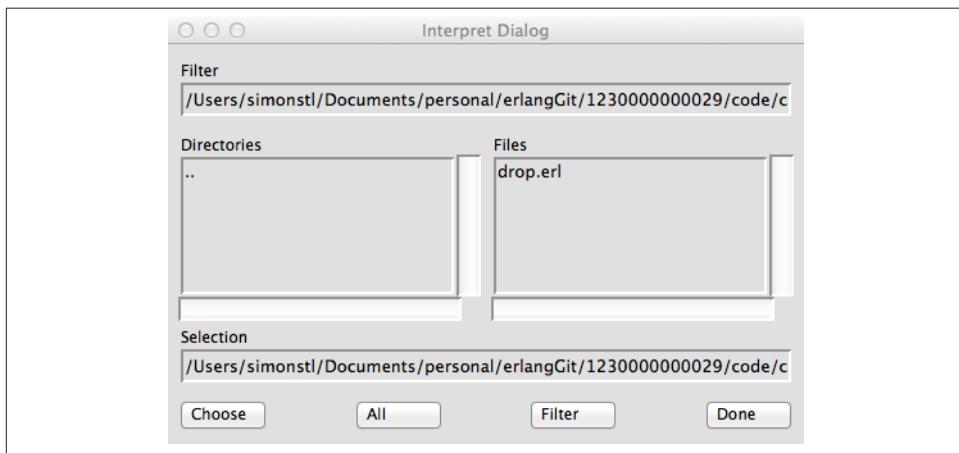


Figure 9-2. Choosing a module

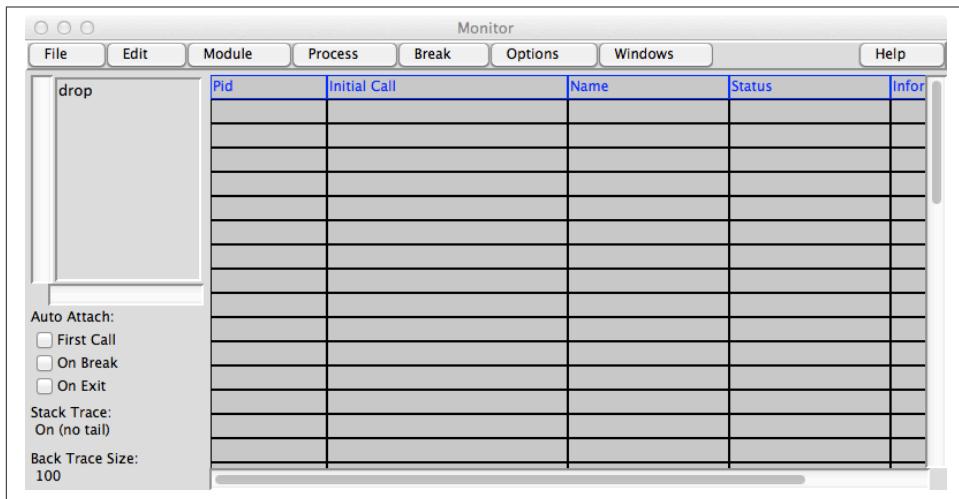
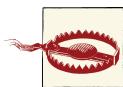


Figure 9-3. A module name appears on the lefthand side



The only feedback you'll get that you actually selected the module is its appearance in the Monitor pane. If your windows are stacked so that you can't see that happen, it's easy to think that nothing happened. It did, really!

Once the module name appears in the lefthand side of the Monitor window, the debugger is ready to watch it. You need to tell the debugger, however, what you want to see. If you double-click on the name of the module (`drop`), you'll get the View Module `drop` window shown in [Figure 9-4](#), showing its code.

You can add a breakpoint by clicking on a line of code, and then choosing Line Break... from the Break menu. You'll see the Line Break dialog shown in [Figure 9-5](#), with reasonable default settings. You can just click OK, and the View Module `drop` window will change to indicate the breakpoint, as shown in [Figure 9-6](#). You can close this window, and just leave the Monitor window open.

Now, if you go back to the shell and request:

```
3> drop:fall_velocity(earth,20).
```

You'll just get a pause. Nothing seems to happen, as the breakpoint stopped execution. However, in the Monitor window, you'll see a new entry in the table on the righthand side, as shown in [Figure 9-7](#). If you double-click that new entry, you'll get to the Attach Process window in [Figure 9-8](#), which lets you step through the code line by line.

Once you have the Attach Process window up, you can work through your code line by line (or tell it to continue) using the buttons in the middle line:

The screenshot shows a software window titled "View Module drop". The menu bar includes "File", "Edit", "Break", "Windows", and "Help". The main area displays the following Prolog code:

```
1: -module(drop).
2: -export([fall_velocity/2]).
3:
4: fall_velocity(Planet, Distance) ->
5:
6: Gravity = case Planet of
7:   earth -> 9.8;
8:   moon -> 1.6;
9:   mars -> 3.71
10: end, % note comma - function isn't done yet
11:
12: try math:sqrt(2 * Gravity * Distance) of
13:   Result -> Result
14: catch
15:   error:Error -> {found, Error}
16: end.
```

Figure 9-4. Examining the code for the drop module

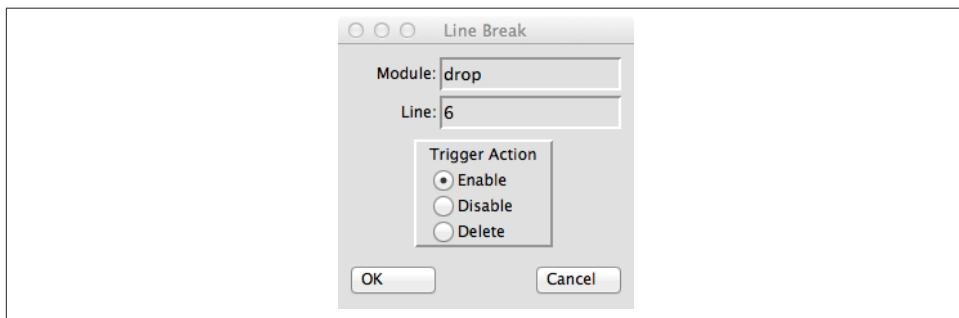


Figure 9-5. The Line Break dialog for setting breakpoints

Step

Execute the current line of code and move into the next line. If the next line of code to be executed is in another function (and that function is in a module compiled for debugging), you'll step through that function's code.

Next

Execute the current line of code and move to the next line of code in *this* module.

Continue

End the line-by-line stepping and just have the code execute as usual.

```

1: -module(drop).
2: -export([fall_velocity/2]).
3:
4: fall_velocity(Planet, Distance) ->
5:
6:     Gravity = case Planet of
7:         earth -> 9.8;
8:         moon -> 1.6;
9:         mars -> 3.71
10:    end, % note comma - function isn't done yet
11:
12:    try math:sqrt(2 * Gravity * Distance) of
13:        Result -> Result
14:    catch
15:        error:Error -> {found, Error}
16:    end.

```

Figure 9-6. The drop module with a breakpoint set on line 6

Monitor																																												
File	Edit	Module	Process	Break																																								
<table border="1"> <thead> <tr> <th>Pid</th><th>Initial Call</th><th>Name</th><th>Status</th><th>Infor</th></tr> </thead> <tbody> <tr> <td><0.31.0></td><td>drop:fall_velocity/2</td><td></td><td>break</td><td>{drop</td></tr> <tr><td></td><td></td><td></td><td></td><td></td></tr> <tr><td></td><td></td><td></td><td></td><td></td></tr> <tr><td></td><td></td><td></td><td></td><td></td></tr> <tr><td></td><td></td><td></td><td></td><td></td></tr> <tr><td></td><td></td><td></td><td></td><td></td></tr> <tr><td></td><td></td><td></td><td></td><td></td></tr> </tbody> </table>					Pid	Initial Call	Name	Status	Infor	<0.31.0>	drop:fall_velocity/2		break	{drop																														
Pid	Initial Call	Name	Status	Infor																																								
<0.31.0>	drop:fall_velocity/2		break	{drop																																								

Figure 9-7. The Monitor window shows activity

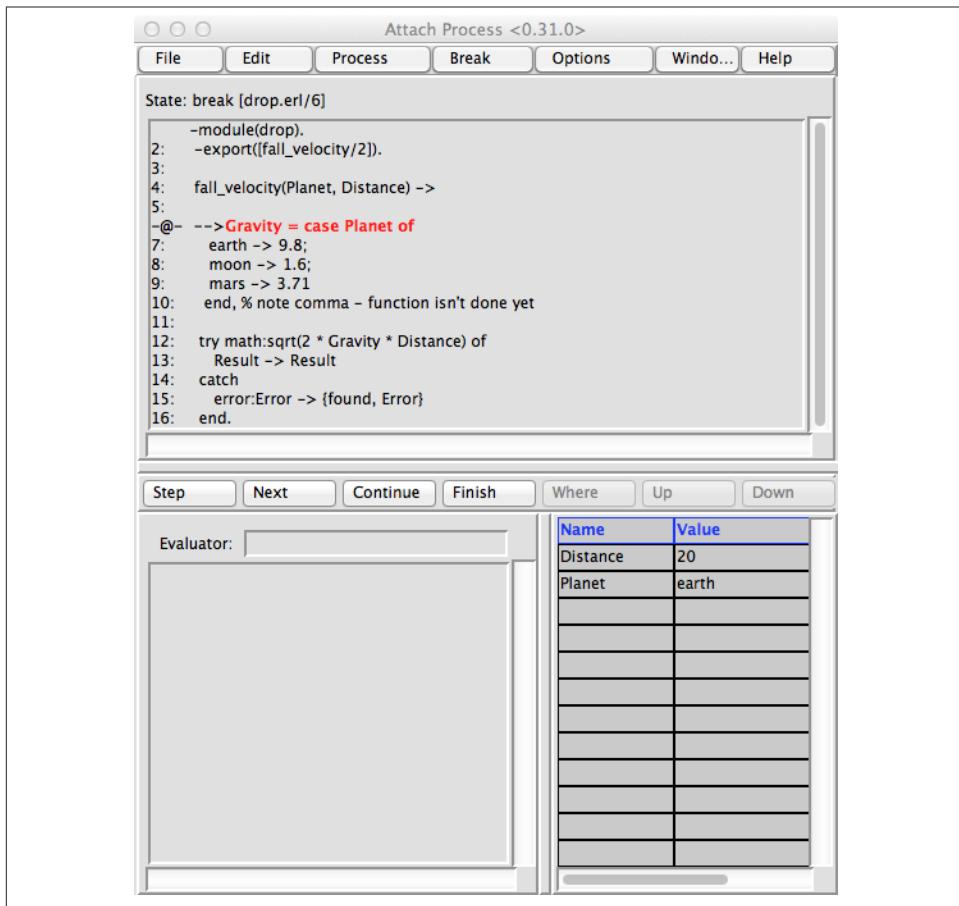


Figure 9-8. Code and bound values in the Attach Process window

Finish

Similar to continue, but continues only for the current function. The debugger can keep working on the code when it returns from this function. (This is useful when you've stepped into a function whose details don't interest you and you don't have the patience to wait.)

Where

Moves the code window to the currently executing line.

Up and Down

Moves the code window up or down a function level in the stack.

Figure 9-9 through **Figure 9-12** show the results of stepping through the code executed by the `drop:fall_velocity(earth,20)` call. Note the changing bound variables and the final return to State: uninterpreted in **Figure 9-12** when the call completes.

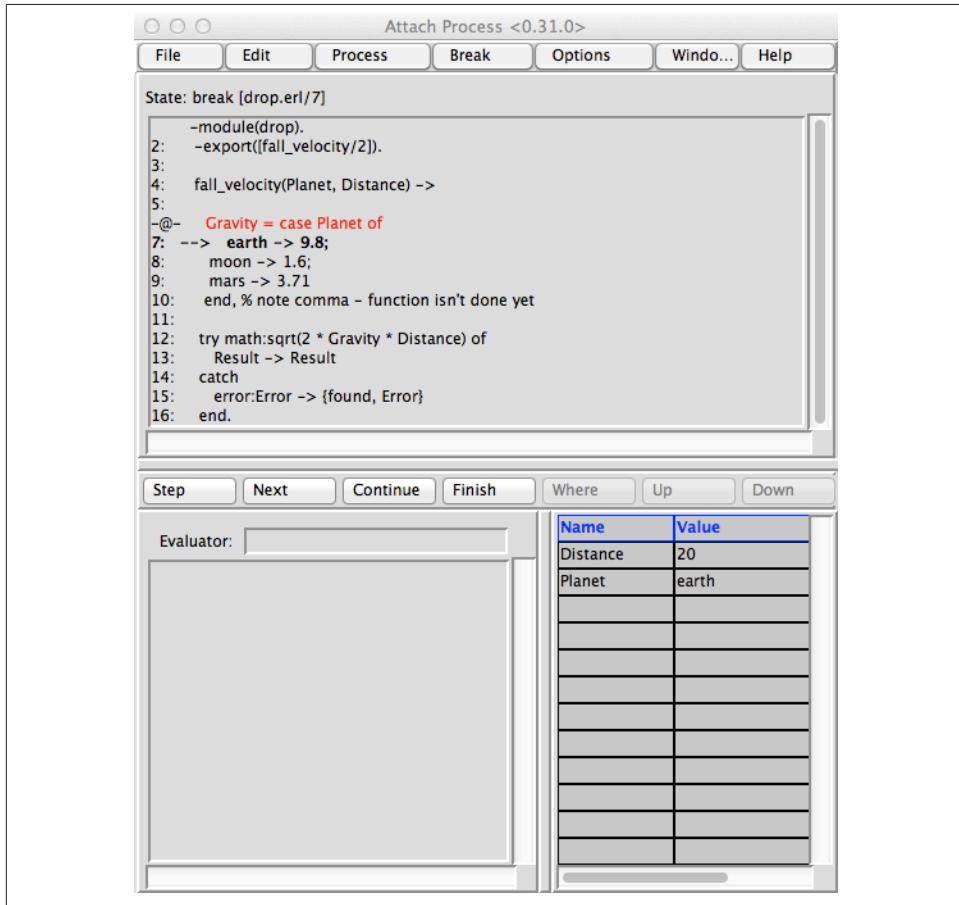


Figure 9-9. Stepping to the match on line 7

You can use the Evaluator pane to make your own calculations using the values and functions available in the current scope at any time the code is paused. Unlike some debuggers, you can't change the value of the bound variables here—because you can't change the values of variables in Erlang generally.

The debugger offers many more features, but this core set will get you started.

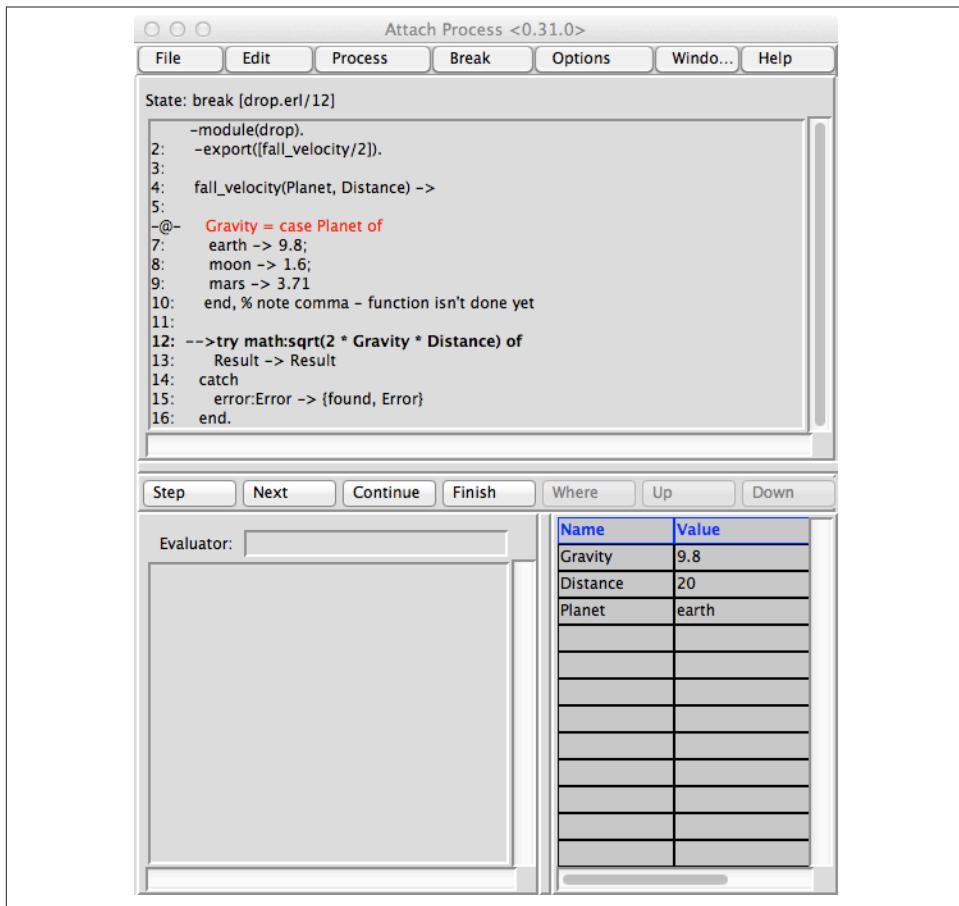


Figure 9-10. Next step: the try statement and calculation on line 12, with additional bound values

Tracing Messages

Erlang also offers a wide variety of tools for tracing code, both with other code (with the `trace` and `trace_pattern` built-in functions) and with a text-based debugger/reporter. The `dbg` module is the easiest place to start into this toolset, letting you specify what you want traced and showing you the results in the shell.

An easy place to get started is tracing messages sent between processes. You can use `dbg:p` to trace the messages sent between the `mph_drop` process defined in [Example 8-8](#) and the `drop` process from [Example 8-6](#). After compiling the modules—the `debug_info`

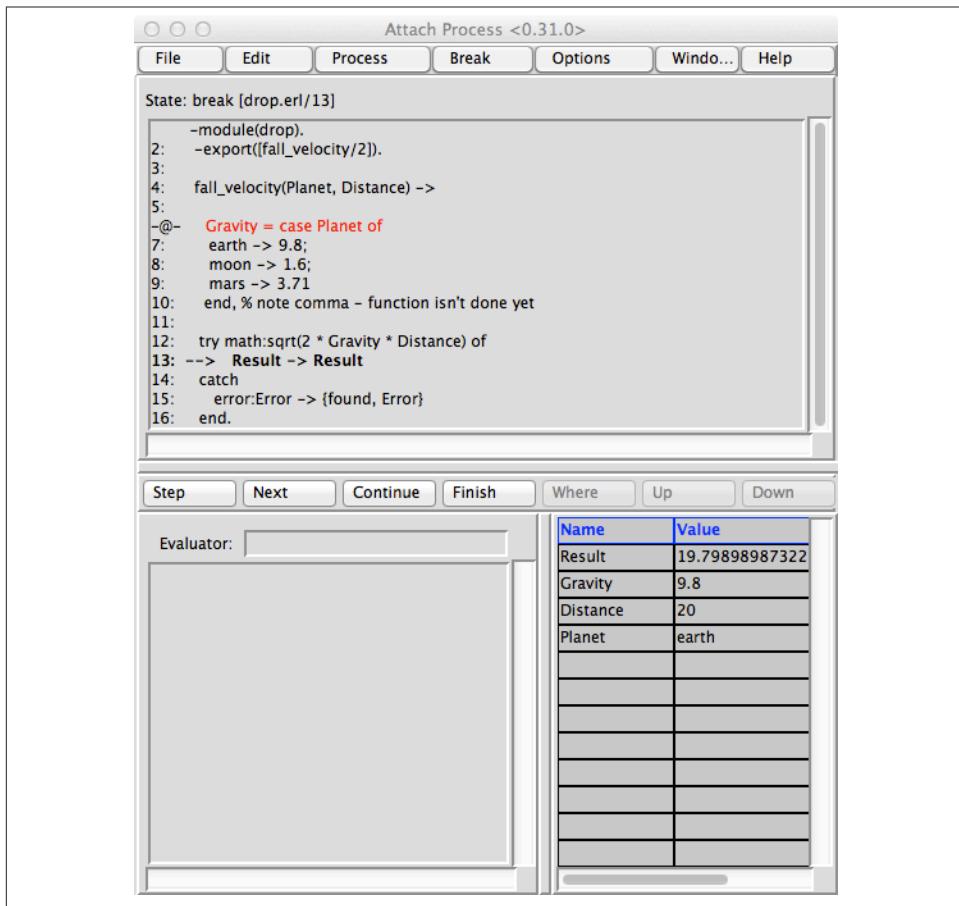


Figure 9-11. A successful calculation leads to line 13, which provides the return value

flag isn't needed here—you call `dbg:tracer()` to start reporting trace information to the shell. Then you `spawn` the `mph_drop` process as usual, and pass that pid to the `dbg:p/2` process. The second argument here will be `m`, meaning that the trace should report the messages.

```
1> c(drop).
{ok,drop}
2> c(mph_drop).
{ok,mph_drop}
3> dbg:tracer().
{ok,<0.43.0>}
```

```

4> Pid1=spawn(mph_drop,mph_drop,[]).
<0.46.0>
5> dbg:p(Pid1,m).
{ok,[{matched,nonode@nohost,1}]}

```

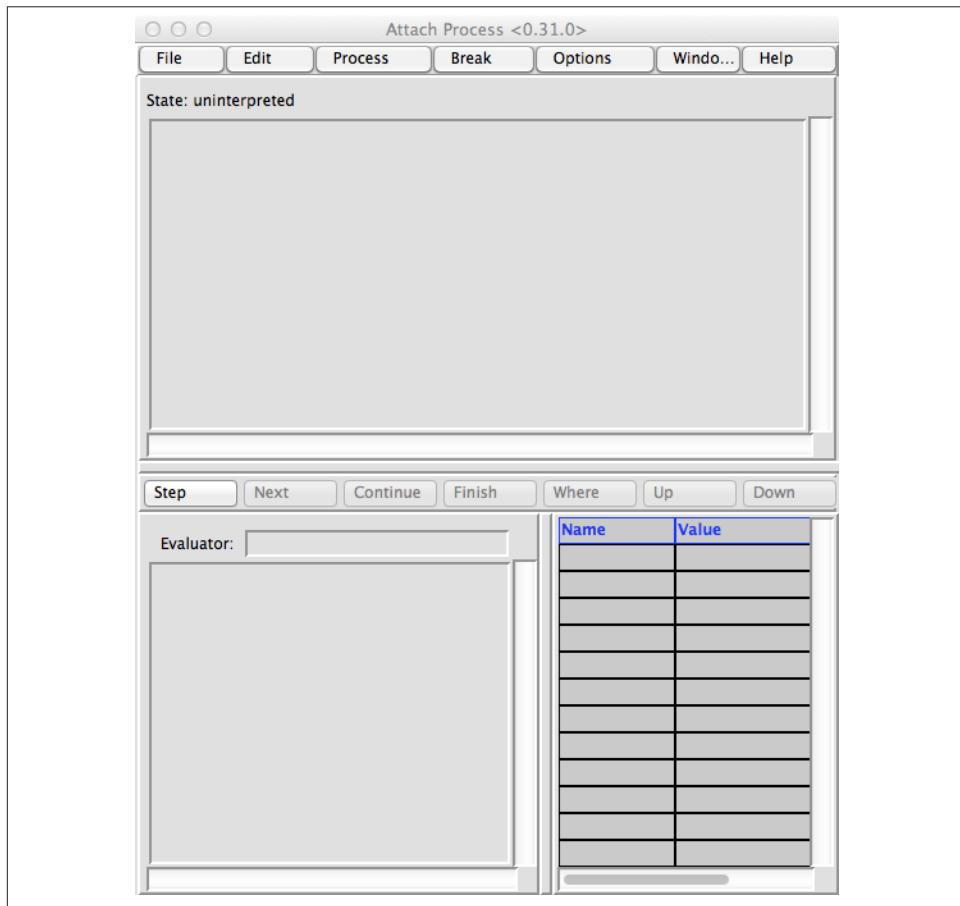


Figure 9-12. The function call complete, the window goes mostly blank

Now when you send a message to the `mph_drop` process, you'll get a set of reports on the resulting flow of messages. (<0.46.0> is the `mph_drop` process, and <0.47.0> is the `drop` process.)

```

6> Pid1 ! {moon,20}.
(<0.46.0>) << {moon,20}
(<0.46.0>) <0.47.0> ! {<0.46.0>,moon,20}
On moon, a fall of 20 meters yields a velocity of 17.89549032 mph.
(<0.46.0>) << {moon,20,8.0}
(<0.46.0>) <0.24.0> ! {io_request,<0.46.0>,<0.24.0>,

```

```

    {put_chars,unicode,io_lib,format,
     ["On ~p, a fall of ~p meters yields a
      velocity of ~p mph.\n",
      [moon,20,17.89549032]]}}}
{moon,20}
(<0.46.0>) << {io_reply,<0.24.0>,ok}
(<0.46.0>) << timeout

```

The `<<` pointing to a pid indicates that that process received a message. Sends are indicated, as usual, with the pid followed by `!` followed by the message. In this case:

- `mph_drop (<0.46.0>)` receives the message tuple `{moon,20}`.
- It sends a further message, the tuple `{<0.46.0>,moon,20}`, to the `drop` process at pid `<0.47.0>`.
- On this run, the report from `mph_drop` that “On moon, a fall of 20 meters yields a velocity of 17.89549032 mph.” comes through faster than the tracing information. The rest of the trace indicates how that report got there.
- `mph_drop` receives a tuple `{moon,20,8.0}` (from `drop`).
- Then it calls `io:format/2`, which triggers another set of process messages to make the report, concluding with a `timeout` that doesn’t do anything.

The trace reports come through a bit after the actual execution of the code, but they make the flow of messages clear. You’ll want to learn to use `dbg` in its many variations to trace your code, and may eventually want to use match patterns and the `trace` functions themselves to create more elegant systems for watching specific code.

Watching Function Calls

If you just want to keep track of arguments moving between function calls, you can use the tracer to report on the sequence of calls. [Chapter 4](#) demonstrated recursion and reported results along the way through `io:format`. There’s another way to see that work, again using the `dbg` module.

[Example 4-11](#), the upward factorial calculator, started with a call to `fact:factorial/1`, which then called `fact:factorial/3` recursively. `dbg` will let you see the actual function calls and their arguments, mixed with the `io:format` reporting. (You can find it in `ch09/ex4-dbg`.)

Tracing functions is a little trickier than tracing messages because you can’t just pass `dbg:p/2` a pid. As shown on line 3 in the following code sample, you need to tell it you want it to report on all processes (`all`), and their calls (`c`). Once you’ve done that, you

have to specify which calls you want it to report, using `dbg:tpl` as shown on line 4. It takes a module name (`fact`), function name (`factorial`), and optionally a match specification that lets you specify arguments more precisely. Variations on this function also let you specify arity.

So turn on the tracer, tell it you want to follow function calls, and specify a function (or functions, through multiple calls to `dbg:tpl`) to watch. Then call the function, and you'll see a list of the calls.

```
1> c(fact).
fact.erl:13: Warning: variable 'Current' is unused
fact.erl:13: Warning: variable 'N' is unused
{ok,fact}
2> dbg:tracer().
{ok,<0.38.0>}
3> dbg:p(all, c).
{ok,[{matched,nonode@nohost,26}]}
4> dbg:tpl(fact, factorial, []).
{ok,[{matched,nonode@nohost,2}]}
5> fact:factorial(4).
1 yields 1!
(<0.31.0>) call fact:factorial(4)
(<0.31.0>) call fact:factorial(1,4,1)
2 yields 2!
(<0.31.0>) call fact:factorial(2,4,1)
3 yields 6!
(<0.31.0>) call fact:factorial(3,4,2)
4 yields 24!
(<0.31.0>) call fact:factorial(4,4,6)
Finished.
(<0.31.0>) call fact:factorial(5,4,24)
24
```

You can see that the sequence is a bit messy here, with the trace reporting coming a little bit after the `io:format` results from the function being traced. Because the trace is running in a separate process (at pid `<0.38.0>`) from the function (at pid `<0.31.0>`), its reporting may not line up smoothly (or at all, though it usually does).

When you're done tracing, call `dbg:stop/0` (if you might want to restart tracing with the same setup) or `dbg:stop_clear/0` (if you know that when you start again you'll want to set things up again).

The `dbg` module and the `erlang:trace` functions on which it builds are incredibly powerful tools.

Storing Structured Data

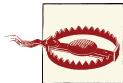
Tuples and lists are powerful tools for creating complex data structures, but there are two key pieces missing from the story so far. First, tuples are relatively anonymous structures. Relying on a specific order and number of components in tuples can create major maintenance headaches. Second, despite Erlang's general preference for avoiding side effects, storing and sharing data is a fundamental side effect needed for a wide variety of projects.

Records will help you create labeled sets of information. Erlang Term Storage (ETS) will help you store and manipulate those sets, and the Mnesia database provides additional features for reliable distributed storage.

From Tuples to Records

Tuples let you build complex data structures, but force you to rely on keeping the order and number of items consistent. If you change the sequence of items in a tuple, or if you want to add an item, you have to check through all of your code to make sure that the change propagates smoothly. As your projects grow, and especially if you need to share data structures with code you don't control, you'll need a safer way to store and address information.

Records let you create data structures that use names to connect with data rather than order. You can read, write, and pattern match data in a record without having to worry about the details of where in a tuple a field lurks or whether someone's added a new field.



There are still tuples underneath records, and occasionally Erlang will expose them to you. Do not attempt to use the tuple representation directly, or you will add all the potential problems of using tuples to the slight extra syntax of using records.

Setting Up Records

Using records requires telling Erlang about them with a special declaration. It looks like a `-module` or `-export` declaration, but is a `-record` declaration:

```
-record(planemo, {name, gravity, diameter, distance_from_sun}).
```

That defines a record type named `planemo`, containing fields named `name`, `gravity`, and `distance_from_sun`. Right now, when you create a new record, the fields will all have the value `undefined`, but you can also specify default values if you prefer, for situations where there is a sensible normal option. This declaration creates records for different towers for dropping objects:

```
-record(tower, {location, height=20, planemo=earth, name}).
```

Unlike `-module` or `-export` declarations, you'll often want to share record declarations across multiple modules and (for the examples in this chapter at least) even use them in the shell. To share record declarations reliably, just put the record declarations in their own file, ending with the extension `.hrl`. You may want to put each record declaration in a separate file or all of them in a single file, depending on your needs. To get started, to see how these behave, you can put both of the declarations into a single file, `records.hrl`, shown in [Example 10-1](#). (You can find it in `ch10/ex1-records`.)

Example 10-1. A records.hrl file containing two rather unrelated record declarations

```
-record(planemo, {name, gravity, diameter, distance_from_sun}).  
-record(tower, {location, height=20, planemo=earth, name}).
```



You may want to put individual record declarations into their own files and import them separately, bringing them in only when you actually need to get data into or out of a particular record type. This can be especially important if you're mixing code in cases where different developers used the same name for a record type but different underlying structures.

The command `rr` (for read records) lets you bring this into the shell:

```
1> rr("records.hrl").  
[planemo,tower]
```

The shell now understands records with the names `planemo` and `tower`.



You can also declare records directly in the shell with the `rd/2` function, but if you're doing anything more than just poking around, it's easier to have them in a more reliable formal imported declaration. You can call `rl/0` if you want to see what records are defined, or `rl/1` if you want to see how a specific record is defined.

Creating and Reading Records

You can now create variables that contain new records. The syntax for referencing records prefaces the name of the record type with a `#`, and encloses name-value pairs in curly brackets. For example, you could create towers with syntax like the following:

```
2> Tower1=#tower{}.
#tower{location = undefined,height = 20,planemo = earth,
      name = undefined}
3> Tower2=#tower{location="Grand Canyon"}.
#tower{location = "Grand Canyon",height = 20,
      planemo = earth,name = undefined}
4> Tower3=#tower{location="NYC", height=241, name="Woolworth Building"}.
#tower{location = "NYC",height = 241,planemo = earth,
      name = "Woolworth Building"}
5> Tower4=#tower{location="Rupes Altai 241", height=500, planemo=moon, name="Pic
   colomini View"}.
#tower{location = "Rupes Altai 241",height = 500,
      planemo = moon,name = "Piccolomini View"}
6> Tower5=#tower{planemo=mars, height=500, name="Daga Vallis", location="Valles
   Marineris"}.
#tower{location = "Valles Marineris",height = 500,
      planemo = mars,name = "Daga Vallis"}
```

These towers (or at least drop sites) demonstrate a variety of ways to use the record syntax to create variables as well as interactions with the default values:

- Line 2 just creates `Tower1` with the default values. You can add real values later.
- Line 3 creates a `Tower2` with a `location`, but otherwise relies on the default values.
- Line 4 overrides the default values for `location`, `height`, and `name`, but leaves the `planemo` alone.
- Line 5 replaces all of the default values with new values.
- Line 6 replaces all of the default values, and also demonstrates that *it doesn't matter in what order you list the name/value pairs*. Erlang will sort it out.

You can read record entries with two different approaches. To extract a single value, you can use a dot (.) syntax that may look familiar from other languages. For example, to find out how which `planemo` `Tower5` is on, you could write:

```
7> Tower5#tower.planemo.  
mars
```

You could also use pattern matching to extract several pieces simultaneously:

```
8> #tower{location=L5, height=H5} = Tower5.  
#tower{location = "Valles Marineris",height = 500,  
      planemo = mars,name = "Daga Vallis"}  
9> L5.  
"Valles Marineris"  
10> H5.  
500
```

The syntax feels a little backward, with the variable being bound on the right side of the equals sign instead of in its usual place on the left.

As always, you can't write a new value to an existing variable, but you can create a new record based on the values of an old one. The syntax used on line 13 is much like that used for assigning the contents of a field to a variable, but with a value in place of the variable name:

```
12> Tower5.  
#tower{location = "Valles Marineris",height = 500,  
      planemo = mars,name = "Daga Vallis"}  
13> Tower5a=Tower5#tower{height=512}.  
#tower{location = "Valles Marineris",height = 512,  
      planemo = mars,name = "Daga Vallis"}
```



Yes, you always need to specify the record type. It's a bit of extra typing.

If you ever want to make the shell forget your record declarations, you can issue the shell command `rf()`. Your record-based variables will still exist, in a raw tuple form you should avoid ever using.

Using Records in Functions and Modules

Records work well in modules as well, using the same declaration files. You *can*, of course, just include the record declaration in every module that uses it, but that will require you to hunt down every declaration and update it if you ever want to change it. The saner approach is to use the files like the ones previously shown. You can do that easily with a single extra declaration near the top of your module:

```
-include("records.hrl").
```

Once you have the record declaration included, you can pattern match against records submitted as arguments. The simplest way to do this is to just match against the type of the record, as shown in [Example 10-2](#), which is also in *ch10/ex1-records*.

Example 10-2. A method that pattern matches a complete record

```
-module(record_drop).
-export([fall_velocity/1]).
-include("records.hrl").

fall_velocity(#tower{} = T) ->
    fall_velocity(T#tower.planemo, T#tower.height).

fall_velocity(earth, Distance) when Distance >= 0 -> math:sqrt(2 * 9.8 * Distance);
fall_velocity(moon, Distance) when Distance >= 0 -> math:sqrt(2 * 1.6 * Distance);
fall_velocity(mars, Distance) when Distance >= 0 -> math:sqrt(2 * 3.71 * Distance).
```

This uses a pattern match that will match only `tower` records, and puts the record into a variable `T`. Once again, the syntax may seem backward, with `T` being on the right of the equals instead of on the left, but it works. Then, like its predecessor in [Example 3-8](#), it passes the individual arguments to `fall_velocity/2` for calculations, this time using the record syntax.



Short variable names suddenly seem more attractive when you have to append the name of the record type on every use. In simple functions this can work, but in more complex functions short names may prove confusing, especially if you have two variables containing the same kind of record.

Because you used the same `-record` declaration in both the shell and the module, you can use the records you created to test the function.

```
14> c(record_drop).
{ok,record_drop}
15> record_drop:fall_velocity(Tower5).
60.909769331364245
16> record_drop:fall_velocity(Tower1).
19.798987322333
```

The `record_drop:fall_velocity/1` function shown in [Example 10-3](#) pulls out the `planemo` and binds it to `Planemo`, and `height` and binds it to `Distance`. Then it returns the velocity of an object dropped from that `Distance` just like earlier examples throughout this book.

You can also extract the specific fields from the record in the pattern match, as shown in [Example 10-3](#), which is in *ch10/ex2-records*.

Example 10-3. A method that pattern matches components of a record

```
-module(record_drop).
-export([fall_velocity/1]).
-include("records.hrl").

fall_velocity(#tower{planemo=Planemo, height=Distance}) ->
    fall_velocity(Planemo, Distance).

fall_velocity(earth, Distance) when Distance >= 0 -> math:sqrt(2 * 9.8 * Distance);
fall_velocity(moon, Distance) when Distance >= 0 -> math:sqrt(2 * 1.6 * Distance);
fall_velocity(mars, Distance) when Distance >= 0 -> math:sqrt(2 * 3.71 * Distance).
```

Again, the syntax may seem backwards, but it lets you extract the individual fields. You can take the records created and feed them into this function, and it will tell you the velocity resulting from a drop from the top of that tower to the bottom.

Finally, you can pattern match against both the fields and the records as a whole. **Example 10-4**, in *ch10/ex3-records*, demonstrates using this mixed approach to create a more detailed response than just the fall velocity.

Example 10-4. A method that pattern matches the whole record as well as components of a record

```
-module(record_drop).
-export([fall_velocity/1]).
-include("records.hrl").

fall_velocity(#tower{planemo=Planemo, height=Distance} = T) ->
    io:format("From ~s's elevation of ~p meters on ~p, the object will reach ~p m/s
before crashing in ~s.~n",[T#tower.name, Distance, Planemo, fall_velocity(Planemo,
Distance), T#tower.location ]).

fall_velocity(earth, Distance) when Distance >= 0 -> math:sqrt(2 * 9.8 * Distance);
fall_velocity(moon, Distance) when Distance >= 0 -> math:sqrt(2 * 1.6 * Distance);
fall_velocity(mars, Distance) when Distance >= 0 -> math:sqrt(2 * 3.71 * Distance).
```

If you pass a `tower` record to `record_drop:fall_velocity/1`, it will match against individual fields it needs to do the calculation, and match the whole record into `T` so that it can produce a more interesting if not necessarily grammatically correct report.

```
17> record_drop:fall_velocity(Tower5).
From Daga Vallis's elevation of 500 meters on mars, the object will reach
60.909769331364245 m/s before crashing in Valles Marineris.
ok
18> record_drop:fall_velocity(Tower3).
From Woolworth Building's elevation of 241 meters on earth, the object will
reach 68.72845116834803 m/s before crashing in NYC.
ok
```



`record_drop:fall_velocity/1` uses the `~s` control sequence for the `io:format/2` call. It just includes the contents of the string, without surrounding quotes.

Storing Records in Erlang Term Storage

Erlang Term Storage (ETS) is a simple but powerful in-memory collection store. It holds tuples, and since records are tuples underneath, they're a natural fit. ETS and its disk-based cousin DETS provide a (perhaps too) simple solution for many data management problems. ETS is not exactly a database, but does similar work, and is useful by itself as well as underneath the Mnesia database you'll see in the next section.

Every entry in an ETS tables is a tuple (or corresponding record), and one piece of the tuple is designated the key. ETS offers a few different structural choices depending on how you want to handle that key. ETS can hold four kinds of collections:

Sets (set)

Can contain only one entry with a given key. This is the default.

Ordered sets (ordered_set)

Same as a set, but also maintains a traversal order based on the keys. Great for anything you want to keep in alphabetic or numeric order.

Bags (bag)

Lets you store more than one entry with a given key. However, if you have multiple entries that have completely identical values, they get combined into a single entry.

Duplicate bags (duplicate_bag)

Not only lets you store more than one entry with a given key, but also lets you store multiple entries with completely identical values.

By default, ETS tables are sets, but you can specify one of the other options when you create a table. The examples here will be sets because they are simpler to figure out, but the same techniques apply to all four table varieties.



There is no requirement in ETS that all of your entries look at all similar. When you're starting out, however, it's much simpler to use the same kind of record, or at least tuples with the same structure. You can also use any kind of value for the key, including complex tuple structures and lists, but again, it's best not to get too fancy at the beginning.

All of the examples in the following section will use the `planemo` record type defined in the previous section, and the data in [Table 10-1](#).

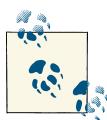
Table 10-1. Planemos for gravitational exploration

Planemo	Gravity (m/s ²)	Diameter (km)	Distance from Sun (10 ⁶ km)
mercury	3.7	4878	57.9
venus	8.9	12104	108.2
earth	9.8	12756	149.6
moon	1.6	3475	149.6
mars	3.7	6787	227.9
ceres	0.27	950	413.7
jupiter	23.1	142796	778.3
saturn	9.0	120660	1427.0
uranus	8.7	51118	2871.0
neptune	11.0	30200	4497.1
pluto	0.6	2300	5913.0
haumea	0.44	1150	6484.0
makemake	0.5	1500	6850.0
eris	0.8	2400	10210.0

Creating and Populating a Table

The `ets:new/2` function lets you create a table. The first argument is a name for the table, and the second argument is a list of options. There are lots and lots of options, including the identifiers for the table types described above, but the two most important for getting started are `named_table` and the tuple starting with `keypos`.

Every table has a name, but only some can be reached using that name. If you don't specify `named_table`, the name is there but visible only inside the database. You'll have to use the value returned by `ets:new/2` to reference the table. If you do specify `named_table`, processes can reach the table as long as they know the name, without needing access to that return value.



Even with a named table, you still have some control over which processes can read and write the table through the `private`, `protected`, and `public` options.

The other important option, especially for ETS tables containing records, is the `keypos` tuple. By default, ETS treats the first value in a tuple as the key. The tuple representation underneath records (which you shouldn't really touch) always uses the first value in a tuple to identify the kind of record, so that approach works very badly as a key for records. Using the `keypos` tuple lets you specify which record value should be the key.

Remember, the record format for a `planemo` looks like the following:

```
-record(planemo, {name, gravity, diameter, distance_from_sun}).
```

Because this table is mostly used for calculations based on a given `planemo`, it makes sense to use the `name` as a key. An appropriate declaration for setting up the ETS table might look like the following:

```
PlanemoTable=ets:new(planemos,[ named_table, {keypos, #planemo.name} ])
```

That gives the table the name `planemos` and uses the `named_table` option to make that table visible to other processes that know the name. Because of the default access level of `protected`, this process can write to that table but other processes can only read it. It also tells ETS to use the `name` field as the key. Because it doesn't specify otherwise, the table will be treated as a `set`—each key maps to only one instance of record, and ETS doesn't keep the list sorted by key.

Once you have the table set up, as shown in [Example 10-5](#), you use the `ets:info/1` function to check out its details. (You can find this in `ch10/ex4-ets`.)

Example 10-5. Setting up a simple ETS table and reporting on what's there

```
-module(planemo_storage).
-export([setup/0]).
-include("records.hrl").

setup() ->
    PlanemoTable=ets:new(planemos, [named_table, {keypos, #planemo.name}]),
    ets:info(PlanemoTable).
```

If you compile and run this, you'll get a report of an empty ETS table with more properties than you probably want to know about at the moment.

```
1> c(planemo_storage).
{ok,planemo_storage}
2> planemo_storage:setup().
[{compressed,false},
 {memory,317},
 {owner,<0.316.0>},
 {heir,none},
 {name,planemos},
 {size,0},
 {node,nonode@nohost},
 {named_table,true},
 {type,set},
 {keypos,2},
 {protection,protected}]
```

Most of this is either more information than you need or unsurprising, but it is good to see the name(planemos), size (0—empty!), and keypos (not 1, the default, but 2, the location of the name in the tuple underneath the record). It is, as the defaults specify, set up as a **protected set**.

You can set up only one ETS table with the same name. If you call `planemo_storage:setup/0` twice, you'll get an error:

```
3> planemo_storage:setup().
** exception error: bad argument
   in function  ets:new/2
      called as ets:new(planemos,[named_table,{keypos,2}])
   in call from planemo_storage:setup/0 (planemo_storage.erl, line 6)
```

To avoid this, at least in these early tests, you'll want to use the `f()` shell command to clear out any previous tables. If you think you're likely to call your initialization code repeatedly after you figure the basics out, you can also test the `ets:info/1` for undefined to make sure the table doesn't already exist, or put a `try...catch` construct around the `ets:new/2` call.

A more exciting ETS table, of course, will include content. The next step is to use `ets:insert/2` to add content to the table. The first argument is the table, referenced either by its name (if you set the `named_table` option), or by the variable that captured the return value of `ets:new/2`. In [Example 10-6](#), which is in `ch10/ex5-ets`, the first call uses the name, to show that it works, and the rest use the variable. The second argument is a record representing one of the rows from [Table 10-1](#).

Example 10-6. Populating a simple ETS table and reporting on what's there

```
-module(planemo_storage).
-export([setup/0]).
-include("records.hrl").

setup() ->
PlanemoTable=ets:new(planemos, [named_table, {keypos, #planemo.name}]),

ets:insert(planemos,
#planemo{ name=mercury, gravity=3.7, diameter=4878, distance_from_sun=57.9 }),
ets:insert(PlanemoTable,
#planemo{ name=venus, gravity=8.9, diameter=12104, distance_from_sun=108.2 }),
ets:insert(PlanemoTable,
#planemo{ name=earth, gravity=9.8, diameter=12756, distance_from_sun=149.6 }),
ets:insert(PlanemoTable,
#planemo{ name=moon, gravity=1.6, diameter=3475, distance_from_sun=149.6 }),
ets:insert(PlanemoTable,
#planemo{ name=mars, gravity=3.7, diameter=6787, distance_from_sun=227.9 }),
ets:insert(PlanemoTable,
#planemo{ name=ceres, gravity=0.27, diameter=950, distance_from_sun=413.7 }),
ets:insert(PlanemoTable,
#planemo{ name=jupiter, gravity=23.1, diameter=142796, distance_from_sun=778.3 }),
```

```

ets:insert(PlanemoTable,
  #planemo{ name=saturn, gravity=9.0, diameter=120660, distance_from_sun=1427.0 }),
ets:insert(PlanemoTable,
  #planemo{ name=uranus, gravity=8.7, diameter=51118, distance_from_sun=2871.0 }),
ets:insert(PlanemoTable,
  #planemo{ name=neptune, gravity=11.0, diameter=30200, distance_from_sun=4497.1 }),
ets:insert(PlanemoTable,
  #planemo{ name=pluto, gravity=0.6, diameter=2300, distance_from_sun=5913.0 }),
ets:insert(PlanemoTable,
  #planemo{ name=haumea, gravity=0.44, diameter=1150, distance_from_sun=6484.0 }),
ets:insert(PlanemoTable,
  #planemo{ name=makemake, gravity=0.5, diameter=1500, distance_from_sun=6850.0 }),
ets:insert(PlanemoTable,
  #planemo{ name=eris, gravity=0.8, diameter=2400, distance_from_sun=10210.0 }),
ets:info(PlanemoTable).

```

Again, the last call is to `ets:info/1`, which now reports that the table has 14 items.

```

4> c(planemo_storage).
{ok,planemo_storage}
5> f().
ok
6> planemo_storage:setup().
[{compressed,false},
 {memory,541},
 {owner,<0.342.0>},
 {heir,none},
 {name,planemos},
 {size,14},
 {node,nonode@nohost},
 {named_table,true},
 {type,set},
 {keypos,2},
 {protection,protected}]

```

If you want to see what's in that table, you have a couple of options. The quick way to do it in the shell is to use the `ets:tab2list/1` function, which will return a list of records (or tuples, if you leave out the record import on line 7):

```

7> rr("records.hrl").
[planemo,tower]
8> ets:tab2list(planemos).
[#planemo{name = pluto,gravity = 0.6,diameter = 2300,
          distance_from_sun = 5913.0},
 #planemo{name = saturn,gravity = 9.0,diameter = 120660,
          distance_from_sun = 1427.0},
 #planemo{name = moon,gravity = 1.6,diameter = 3475,
          distance_from_sun = 149.6},
 #planemo{name = mercury,gravity = 3.7,diameter = 4878,
          distance_from_sun = 57.9},
 #planemo{name = earth,gravity = 9.8,diameter = 12756,
          distance_from_sun = 149.6},
 #planemo{name = neptune,gravity = 11.0,diameter = 30200,
          distance_from_sun = 4497.1}]

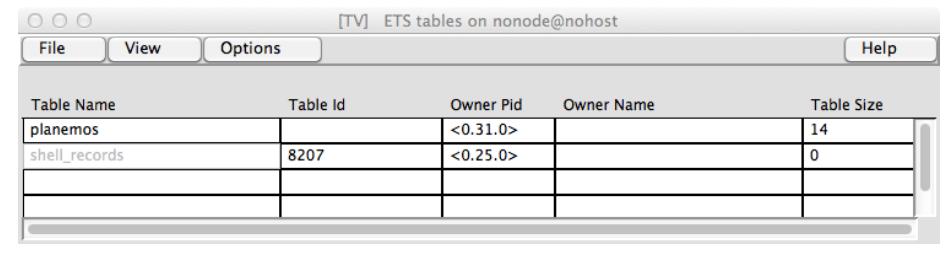
```

```

        distance_from_sun = 4497.1},
#planemo{name = makemake,gravity = 0.5,diameter = 1500,
        distance_from_sun = 6850.0},
#planemo{name = uranus,gravity = 8.7,diameter = 51118,
        distance_from_sun = 2871.0},
#planemo{name = ceres,gravity = 0.27,diameter = 950,
        distance_from_sun = 413.7},
#planemo{name = venus,gravity = 8.9,diameter = 12104,
        distance_from_sun = 108.2},
#planemo{name = mars,gravity = 3.7,diameter = 6787,
        distance_from_sun = 227.9},
#planemo{name = eris,gravity = 0.8,diameter = 2400,
        distance_from_sun = 10210.0},
#planemo{name = jupiter,gravity = 23.1,diameter = 142796,
        distance_from_sun = 778.3},
#planemo{name = haumea,gravity = 0.44,diameter = 1150,
        distance_from_sun = 6484.0}]

```

If you'd rather keep track of the table in a separate window, Erlang's table visualizer shows the same information in a slightly more readable form. You can start it from the shell with `tv:start()`, and you'll see something like [Figure 10-1](#). Double-click on the planemos table, and you'll see a more detailed report on its contents like the one shown in [Figure 10-2](#).



The screenshot shows a window titled "[TV] ETS tables on nonode@nohost". The window has a menu bar with "File", "View", "Options", and "Help". The main area is a table with the following data:

Table Name	Table Id	Owner Pid	Owner Name	Table Size
planemos		<0.31.0>		14
shell_records	8207	<0.25.0>		0

Figure 10-1. Opening the table visualizer

The visualizer doesn't know about your record declarations, but the skeleton key icon over the second column indicates that it is the key. The Options menu lets you poll the table to be sure you have its latest contents, and set a polling interval if you want it to refresh automatically. If you declare tables public, you can even edit their contents in the visualizer.



If you want to see a table of all the current ETS tables, try issuing `ets:i()` in the shell. You'll see the tables you've created (probably) near the bottom.

1	2	3	4	5	6
1	planemo	pluto	0.600000	2300	5913.00
2	planemo	saturn	9.00000	120660	1427.00
3	planemo	moon	1.60000	3475	149.600
4	planemo	mercury	3.70000	4878	57.9000
5	planemo	earth	9.80000	12756	149.600
6	planemo	neptune	11.0000	30200	4497.10
7	planemo	makemake	0.500000	1500	6850.00
8	planemo	uranus	8.70000	51118	2871.00
9	planemo	ceres	0.270000	950	413.700
10	planemo	venus	8.90000	12104	108.200
11	planemo	mars	3.70000	6787	227.900
12	planemo	eris	0.800000	2400	1.02100e+4
13	planemo	jupiter	23.1000	142796	778.300
14	planemo	haumea	0.440000	1150	6484.00
15					
16					
17					

Figure 10-2. Reviewing the planemos table in the visualizer

Simple Queries

The easiest way to look up records in your ETS table is with the `ets:lookup/2` function and the key. You can test this easily from the shell:

```
9> ets:lookup(planemos,eris).
[#planemo{name = eris,gravity = 0.8,diameter = 2400,
           distance_from_sun = 10210.0}]
```

The return value is always a list. This is true despite Erlang's knowing that this ETS table has the `set` type, so only one value can match the key, and despite there being only one value. In situations like this where you know that there will only be one returned value, the `hd/1` function, which [Example 5-5](#) showed for use with user inputs, can get you the head of a list quickly. Since there is only one item, the head is just that item.

```
10> hd(ets:lookup(planemos,eris)).
#planemo{name = eris,gravity = 0.8,diameter = 2400,
           distance_from_sun = 10210.0}
```

The square brackets are gone, which means that you can now extract, say, the gravity of a planemo:

```
11> Result=hd(ets:lookup(planemos,eris)).  
#planemo{name = eris,gravity = 0.8,diameter = 2400,  
         distance_from_sun = 10210.0}  
12> Result#planemo.gravity.  
0.8
```

A Key Feature: Overwriting Values

Up until now, you've had to work with (or around) Erlang's single-assignment paradigm. You can't overwrite the value of a variable, or change the value of an item in a list directly. However, ETS doesn't have that restriction. If you want to change the value of gravity on mercury, you can:

```
13> ets:insert(planemos, #planemo{ name=mercury,  
                                     gravity=3.9, diameter=4878, distance_from_sun=57.9 }).  
true  
14> ets:lookup(planemos, mercury).  
[#planemo{name = mercury,gravity = 3.9,diameter = 4878,  
          distance_from_sun = 57.9}]
```

Just because you *can* change values in an ETS table, however, doesn't mean that you should rewrite your code to replace immutable variables with flexible ETS table contents. Nor should you make all your tables public so that various processes can read and write whatever they like to the ETS table, making it a different form of shared memory.

Try to remember the discipline you've had to learn up until this point. Ask yourself when making changes is going to be useful, and when it might introduce tricky bugs. You probably won't have to change the gravity of Mercury, but it certainly could make sense to change a shipping address. If you have doubts, lean toward caution.

ETS Tables and Processes

Now that you can extract gravitational constants for planemos, you can expand the drop module to calculate drops in many more locations. [Example 10-7](#) combines the drop module from [Example 8-6](#) with the ETS table built in [Example 10-6](#) to create a more powerful drop calculator. (You can find this in *ch10/ex6-ets-calculator*.)

Example 10-7. Calculating drop velocities using an ETS table of planemo properties

```
-module(drop).  
-export([drop/0]).  
-include("records.hrl").  
  
drop() ->  
    setup(),  
    handle_drops().
```

```

handle_drops() ->
receive
    {From, Planemo, Distance} ->
        From ! {Planemo, Distance, fall_velocity(Planemo, Distance)},
        handle_drops()
end.

fall_velocity(Planemo, Distance) when Distance >= 0 ->
    P=hd(ets:lookup(planemos,Planemo)),
    math:sqrt(2 * P#planemo.gravity * Distance).

setup() ->
    ets:new(planemos, [named_table, {keypos, #planemo.name}]),

    ets:insert(planemos,
        #planemo{ name=mercury, gravity=3.7, diameter=4878, distance_from_sun=57.9 }),
    ets:insert(planemos,
        #planemo{ name=venus, gravity=8.9, diameter=12104, distance_from_sun=108.2 }),
    ets:insert(planemos,
        #planemo{ name=earth, gravity=9.8, diameter=12756, distance_from_sun=149.6 }),
    ets:insert(planemos,
        #planemo{ name=moon, gravity=1.6, diameter=3475, distance_from_sun=149.6 }),
    ets:insert(planemos,
        #planemo{ name=mars, gravity=3.7, diameter=6787, distance_from_sun=227.9 }),
    ets:insert(planemos,
        #planemo{ name=ceres, gravity=0.27, diameter=950, distance_from_sun=413.7 }),
    ets:insert(planemos,
        #planemo{ name=jupiter, gravity=23.1, diameter=142796, distance_from_sun=778.3 }),
    ets:insert(planemos,
        #planemo{ name=saturn, gravity=9.0, diameter=120660, distance_from_sun=1427.0 }),
    ets:insert(planemos,
        #planemo{ name=uranus, gravity=8.7, diameter=51118, distance_from_sun=2871.0 }),
    ets:insert(planemos,
        #planemo{ name=neptune, gravity=11.0, diameter=30200, distance_from_sun=4497.1 }),
    ets:insert(planemos,
        #planemo{ name=pluto, gravity=0.6, diameter=2300, distance_from_sun=5913.0 }),
    ets:insert(planemos,
        #planemo{ name=haumea, gravity=0.44, diameter=1150, distance_from_sun=6484.0 }),
    ets:insert(planemos,
        #planemo{ name=makemake, gravity=0.5, diameter=1500, distance_from_sun=6850.0 }),
    ets:insert(planemos,
        #planemo{ name=eris, gravity=0.8, diameter=2400, distance_from_sun=10210.0 }).

```

The `drop/0` function changes a little to call the initialization separately and avoid setting up the table on every call. This moves the message handling to a separate function, `handle_drop/0`. The `fall_velocity/2` function also changes, as it now looks up planemo names in the ETS table and gets their gravitational constant from that table rather than hardcoding those contents into the function. (While it would certainly be possible to pass the `PlanemoTable` variable from the previous example as an argument to the recursive message handler, it's simpler to just use it as a named table.)



If this process crashes and needs to be restarted, restarting it will trigger the `setup/0` function, which currently doesn't check to see if the ETS table exists. That could cause an error, except that ETS tables vanish when the processes that created them die. ETS offers an `heir` option and an `ets:give_away/3` function if you want to avoid that behavior, but for now it works well.

If you combine this module with the `mph_drop` module from [Example 8-7](#), you'll be able to calculate drop velocities on all of these planemos:

```
1> c(drop).
{ok,drop}
2> c(mph_drop).
{ok,mph_drop}
3> Pid1=spawn(mph_drop,mph_drop,[]).
<0.33.0>
4> Pid1 ! {earth,20}.
On earth, a fall of 20 meters yields a velocity of 44.289078952755766 mph.
{earth,20}
5> Pid1 ! {eris,20}.
On eris, a fall of 20 meters yields a velocity of 12.65402255793022 mph.
{eris,20}
6> Pid1 ! {makemake,20}.
On makemake, a fall of 20 meters yields a velocity of 10.003883211552367 mph.
{makemake,20}
```

That's a lot more variety than its `earth`, `moon`, and `mars` predecessors!

Next Steps

While many applications just need a fast key/value store, ETS tables are far more flexible than the examples so far demonstrate. You can use Erlang's match specifications and `ets:fun2ms` to create more complex queries with `ets:match` and `ets:select`. You can delete rows (and tables) with `ets:delete`. The `ets:first`, `ets:next`, and `ets:last` functions let you traverse tables recursively.

Perhaps most important, you can also explore DETS, the Disk-Based Term Storage, which offers similar features but with tables stored on disk. It's slower, with a 2GB limit, but the data doesn't vanish when the controlling process stops.

You can dig deeper into ETS and DETS, but if your needs are more complex, and especially if you need to split data across multiple nodes, you should probably explore the Mnesia database.

Storing Records in Mnesia

Mnesia is a database management system (DBMS) that comes with Erlang. It uses ETS and DETS underneath, but provides many more features than those components.

You should consider shifting from ETS (and DETS) tables to the Mnesia database if:

- You need to store and access data across a set of nodes, not just a single node.
- You don't want to have to think about whether you're going to store data in memory or on a disk or both.
- You need to be able to roll back transactions if something goes wrong.
- You'd like a more approachable syntax for finding and joining data.
- Management prefers the sound of "database" to the sound of "tables".

You may even find that you use ETS for some aspects of a project and Mnesia for others.



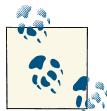
That isn't "amnesia," the forgetting, but "mnesia," the Greek word for memory.

Starting up Mnesia

If you want to store data on disk, you need to give Mnesia some information. Before you turn Mnesia on, you need to create a database, using the `mnesia:create_schema/1` function. For now, because you'll be getting started using only the local node, that will look like the following:

```
1> mnesia:create_schema([node()]).  
ok
```

By default, when you call `mnesia:create_schema/1`, Mnesia will store schema data in the directory where you are when you start it. If you look in the directory where you started Erlang, you'll see a new directory with a name like `Mnesia.nonode@nohost`. Initially, it holds a `LATEST.LOG` file and a `schema.DAT` file. The `node()` function just returns the identifier of the node you're on, which is fine when you're getting started. (If you want to change where Mnesia stores data, you can start Erlang with some extra options: `erl -mnesia dir " path "`. The `path` will be the location Mnesia keeps any disk-based storage.)



If you start Mnesia without calling `mnesia:create_schema/1`, Mnesia will keep its schema in memory, and it will vanish if and when Mnesia stops.

Unlike ETS and DETS, which are always available, you need to turn Mnesia on:

```
2> mnesia:start().  
ok
```

There's also an `mnesia:stop/0` function if you want to stop it.



If you run Mnesia on a computer that goes to sleep, you may get odd messages like `Mnesia(nonode@nohost): ** WARNING ** Mnesia is overloaded: {dump_log, time_threshold}` when it wakes up. Don't worry, it's a side-effect of waking up, and your data should still be safe. You probably shouldn't run production systems on devices that go to sleep, of course.

Creating Tables

Like ETS, Mnesia's basic concept of a table is a collection of records. It also offers `set`, `ordered_set`, and `bag` options, just like those in ETS, but doesn't offer `duplicate_bag`.

Mnesia wants to know more about your data than ETS, too. ETS pretty much takes data in tuples of any shape, counting only on there being a key it can use. The rest is up to you to interpret. Mnesia wants to know more about what you store, and takes a list of field names. The easy way to handle this is to define records and consistently use the field names from the records as Mnesia field names. There's even an easy way to pass the record names to Mnesia, using `record_info/2`.

The `planemos` table can work just as easily in Mnesia as in ETS, and some aspects of dealing with it will be easier. [Example 10-8](#), which is in `ch10/ex7-mnesia`, shows how to set up the `planemo` table in Mnesia. The `setup/0` method creates a schema, then starts Mnesia, and then creates a table based on the `planemo` record type. Once the table is created, it writes the values from [Table 10-1](#) to it.

Example 10-8. Setting up an Mnesia table of planemo properties

```
-module(drop).  
-export([setup/0]).  
-include("records.hrl").  
  
setup() ->  
    mnesia:create_schema([node()]),  
    mnesia:start(),  
    mnesia:create_table(planemo, [{attributes, record_info(fields, planemo)}]),  
  
    F = fun() ->  
        mnesia:write(  
            #planemo{ name=mercury, gravity=3.7, diameter=4878, distance_from_sun=57.9 }),  
        mnesia:write(  
            #planemo{ name=venus, gravity=8.9, diameter=12104, distance_from_sun=108.2 }) ,
```

```

mnesia:write(
  #planemo{ name=earth, gravity=9.8, diameter=12756, distance_from_sun=149.6 }),
mnesia:write(
  #planemo{ name=moon, gravity=1.6, diameter=3475, distance_from_sun=149.6 }),
mnesia:write(
  #planemo{ name=mars, gravity=3.7, diameter=6787, distance_from_sun=227.9 }),
mnesia:write(
  #planemo{ name=ceres, gravity=0.27, diameter=950, distance_from_sun=413.7 }),
mnesia:write(
  #planemo{ name=jupiter, gravity=23.1, diameter=142796, distance_from_sun=778.3 }),
mnesia:write(
  #planemo{ name=saturn, gravity=9.0, diameter=120660, distance_from_sun=1427.0 }),
mnesia:write(
  #planemo{ name=uranus, gravity=8.7, diameter=51118, distance_from_sun=2871.0 }),
mnesia:write(
  #planemo{ name=neptune, gravity=11.0, diameter=30200, distance_from_sun=4497.1 }),
mnesia:write(
  #planemo{ name=pluto, gravity=0.6, diameter=2300, distance_from_sun=5913.0 }),
mnesia:write(
  #planemo{ name=haumea, gravity=0.44, diameter=1150, distance_from_sun=6484.0 }),
mnesia:write(
  #planemo{ name=makemake, gravity=0.5, diameter=1500, distance_from_sun=6850.0 }),
mnesia:write(
  #planemo{ name=eris, gravity=0.8, diameter=2400, distance_from_sun=10210.0 })
end,
mnesia:transaction(F).

```

Apart from the setup, the key thing to note is that all of the writes are contained in a `fun` that is then passed to `mnesia:transaction` to be executed as a transaction. Mnesia will restart the transaction if there is other activity blocking it, so the code may get executed repeatedly before the transaction happens. Because of this, do not include any calls that create side effects to the function you'll be passing to `mnesia:transaction`, and don't try to catch exceptions on Mnesia functions within a transaction. If your function calls `mnesia:abort/1` (probably because some condition for executing it wasn't met), the transaction will be rolled back, returning a tuple beginning with `aborted` instead of `atomic`.



You may also want to explore the more flexible `mnesia:activity/2` when you need to mix more kinds of tasks in a transaction.

Your interactions with Mnesia should be contained in transactions, especially when your database is shared across multiple nodes. The main `mnesia:write`, `mnesia:read`, and `mnesia:delete` methods work only within transactions, period. There are `dirty_` methods, but every time you use them, especially to write data to the database, you're taking a risk.



Just as in ETS, you can overwrite values by writing a new value with the same key as a previous entry.

If you want to check on how this function worked out, try the `mnesia:table_info` function, which can tell you more than you want to know. The listing below is abbreviated to focus on key results.

```
1> c(drop).
{ok,drop}
2> rr("records.hrl").
[planemo,tower]
3> drop:setup().
{atomic,ok}
4> mnesia:table_info(planemo,all).
[{access_mode,read_write},
 {active_replicas,[nonode@nohost]},
 {all_nodes,[nonode@nohost]},
 {arity,5},
 {attributes,[name,gravity,diameter,distance_from_sun]},
 ...
 {memory,541},
 {ram_copies,[nonode@nohost]},
 {record_name,planemo},
 {record_validation,{planemo,5,set}},
 {type,set},
 {size,14},
 ...]
```

You can see which nodes are involved in the table (`nonode@nohost` is the default for the current node). `arity` in this case is the count of fields in the record, and `attributes` tells you what their names are. `ram_copies` plus the name of the current node tells you that this table is stored in memory locally. It is, as in the ETS example, of type `set`, and there are 14 records.



By default, Mnesia will store your table in RAM only (`ram_copies`) on the current node. This is speedy, but it means the data vanishes if the node crashes. If you specify `disc_copies` (note the spelling), Mnesia will keep a copy of the database on disk, but still use RAM for speed. You can also specify `disc_only_copies`, which will be slow. Unlike ETS, the table you create will still be around if the *process* that created it crashes, and will likely survive even a node crash so long as it wasn't only in RAM on a single node. By combining these options and (eventually) multiple nodes, you should be able to create fast and resilient systems.

The table is now set up, and you can start to use it. If you're running the Table Viewer, or start it with `tv:start()`, you can take a look at the contents of your Mnesia tables as well as your ETS tables. In the View menu, choose Mnesia Tables. The interface is similar to that for ETS tables.

Reading Data

Just like writes, you should wrap `mnesia:read` calls in a `fun`, which you then pass to `mnesia:transaction`. You can do that in the shell if you want to explore:

```
5> mnesia:transaction(fun() -> mnesia:read(planemo,neptune) end).  
{atomic,[#planemo{name = neptune,gravity = 11.0,  
diameter = 30300,distance_from_sun = 4497.1}]}
```

The result arrives as a tuple, which when successful contains `atomic` plus a list with the data from the table. The table data is packaged as a record, and you can get to its fields easily.

You can rewrite the `fall_velocity/2` function from [Example 10-8](#) to use an Mnesia transaction instead of an ETS call. The ETS version looked like the following:

```
fall_velocity(Planemo, Distance) when Distance >= 0 :-  
P=hd(ets:lookup(planemos,Planemo)),  
math:sqrt(2 * P#planemo.gravity * Distance).
```

Line 2 of the Mnesia version is a bit different.

```
fall_velocity(Planemo, Distance) when Distance >= 0 :-  
{atomic, [P | _]} = mnesia:transaction(fun() -> mnesia:read(planemo,Planemo)  
end),  
math:sqrt(2 * P#planemo.gravity * Distance).
```

Because Mnesia returns a tuple rather than a list, this uses pattern matching to extract the first item in the list contained in the second item of the tuple (and throws away the tail of that list with `_`). This table is a set, so there will always be only one item there. Then the data, contained in `P`, can be used for the same calculation as before.

If you compile and run this, you'll see a familiar result:

```
6> c(drop).
{ok,drop}
7> drop:fall_velocity(earth,20).
19.79898987322333
8> Pid1=spawn(mph_drop,mph_drop,[]).
<0.120.0>
9> Pid1 ! {earth,20}.
{earth,20}
On earth, a fall of 20 meters yields a velocity of 44.289078952755766 mph.
```

For these purposes, the simple `mnesia:read` is enough. You can tell Mnesia to build indexes for fields other than the key, and query those with `mnesia:index_read` as well.



If you want to delete records, you can run `mnesia:delete/2`, also inside of a transaction.

Query List Comprehensions

If Mnesia is really a database, it should be able to do more than key-value querying, right? It definitely can. You can use Erlang match specifications (as you can with ETS), but Query List Comprehensions (QLCs) are much more readable. They look like list comprehensions, which you saw in [Chapter 7](#), but operate on Mnesia tables rather than lists.

Suppose you want to find all the planemos with gravity less than that of earth. You could traverse the table with the `mnesia:first` and `mnesia:next` methods, but that seems like a lot of extra work. Instead, you can use the `qlc:q` function to hold a list comprehension and the `qlc:e` (or the equivalent but longer `qlc:eval`) function to process it. Then you run that inside of an `mnesia:transaction` call.



You can run query list comprehensions in the shell, but if you want to use them in modules you need to add `-include_lib("stdlib/include/qlc.hrl")`. to the declarations at the top of your module.

The simplest query list comprehension just returns all the values in the table. I've broken it out here on separate lines so that you can see how they interact:

```
mnesia:transaction(
  fun() ->
    qlc:e(
```

```

    qlc:q( [X || X <- mnesia:table(planemo)] )
)
end
)

```

As always, the `mnesia:transaction` function takes a fun as its argument. In this case, the fun contains a `qlc:e` function, which then contains a `qlc:q` function, where the real query is. It will build a list from the contents of the `planemo` table.

If you compact this a bit and run it in the shell, you'll see that the resulting list—wrapped in a transaction result tuple—contains the entire table.

```

10> mnesia:transaction( fun() -> qlc:e(qlc:q([X || X <- mnesia:table(planemo)])))
end).

{atomic,[#planemo{name = pluto,gravity = 0.6,
                  diameter = 2300,distance_from_sun = 5913.0},
        #planemo{name = saturn,gravity = 9.0,diameter = 120660,
                  distance_from_sun = 1427.0},
        #planemo{name = moon,gravity = 1.6,diameter = 3475,
                  distance_from_sun = 149.6},
        #planemo{name = mercury,gravity = 3.7,diameter = 4878,
                  distance_from_sun = 57.9},
        #planemo{name = earth,gravity = 9.8,diameter = 12756,
                  distance_from_sun = 149.6},
        #planemo{name = neptune,gravity = 11.0,diameter = 30200,
                  distance_from_sun = 4497.1},
        #planemo{name = makemake,gravity = 0.5,diameter = 1500,
                  distance_from_sun = 6850.0},
        #planemo{name = uranus,gravity = 8.7,diameter = 51118,
                  distance_from_sun = 2871.0},
        #planemo{name = ceres,gravity = 0.27,diameter = 950,
                  distance_from_sun = 413.7},
        #planemo{name = venus,gravity = 8.9,diameter = 12104,
                  distance_from_sun = 108.2},
        #planemo{name = mars,gravity = 3.7,diameter = 6787,
                  distance_from_sun = 227.9},
        #planemo{name = eris,gravity = 0.8,diameter = 2400,
                  distance_from_sun = 10210.0},
        #planemo{name = jupiter,gravity = 23.1,diameter = 142796,
                  distance_from_sun = 778.3},
        #planemo{name = haumea,gravity = 0.44,diameter = 1150,
                  distance_from_sun = 6484.0}]}

```

You can add conditions to the query list comprehension. You might want to know all of the planemos with gravity less than that of Earth's 9.8. That could look like the following:

```

mnesia:transaction(
fun() ->
qlc:e(
qlc:q( [X || X <- mnesia:table(planemo),

```

```

        X#planemo.gravity < 9.8] )
    )
end
)

```

Compress and run that in the shell, and you'll get a shorter list of planemos where everything feels a little lighter.

```

11> mnesia:transaction( fun() -> qlc:e(qlc:q( [X || X <- mnesia:table(planemo),
X#planemo.gravity < 9.8] )) end).

{atomic,[#planemo{name = pluto,gravity = 0.6,
diameter = 2300,distance_from_sun = 5913.0},
#planemo{name = saturn,gravity = 9.0,diameter = 120660,
distance_from_sun = 1427.0},
#planemo{name = moon,gravity = 1.6,diameter = 3475,
distance_from_sun = 149.6},
#planemo{name = mercury,gravity = 3.7,diameter = 4878,
distance_from_sun = 57.9},
#planemo{name = makemake,gravity = 0.5,diameter = 1500,
distance_from_sun = 6850.0},
#planemo{name = uranus,gravity = 8.7,diameter = 51118,
distance_from_sun = 2871.0},
#planemo{name = ceres,gravity = 0.27,diameter = 950,
distance_from_sun = 413.7},
#planemo{name = venus,gravity = 8.9,diameter = 12104,
distance_from_sun = 108.2},
#planemo{name = mars,gravity = 3.7,diameter = 6787,
distance_from_sun = 227.9},
#planemo{name = eris,gravity = 0.8,diameter = 2400,
distance_from_sun = 10210.0},
#planemo{name = haumea,gravity = 0.44,diameter = 1150,
distance_from_sun = 6484.0}]}

```

That still contains more information than might be necessary. You can modify the left side of the comprehension to cut things down, creating a tuple that is just the name and gravity of the planemo:

```

mnesia:transaction(
fun() ->
qlc:e(
qlc:q( [{X#planemo.name, X#planemo.gravity} ||
X <- mnesia:table(planemo),
X#planemo.gravity < 9.8] )
)
end
)

```

The result is much trimmer:

```

12> mnesia:transaction( fun() -> qlc:e(qlc:q( [ {X#planemo.name, X#planemo.gravi-
ty}] ||
X <- mnesia:table(planemo), X#planemo.gravity < 9.8] )) end).

{atomic,[{pluto,0.6},

```

```
{saturn,9.0},  
{moon,1.6},  
{mercury,3.7},  
{makemake,0.5},  
{uranus,8.7},  
{ceres,0.27},  
{venus,8.9},  
{mars,3.7},  
{eris,0.8},  
{haumea,0.44}]]}
```

There are ways to reduce at least some of the syntax overhead here. It's not difficult, for example, to move the `mnesia:transaction`, fun definition, and `qlc:e` call to a function that takes the `qlc:q` function as its argument. In *Programming Erlang*, Joe Armstrong does just that to create a `do` function. You may want to break things up differently depending on how you're working with data and your coding style.



You can use query list comprehensions on more than one table at a time, which is how you can create the equivalent of joins between tables, and it is also possible to use them on ETS tables.

This is just a brief introduction to Mnesia. It gets some coverage in all of the Erlang books, but eventually I hope it will get a book of its own, about as long as this one.

Getting Started with OTP

At this point, it might seem like you have all you need to create process-oriented projects with Erlang. You know how to create useful functions, can work with recursion, know the data structures Erlang offers, and probably most important, know how to create and manage processes. What more could you need?

Process-oriented programming is great, but the details matter. The basic Erlang tools are powerful, but can also bring you to frustrating mazes debugging race conditions that happen only once in a while. Mixing different programming styles can lead to incompatible expectations, and code that worked well in one environment may prove harder to integrate in another.

Ericsson encountered these problems early, and created a set of libraries that eases them. OTP, the Open Telecom Platform, is useful for pretty much any large-scale project you want to do with Erlang, not just telecom work. It's included with Erlang, and though it isn't precisely part of the language, it is definitely part of Erlang culture. The boundaries of where Erlang ends and OTP begins aren't always clear, but the entrypoint is definitely behaviors. You'll combine processes built with behaviors and managed by supervisors into an OTP application.

So far, the lifecycle of the processes shown in the previous chapters has been pretty simple. If needed, they set up other resources or processes to get started. Once running, they listen for messages and process them, collapsing if they fail. Some of them might restart a failed process if needed.

OTP formalizes those activities, and a few more, into a set of behaviors (or behaviours —this was originally created with British spelling). The most common behaviors are `gen_server` (generic server) and `supervisor`, though `gen_fsm` (finite state machine) and `gen_event` are also available. The `application` behavior lets you package your OTP code into a single runnable (and updatable) system.

The behaviors pre-define the mechanisms you'll use to create and interact with processes, and the compiler will warn you if you're missing some of them. Your code will handle the callbacks, specifying how to respond to particular kinds of events, and you will need to decide upon a structure for your application.



If you'd like a free one-hour video introduction to OTP, see Steve Vinoski's "Erlang's Open Telecom Platform (OTP) Framework" at <http://www.infoq.com/presentations/Erlang-OTP-Behaviors>. You probably already know the first half hour or so of it, but the review is excellent. In a very different style, if you'd like an explanation of why it's worth learning OTP and process-oriented development in general, Francesco Cesarini's slides at <https://www.erlang-factory.com/upload/presentations/719/francesco-otp.pdf> work even without narration, especially the second half.

Creating Services with `gen_server`

Much of the work you think of as the core of a program—calculating results, storing information, and preparing replies—will fit neatly into the `gen_server` behavior. It provides a core set of methods that let you set up a process, respond to requests, end the process gracefully, and even pass state to a new process if this one needs to be upgraded in place.

[Table 11-1](#) shows the methods you need to implement in a service that uses the `gen_server` behavior. For a simple service, the first two or three are the most important, and you may just use placeholder code for the rest.

Table 11-1. What calls and gets called in `gen_server`

Method	Triggered by	Does
<code>init/1</code>	<code>gen_server:start_link</code>	Sets up the process
<code>handle_call/3</code>	<code>gen_server:call</code>	Handles synchronous calls
<code>handle_cast/2</code>	<code>gen_server:cast</code>	Handles asynchronous calls
<code>handle_info/2</code>	random messages	Deals with non-OTP messages
<code>terminate/2</code>	failure or shutdown signal from supervisor	Cleans up the process
<code>code_change/3</code>	system libraries for code upgrades	Lets you switch out code without losing state

[Appendix B](#) shows a complete `gen_server` template from the Erlang emacs-mode, which is worth exploring in particular for the models it offers for the return value. However, it's pretty big. [Example 11-1](#), which you can find in `ch11/ex1-drop`, shows a less verbose example (based on the template) that you can use to get started. It mixes a simple calculation from way back in [Example 2-1](#) with a counter like that in [Example 8-4](#).

Example 11-1. A simple gen_server example based on the template from the Erlang mode for Emacs

```
-module(drop).
-behaviour(gen_server).
-export([start_link/0]). % convenience call for startup
-export([init/1,
        handle_call/3,
        handle_cast/2,
        handle_info/2,
        terminate/2,
        code_change/3]). % gen_server callbacks
#define(SERVER, ?MODULE). % macro that just defines this module as server
-record(state, {count}). % simple counter state

%%% convenience method for startup
start_link() ->
    gen_server:start_link({local, ?SERVER}, ?MODULE, [], []).

%%% gen_server callbacks
init([]) ->
    {ok, #state{count=0}}.

handle_call(_Request, _From, State) ->
    Distance = _Request,
    Reply = {ok, fall_velocity(Distance)},
    NewState=#state{ count = State#state.count+1 },
    {reply, Reply, NewState}.

handle_cast(_Msg, State) ->
    io:format("So far, calculated ~w velocities.\n", [State#state.count]),
    {noreply, State}.

handle_info(_Info, State) ->
    {noreply, State}.

terminate(_Reason, _State) ->
    ok.

code_change(_OldVsn, State, _Extra) ->
    {ok, State}.

%%% Internal functions

fall_velocity(Distance) -> math:sqrt(2 * 9.8 * Distance).
```

The module name (`drop`) should be familiar from past examples. The second line is a `-behaviour` declaration specifying that this is going to be using the `gen_server` behavior. That declaration tells Erlang that it can expect this code to support the core callback functions of that behavior.



You can spell the `-behaviour` declaration `-behavior` if you prefer American spellings. Erlang doesn't mind.

The `-export` declarations are pretty standard, though they break out the `start_link/0` method into a separate declaration from the core `gen_server` methods. This isn't necessary, but it's a nice reminder that `start_link` isn't required for the `gen_server` behavior to work. (It *calls* `gen_server` code, but isn't a callback itself.)

The `-define` declaration is probably unfamiliar. Erlang lets you declare macros using `-define`. Macros are simple text replacements. This declaration tells the compiler that any time it encounters `?SERVER`, it should replace it with `?MODULE`. What is `?MODULE`? That's a built-in macro that always refers to the name of the module it appears in. In this case, that means it will be processed into `drop`. (You may find cases where you might want to register the server under a name other than the module name, but this is a workable default.)

The `-record` declaration should be familiar, though it contains only one field, to keep a count of the number of calls made. Many services will have more fields, including things like database connections, references to other processes, perhaps network information, and metadata specific to this particular service. It is also possible to have services with no state, which would be represented by an empty tuple here. As you'll see further down, every single `gen_server` function will reference the state.



The `state` record declaration is a good example of a record declaration you should make inside of a module and not declare through an included file. It is possible that you'll want to share state models across different `gen_server` processes, but it's easier to see what `State` should contain if the information is right there.

The first function in the sample, `start_link/0`, is *not* one of the required `gen_server` functions. Instead, it calls `gen_server`'s `start_link` function to start up the process. When you're just getting started, this is useful for testing. As you move toward production code, you may find it easier to leave these out and use other mechanisms.

The `start_link/0` function uses the `?SERVER` macro defined in the `-define` declaration as well as the built-in `?MODULE` declaration.

```
%%% convenience method for startup
start_link() ->
    gen_server:start_link({local, ?SERVER}, ?MODULE, [], []).
```

The first argument, a tuple, opens with an atom that must be `local` or `global`, depending on whether you want the name of the process registered just with the local Erlang

instance or with all associated nodes. The `?SERVER` macro will be expanded to `?MODULE` which will itself be expanded to the name of the current module, and that will be used as the name for this process. The second argument is the name of the module, here identified with the `?MODULE` macro, and then lists for arguments and options follow. In this case, they're both empty. Options can specify things like debugging, timeouts, and options for spawning the process.



You may also see a form of `gen_server:start_link` with `via` as the atom in the first tuple. This lets you set up custom process registries, of which `gproc` is the best known. For more on that, see <https://github.com/uwiger/gproc>.

All of the remaining functions are part of the `gen_server` behavior. `init/1` creates a new state record instance and sets its `count` field to zero—no velocities have yet been calculated. The two functions that do much here are `handle_call/3` and `handle_cast/2`. For this demonstration, `handle_call/3` expects to receive a distance in meters and returns a velocity for a fall from that height on earth, while `handle_cast/2` is a trigger to report the number of velocities calculated.

`handle_call/3` makes synchronous communications between Erlang processes simple.

```
handle_call(_Request, _From, State) ->
    Distance = _Request,
    Reply = {ok, fall_velocity(Distance)},
    NewState = #state{ count = State#state.count+1 },
    {reply, Reply, NewState}.
```

This extracts the `Distance` from the `_Request`, which isn't necessary except that I wanted to leave the variable names for the function the same as they were in the template. (`(handle_call(Distance, _From, State)` would have been fine.) Your `_Request` is more likely to be a tuple or a list rather than a bare value, but this works for simple calls.

It then creates a reply based on sending that `Distance` to the simple `fall_velocity/1` function at the end of the module. It then creates a `NewState` containing an incremented count. Then the atom `reply`, the `Reply` tuple containing the velocity, and the `NewState` containing the count get passed back.

Because the calculation is really simple, treating the drop as a simple synchronous call is perfectly acceptable. For more complex situations where you can't predict how long a response might take, you may want to consider responding with a `noreply` response and using the `_From` argument to send a response later. (There is also a `stop` response available that will trigger the `terminate/2` method and halt the process.)



By default, Erlang will time out any synchronous calls that take longer than five seconds to calculate. You can override this by making your call using `gen_server:call/3` to specify a timeout (in milliseconds) explicitly, or by using the atom `infinity`.

The `handle_cast/2` function supports asynchronous communications. It isn't supposed to return anything directly, though it does report `noreply` (or `stop`) and updated state. In this case, it takes a very weak approach, but one that does well for a demonstration, calling `io:format/2` to report on the number of calls.

```
handle_cast(_Msg, State) ->
    {io:format("So far, calculated ~w velocities.~n", [State#state.count]),
     {noreply, State}}.
```

The state doesn't change, because asking for the number of times the process has calculated a fall velocity is not the same thing as actually calculating a fall velocity.

Until you have good reason to change them, you can leave `handle_info/2`, `terminate/2`, and `code_change/3` alone.

Making a `gen_server` process run and calling it looks a little different than starting the processes you saw in [Chapter 8](#).

```
1> c(drop).
{ok,drop}
2> drop:start_link().
{ok,<0.33.0>}
3> gen_server:call(drop, 20).
{ok,19.79898987322333}
4> gen_server:call(drop, 40).
{ok,28.0}
5> gen_server:call(drop, 60).
{ok,34.292856398964496}
6> gen_server:cast(drop, {}).
So far, calculated 3 velocities.
ok
```

The call to `drop:start_link()` sets up the process and makes it available. Then, you're free to use `gen_server:call` or `gen_server:cast` to send it messages and get responses.



While you can capture the pid, you don't have to keep it around to use the process. Because `start_link` returns a tuple, if you want to capture the pid you can do something like `{ok, Pid} = drop:start_link()`.

Because of the way OTP calls `gen_server` functions, there's an additional bonus—or perhaps a hazard—in that you can update code on the fly. For example, I tweaked the `fall_velocity/1` function to lighten Earth's gravity a little, using 9.1 as a constant instead of 9.8. Recompiling the code and asking for a velocity returns a different answer:

```
7> c(drop).
{ok,drop}
8> gen_server:call(drop, 60).
{ok,33.04542328371661}
```

This can be very convenient during the development phase, but be careful doing anything like this on a production machine. OTP has other mechanisms for updating code on the fly. There is also a built-in limitation to this approach: `init` gets called only when `start_link` sets up the service. It does not get called if you recompiled the code. If your new code requires any changes to the structure of its state, your code will break the next time it's called.

A Simple Supervisor

When you started the `drop` module from the shell, you effectively made the shell the supervisor for the module—though the shell doesn't really do any supervision. You can break the module easily:

```
9> gen_server:call(drop, -60).

=ERROR REPORT==== 2-Dec-2012::21:14:51 ===
** Generic server drop terminating
** Last message in was -60
** When Server state == {state,0}
** Reason for termination ==
** {badarith,[{math,sqrt,[-1176.0],[]},
             {drop,fall_velocity,1,[{file,"drop.erl"},{line,42}]},
             {drop,handle_call,3,[{file,"drop.erl"},{line,23}]},
             {gen_server,handle_msg,5,[{file,"gen_server.erl"},{line,588}]},
             {proc_lib,init_p_do_apply,3,
              [{file,"proc_lib.erl"},{line,227}]}]}
** exception exit: badarith
  in function  math:sqrt/1
    called as math:sqrt(-1176.0)
  in call from drop:fall_velocity/1 (drop.erl, line 42)
  in call from drop:handle_call/3 (drop.erl, line 23)
  in call from gen_server:handle_msg/5 (gen_server.erl, line 588)
  in call from proc_lib:init_p_do_apply/3 (proc_lib.erl, line 227)
10> gen_server:call(drop, 60).
** exception exit: {noproc,[{gen_server,call,[drop,60]}]}
  in function  gen_server:call/2 (gen_server.erl, line 180)
```

The error message is nicely complete, even telling you the last message and the state, but when you go to call the service again on line 10, it isn't there. You can restart it with `drop:start_link/0` again, but you're not always going to be watching your processes personally.

Instead, you want something that can watch over your processes and make sure they restart (or not) as appropriate. OTP formalizes the process management you saw in [Example 8-10](#) with its supervisor behavior. [Example B-2](#) in [Appendix B](#) shows a full template (again, from the Erlang mode for Emacs), but you can create a less verbose supervisor.

A basic supervisor needs to support only one callback function, `init/1`, and can also have a `start_link` function to fire it up. The return value of that `init/1` function tells OTP which child processes your supervisor manages and how you want to handle their failures. A supervisor for the `drop` module might look like [Example 11-2](#), which is in `ch11/ex2-drop-sup`.

Example 11-2. A simple supervisor

```
-module(drop_sup).
-behaviour(supervisor).
-export([start_link/0]). % convenience call for startup
-export([init/1]). % supervisor calls
-define(SERVER, ?MODULE).

%%% convenience method for startup
start_link() ->
    supervisor:start_link({local, ?SERVER}, ?MODULE, []).

%%% supervisor callback
init([]) ->
    RestartStrategy = one_for_one,
    MaxRestarts = 1, % one restart every
    MaxSecondsBetweenRestarts = 5, % five seconds

    SupFlags = {RestartStrategy, MaxRestarts, MaxSecondsBetweenRestarts},

    Restart = permanent, % or temporary, or transient
    Shutdown = 2000, % milliseconds, could be infinity or brutal_kill
    Type = worker, % could also be supervisor

    Drop = {drop, {drop, start_link, []},
            {Restart, Shutdown, Type, [drop]}},

    {ok, {SupFlags, [Drop]}}.

%%% Internal functions (none here)
```

The `init/1` function's job is to assemble a fairly complex data structure. The first few variables from the template define how the supervisor should handle failure.

The `RestartStrategy` of `one_for_one` tells OTP that it should create a new child process every time a process that is supposed to be `permanent` fails. You can also go with `one_for_all`, which terminates and restarts all of the processes the supervisor oversees when one fails, or `rest_for_one`, which restarts the process and any processes that began after the failed process had started.



When you're ready to take more direct control of how your processes respond to their environment, you might explore working with the dynamic functions `supervisor:start/2`, `supervisor:terminate_child/2`, `supervisor:restart_child/2`, and `supervisor:delete_child/2`, as well as the restart strategy `simple_one_for_one`.

The next two values define how often the worker processes can crash before terminating the supervisor itself. In this case, it's one restart every five seconds. Customizing these values lets you handle a variety of conditions, but probably won't affect you much initially. (Setting `MaxRestarts` to zero means that the supervisor will just terminate if a worker has an error.)

Those values, which here get combined into the tuple contained in `SupFlags`, apply to *all* of the workers managed by this supervisor. The next few lines define properties that apply to only one worker process, in this case the `gen_server` specified by `Drop`. It is to be a `permanent` service, so the supervisor should restart it when it fails. The supervisor can wait two seconds before shutting it off completely, and this worker is *only* a worker, not itself a supervisor. More complex OTP applications can contain trees of supervisors managing other supervisors, which themselves manage other supervisors or workers.

The `Drop` variable assignment might seem a bit repetitive, but it creates a complete set of information for that process. First it specifies a name, and then a tuple containing the name of the module containing the code, the function to use to start the process and a list of arguments. (Here there aren't any arguments.) Then the `Restart`, `Shutdown`, and `Type` are specified, and the final list identifies all the modules on which this process will depend. In this case, it all fits into a single module, so the list contains only the name of that module.



OTP wants to know the dependencies so that it can help you upgrade software in place. It's all part of the magic of keeping systems running without ever bringing them to a full stop.

Now that you have a supervisor process, you can set up the drop function by just calling the supervisor. However, running a supervisor from the shell using the `start_link/0` function call creates its own set of problems; the shell is itself a supervisor, and will terminate processes that report errors. After a long error report, you'll find that both your worker and the supervisor have vanished.

In practice this means that there are two ways to test supervised OTP processes (that aren't yet part of an application) directly from the shell. The first explicitly breaks the bond between the shell and the supervisor process by catching the pid of the supervisor (line 2) and then using the `unlink/1` function to remove the link (line 3). Then you can call the process as usual with `gen_server:call/2` and get answers. If you get an error (line 6), it'll be okay. The supervisor will restart the worker, and you can make new calls (line 7) successfully. The calls to `whereis(drop)` on lines 5 and 8 demonstrate that the supervisor has restarted `drop` with a new pid.

```
1> c(drop_sup).
{ok,drop_sup}
2>{ok, Pid} = drop_sup:start_link().
{ok,<0.38.0>}
3> unlink(Pid).
true
4> gen_server:call(drop, 60).
{ok,34.292856398964496}
5> whereis(drop).
<0.39.0>
6> gen_server:call(drop, -60).

=ERROR REPORT==== 2-Dec-2012::21:17:19 ===
** Generic server drop terminating
** Last message in was -60
** When Server state == {state,1}
** Reason for termination ==
** {{badarith,[{math,sqrt,[-1176.0],[]}],
     [{drop,fall_velocity,1,[{file,"drop.erl"},{line,42}]}],
     [{drop,handle_call,3,[{file,"drop.erl"},{line,23}]}],
     [{gen_server,handle_msg,5,[{file,"gen_server.erl"},{line,588}]}],
     [{proc_lib,init_p_do_apply,3,
       [{file,"proc_lib.erl"},{line,227}]}]}}
** exception exit: {{badarith,
     [{math,sqrt,[-1176.0],[]},
      {drop,fall_velocity,1,
        [{file,"drop.erl"},{line,42}]},
      {drop,handle_call,3,
        [{file,"drop.erl"},{line,23}]},
      {gen_server,handle_msg,5,
        [{file,"gen_server.erl"},{line,588}]}],
      [{proc_lib,init_p_do_apply,3,
        [{file,"proc_lib.erl"},{line,227}]}]}},
     {gen_server,call,[drop,-60]}]
in function  gen_server:call/2 (gen_server.erl, line 180)
```

```

7> gen_server:call(drop, 60).
{ok,34.292856398964496}
8> whereis(drop).
<0.44.0>

```

The other approach leaves the link in place, but wraps the calls to `gen_server/2` in a `catch` statement. In this case, using `catch` just keeps the shell from ever receiving the exception, so the supervisor remains untouched. You don't *have* to use `catch` to make a call, as line 8 shows, but if the call fails, you'll have to restart the supervisor process yourself. (Line 6 is also a bit split by the error message. Sometimes the timing will make it look like the prompt disappeared. Don't worry.)

```

1> c(drop_sup).
{ok,drop_sup}
2> drop_sup:start_link().
{ok,<0.38.0>}
3> whereis(drop).
<0.39.0>
4> catch gen_server:call(drop, 60).
{ok,34.292856398964496}
5>
5> catch gen_server:call(drop, -60).
{'EXIT',{badarith,[{math,sqrt,[-1176.0],[]}],
         [{drop,fall_velocity,1,[{file,"drop.erl"},{line,42}]}],
         [{drop,handle_call,3,[{file,"drop.erl"},{line,23}]}],
         [{gen_server,handle_msg,5,
           [{file,"gen_server.erl"},{line,588}]}],
         [{proc_lib,init_p_do_apply,3,
           [{file,"proc_lib.erl"},{line,227}]}]},
          {gen_server,call,[drop,-60]}}}
6>
=ERROR REPORT==== 2-Dec-2012::21:21:10 ===
** Generic server drop terminating
** Last message in was -60
** When Server state == {state,1}
** Reason for termination ==
** {badarith,[{math,sqrt,[-1176.0],[]},
           [{drop,fall_velocity,1,[{file,"drop.erl"},{line,42}]}],
           [{drop,handle_call,3,[{file,"drop.erl"},{line,23}]}],
           [{gen_server,handle_msg,5,[{file,"gen_server.erl"},{line,588}]}],
           [{proc_lib,init_p_do_apply,3,
             [{file,"proc_lib.erl"},{line,227}]}]}
catch gen_server:call(drop, 60).
{ok,34.292856398964496}
7> whereis(drop).
<0.43.0>
8> gen_server:call(drop, 60).
{ok,34.292856398964496}

```



You can also tell the shell to stop worrying about such exceptions by issuing the shell command `catch_exception(true)`. However, that turns off the behavior for the entire shell, which may not be what you want. (It will return `false`, the previous setting for that property. Don't worry, it did set it to `true`.)

You can also open Process Manager or Observer and whack away at worker processes through the Kill option on the Trace menu and watch them reappear.

This works, but is the tiniest taste of what supervisors can do. They can create child processes dynamically, and manage their lifecycle in greater detail.

Packaging an Application

OTP also lets you package sets of components into an application. While stopping and starting OTP workers and supervisors may be easier than dealing with processes directly, OTP's facilities for describing applications will lead you down a path to much easier starting, stopping (if you must), updating, and administering your projects.

Erlang applications include two extra components beyond the workers, supervisors, and related files they need. The application resource file, or app file, provides a lot of metadata about your application. You'll also need a module with the behavior `application` to define starting and stopping.



If you're on a Mac, the file extension for the `.app` file will disappear and the operating system will think it's some kind of broken Mac application. Don't worry. It'll still work in Erlang, though the Mac won't know what to do if you double-click on it.

The app file is a large tuple, though easier to read than the one returned by a supervisor's `init/1` functions. [Example 11-3](#), in `ch11/ex3-drop-app`, shows a minimal app file, placed in an `ebin` subdirectory, that sets up this simple drop application.

Example 11-3. drop.app, an application resource file, or app file, for the drop program

```
{application, drop,
[{description, "Dropping objects from towers"},
{vsn, "0.0.1"}, {modules, [drop, drop_sup, drop_app]}, {registered, [drop, drop_sup]}, {applications, [kernel, stdlib]}, {mod, {drop_app, []}}]}.
```

The first line identifies this as an application named `drop`, and then a list of arguments provides more information:

- The `description` is a (sometimes) human-friendly description of what's here. `vsn` is a version number, which in this case is tiny.
- `modules` lists the modules that make up the application, in this case `drop`, `drop_sup`, and `drop_app`.
- `registered` lists modules that are publicly visible, again `drop` and `drop_sup`.
- `applications` lists the required applications on which this application depends, and the `kernel` and `stdlib` seem to be the minimal standard set.
- `mod` has a tuple that points to the module with the `application` behavior. It can take a list of arguments that will go to the `start/2` function of the module, though there aren't any here.

That module is trivial even compared to the other OTP code you've seen, as shown in [Example 11-4](#), which is also in *ch11/ex3-drop-app*. ([Example B-3](#) shows a fuller template.)

Example 11-4. The application module for the drop program

```
-module(drop_app).
-behaviour(application).
-export([start/2, stop/1]).

start(_Type, _StartArgs) ->
    drop_sup:start_link().

stop(_State) ->
    ok.
```

The only thing you really have to do is start up the supervisors for your application in the `start/2` function. In this case there's only one, and the `_Type` and `_StartArgs` don't matter.

Running this application from the shell will require one bit of extra effort on your part. You'll need to compile `drop_app`, of course, but you'll also need to tell Erlang about the `ebin` directory containing the `drop.app` file, as shown on line 2. (OTP expects it to be there, but will give you "no such file or directory" errors if you don't tell Erlang about the directory.)

```
1> c(drop_app).
{ok,drop_app}
2> code:add_path("ebin/").
true
3> application:load(drop).
ok
4> application:loaded_applications().
```

```

[{kernel,"ERTS CXC 138 10","2.15.2"},  

 {drop,"Dropping objects from towers","0.0.1"},  

 {stdlib,"ERTS CXC 138 10","1.18.2"}]  

5> application:start(drop).  

ok  

6> gen_server:call(drop, 60).  

{ok,34.292856398964496}

```

Once Erlang knows where to look, you can use the `application` module's functions to load the application and check that Erlang found it. Once you `start` the application, you can go ahead and make calls to it with `gen_server:call`. Because the supervisor is bound to an application, you don't need to worry about the shell shutting you down. You can go ahead and break the drop calculation process with a negative value, and the supervisor will just fire it back up.

```

7> whereis(drop).  

<0.45.0>  

8> gen_server:call(drop, -60).  

=ERROR REPORT==== 2-Dec-2012::21:25:38 ===  

** Generic server drop terminating  

** Last message in was -60  

** When Server state == {state,1}  

** Reason for termination ==  

** {badarith,[{math,sqrt,[-1176.0],[]},  

           {drop,fall_velocity,1,[{file,"drop.erl"},{line,42}]},  

           {drop,handle_call,3,[{file,"drop.erl"},{line,23}]},  

           {gen_server,handle_msg,5,[{file,"gen_server.erl"},{line,588}]},  

           {proc_lib,init_p_do_apply,3,  

            [{file,"proc_lib.erl"},{line,227}]}]}  

** exception exit: {badarith,  

                  [{math,sqrt,[-1176.0],[]},  

                   {drop,fall_velocity,1,  

                    [{file,"drop.erl"},{line,42}]},  

                   {drop,handle_call,3,  

                    [{file,"drop.erl"},{line,23}]},  

                   {gen_server,handle_msg,5,  

                    [{file,"gen_server.erl"},{line,588}]},  

                   {proc_lib,init_p_do_apply,3,  

                    [{file,"proc_lib.erl"},{line,227}]}]},  

                  {gen_server,call,[drop,-60]}}  

in function  gen_server:call/2 (gen_server.erl, line 180)  

9> gen_server:call(drop, 60).  

{ok,34.292856398964496}  

10> whereis(drop).  

<0.49.0>

```

There is much, much more to learn. OTP deserves a book or several all on its own. Hopefully this chapter provides you with enough information to try some things out and understand those books. However, the gap between what this chapter can reasonably present and what you need to know to write solid OTP-based programs is... vast.

Next Steps Through Erlang

Hopefully you now feel comfortable writing basic Erlang programs, and understand roughly how modules and processes build into programs. You should be ready to experiment with writing Erlang code, but more importantly, you should be ready to explore other resources for mastering Erlang and its many powerful libraries. There's a lot to explore!

Moving Beyond the Shell

The Erlang shell is a great place to test code, and to poke and prod Erlang code. You'll likely spend a lot more time in the shell if you keep using Erlang, but the way you use it may change.

You can compile and run Erlang code outside of the shell, and this makes it much easier to integrate Erlang work with tools you typically use to manage code and related resources. Erlang's `make` module is a common place to start, letting you create *Emakefile* files that provide instruction to the `erl -make` command.

If you want to automate your Erlang builds further, there are many ways to integrate Erlang with the classic Make toolset. For a more thoroughly Erlang-centric build experience, you may want to explore rebar, at <https://github.com/rebar/rebar>. You can even mix Make and rebar to apply the strengths of each.

If you want to use Erlang from an IDE, you may want to explore <http://erlide.org/>, a set of tools for working with Erlang in Eclipse. Emacs users will want to explore the Erlang mode.

Distributed Computing

Almost everything you've learned in this book points toward a computing model that makes it easy to distribute programs across a network of Erlang nodes. Setting up a set of nodes isn't that difficult. It may even be too easy in some ways to let security-obsessed administrators sleep easily.

Before you set up large sets of nodes, you'll want to know much more about how Erlang schedules code to run, how messages get passed among nodes, and how to administer Erlang nodes remotely. All of that is out there, providing the foundation on which tools like OTP and Mnesia can build.

Processing Binary Data

Erlang includes a binary data type, binary operators, and a variety of libraries for processing binary data. If you need to build network protocols that handle bits on the wire, or ASN.1 data, for example, Erlang offers powerful tools for getting information into and out of binary form. If you see numbers or strings enclosed in <> and >>, you've encountered Erlang's tools for handling binary data. They allow you to specify, pattern match, and process binary data structures.

Input and Output

Chapter 4 and Chapter 5 introduced you to the `io:format` functions in the context of presenting information in the shell. The `io` module offers much more, however, for reading and writing data, and the `file`, `filename`, `filelib`, and `io_lib` modules give you the tools you need to get into and out of files. If networking is more your style, you'll want to explore the `gen_tcp`, `gen_udp`, and `inet` modules.

Testing, Analyzing, and Refactoring

Functional programming approaches should, once you get used to them, make it easier to create clean code. However, it is always possible to tie yourself in knots, especially as you move toward solving more complex problems than the ones presented in this book.

Unit testing is one approach to making sure that your code keeps working as you move forward. Focused on small components of your programs that should be able to reliably return a set of correct outputs from a given set of inputs, unit testing can both help to tell you when you've made your code work and warn you when it breaks. Erlang includes the EUnit framework for unit testing and the Common Test framework for system testing.

Erlang also includes Dialyzer, the Discrepancy Analyzer for Erlang, which can help you catch basic errors of sending wrongly typed data, code that never gets called, and similar issues that compilers for statically typed languages are usually good at catching. Erlang also includes profilers and coverage tools; you should explore the `erlprof`, `fprof`, and `cover` modules, and the `cprof` tool.

If you're excited about refactoring code—once you have more of it—you may want to explore Wrangler, at <http://www.cs.kent.ac.uk/projects/forse/wrangler/doc/overview-summary.html>. It allows you to explore your code and automate a wide variety of common program modifications.

Networking and the Web

Erlang is a natural match for a world of web programming in which the number of users is constantly increasing, data needs to move smoothly among nodes, and customers are starting to expect critical web applications to work as reliably as the telephone system. Most web applications are large pipelines of data, well-suited for Erlang's strengths.

Several Erlang frameworks let you build web applications. Yaws (<http://yaws.hyber.org/>), a web server itself written in Erlang, offers an environment that should be fairly familiar to anyone building web applications. Yaws lets you mix Erlang code with HTML (and other) templates, making it fairly easy to get a website or application together on an Erlang foundation. If Yaws doesn't look like it's for you, you may also want to explore Mochiweb (<https://github.com/mochi/mochiweb>) or Cowboy (<https://github.com/extend/cowboy>). For a more comprehensive framework that can run with any of these, take a look at Nitrogen (<http://nitrogenproject.com/>).

If you have fallen for the siren song of REST-based service development, you may also want to explore Webmachine (<http://wiki.basho.com/Webmachine.html>), a toolkit for HTTP processing that brings you very close to the core of the Web's foundation protocol. Even if you don't end up using it, exploring its flow diagram at <http://wiki.basho.com/Webmachine-Diagram.html> will teach you a lot about what's involved in processing a web request.

Data Storage

You already have ETS, DETS, and Mnesia. What else might you need?

Many people are using Erlang without knowing they are using it, as they interact with the popular NoSQL databases CouchDB (<http://couchdb.apache.org/>) and Riak (<http://basho.com/products/riak-overview/>). Their Erlang underpinnings make them easy to distribute and manage, and they've both reached large and growing audiences. For a

great brief intro to them (and five other database options), you should explore *Seven Databases in Seven Weeks* (Pragmatic Programmers). It won't teach you much about using them with Erlang, but it will give you a solid foundation that will help you explore their Erlang interfaces once their broad approaches make sense to you.

Many other databases have Erlang interfaces, and there is support for the classic ODBC connections.

Extending Erlang

If you need to wring out every drop of performance on a complex task, or want to avoid rewriting a library written in a language other than Erlang, you'll want to explore Erlang's tools for connecting with other programming languages. *Erlang Programming* explores Java, C, and Ruby connections, but also notes approaches you can use to connect with .NET languages, Python, Perl, PHP, Haskell, Scheme, and Emacs Lisp. You'll want to examine native implemented functions (NIFs) and drivers.

Languages Built on Erlang

Erlang may put the fun in functional programming, but its structures may feel brittle if you're used to the focus on flexibility that many other languages provide. Elixir (<http://elixir-lang.org/>) combines the Erlang Runtime System with a very different syntax more focused on polymorphism, meta-programming, and associative data structures. (One of Erlang's creators, Joe Armstrong, is also working in similar directions with his Erl2 work, which you can find at <https://github.com/joeearns/erl2>.)

Erlang's runtime model and tools are powerful and unique, and there may be other great ideas coming that will let you apply them to work that might not seem on the surface to be written in Erlang.

Community

As you learn more about Erlang, you'll find a community happy to help you at every level. The erlang-questions mailing list is at its heart, and welcomes beginners. During the writing of this book, I found its archives incredibly valuable. You can find subscription and archive information at <http://erlang.org/mailman/listinfo/erlang-questions>, and you will probably run into its archives regularly if you do searches on Erlang. If you prefer live chat to email, there is also an #erlang IRC channel on freenode (<http://freenode.net>).

If you prefer “real” live chat to the Internet, Erlang has a thriving conference circuit. For completely Erlang-focused venues, check out Erlang Factory (<http://www.erlang-factory.com/>), which produces a number of shows around the world, including the Erlang User Conference. The Association for Computing Machinery (ACM)’s Special Interest Group on Programming Languages (SIGPLAN) also holds an Erlang Workshop, about which you can learn more at <http://www.erlang.org/workshop/>.

There are also a lot of more informal “Erlounges” in a wide variety of locations, and many larger conferences, like O’Reilly’s own Open Source Conference (OSCON), include Erlang sessions and tutorials.

If you want to explore Erlang code, there’s lots of it on github; you can look around at the most active projects by visiting <https://github.com/languages/Erlang>.

Sharing the Gospel of Erlang

It may seem easy to argue for Erlang. The broad shift from single computers to networked and distributed systems of multiprocessor-based computing gives the Erlang environment a tremendous advantage over practically every other environment out there. More and more of the computing world is starting to face exactly the challenges that Erlang was built to address. Veterans of those challenges may find themselves breathing a sigh of relief as they can stop pondering toolsets that tried too hard to carry single-system approaches into a multi-system world.

At the same time, though, I’d encourage you to consider a bit of wisdom from Joe Armstrong: “New technologies have their best chance a) immediately after a disaster or b) at the start of a new project.” (<http://erlang.org/pipermail/erlang-questions/2012-October/069626.html>)

While it is possible you’re reading this because a project you’re working on has had a disaster (or you suspect it will have one soon), it’s easiest to apply Erlang to new projects, preferably projects where the inevitable beginner’s mistakes won’t create new disasters.

Find projects that look like fun to you, and that you can share within your organization or with the world. There’s no better way to show off the power of a programming language and environment than to build great things with it!

APPENDIX A

An Erlang Parts Catalog

Like every language, Erlang has drawers full of parts that are fun to peruse. These are a very few of the more common ones. If you want much much more, see http://www.erlang.org/doc/reference_manual/users_guide.html.

Shell Commands

You can use most Erlang functions from the shell, but these are ones that are exclusive to the shell.

Table A-1. Erlang shell commands

Command	Action
q()	Quits the shell <i>and</i> the Erlang runtime.
c(file)	Compiles the specified Erlang file.
b()	Displays all variable bindings.
f()	Clears all variable bindings.
f(X)	Clears specified variable binding.
h()	Prints the history list of commands.
e(N)	Repeats the command on line N.
v(N)	The return value of line N.
catch_exception(boolean)	Sets how strict the shell will be in passing errors.
rd(Name,Definition)	Defines a record type <i>Name</i> with contents specified by <i>Definition</i> .
rr(File)	Defines record types based on the contents of <i>File</i> .
rf()	Clears all record definitions. Can also clear specific definitions.
rl()	Lists all current record definitions.
pwd()	Gets the present working directory.

Command	Action
<code>ls()</code>	Lists files at the current location.
<code>cd(Directory)</code>	Changes to the specified <code>Directory</code> .

Reserved Words

There are a few Erlang terms you can't use outside of their intended context.

The Erlang compiler will wonder what you're trying to do if you use certain keywords as atoms or function names. It will try to treat your atoms as if they were code, and you can get very strange errors. After all, you should be able to have something called `band`, right?

Table A-2. Reserved words, which require careful use

```
after and andalso band begin bnot bor bsl bsr bxor
case catch cond div end fun if let not of
or orelse query receive rem try when xor
```

For function names, the answer is simple: use something else. If you want to use these as atoms, however, you can. You just need to enclose the offending reserved word in single quotes: '`'received'`', for example.

While they aren't reserved words, there are also a few atoms commonly used in return values. It's probably best to use them only in the circumstances where they're normally expected.

Table A-3. Commonly used return value atoms

Atom	Means
<code>ok</code>	Normal exit to a method. (Does <i>not</i> mean that whatever you asked for succeeded.)
<code>error</code>	Something went wrong. Typically accompanied by a larger explanation.
<code>undefined</code>	A value hasn't been assigned yet. Common in record instances.
<code>reply</code>	A reply is included with some kind of return value.
<code>noreply</code>	No return value is included. A response of some sort may come, however, from other communication.
<code>stop</code>	Used in OTP to signal that a server should stop, and triggers the <code>terminate</code> function.
<code>ignore</code>	Returned by OTP supervisor process that can't start a child.

Operators

Table A-4. Logical (Boolean) Operators

Operator	Description
and	logical and
or	logical or
xor	logical xor
not	unary logical not

The `not` operator is processed first.

`andalso` and `orelse` are also boolean operators for logical and and logical or, but they are short-circuit operators. If they don't need to process all the possibilities in their arguments, they stop at the first one that gives them a definite answer.

Table A-5. Term Comparison Operators

Operator	Description
<code>==</code>	equal to
<code>/=</code>	not equal to
<code>=<</code>	less than or equal to
<code><</code>	less than
<code>=>=</code>	greater than or equal to
<code>></code>	greater than
<code>=:=</code>	exactly equal to
<code>=/=</code>	exactly not equal to

You can compare elements of different types in Erlang. The relationship of types from “least” to “greatest” is:

`number < atom < reference < fun < port < pid < tuple < list < bit string`

Within `number`, you can compare integers and floats except with the more specific `=:=` and `=/=` operators, both of which will return `false` when you compare numbers of different types.

You can also compare tuples even when they contain different numbers of values. Erlang will go through the tuples from left to right and evaluate on the first value that returns a clear answer.

Table A-6. Arithmetic Operators

Operator	Description
+	unary + (positive)
-	unary - (negative)
+	addition
-	subtraction
*	multiplication
/	floating point division
div	integer division
rem	integer remainder of X/Y

Table A-7. Binary Operators

Operator	Description
bnot	unary bitwise not
band	bitwise and
bor	bitwise or
bxor	arithmetic bitwise xor
bsl	arithmetic bitshift left
bsr	bitshift right

Table A-8. Operator Precedence, from highest to lowest

Operator	Associativity
:	
#	
Unary + - bnot not	
/ * div rem band and	Left associative
+ - bor bxor bsl bsr or xor	Left associative
++ --	Right associative
==/===<< >= >=:==/+=	
andalso	
orelse	
= !	Right associative
catch	

The highest priority operator in an expression is evaluated first. Erlang evaluates operators with the same priority by following associative paths. (Left associative operators go left to right, right associative operators go right to left.)

Guard Components

Erlang allows only a limited subset of functions and other features in guard expressions, going well beyond a “no side effects” rule to keep a simple subset of possibilities. The list of allowed components includes the following:

- `true`
- Other constants (regarded as `false`)
- Term comparisons ([Table A-5](#))
- Arithmetic expressions ([Table A-6](#) and [Table A-7](#))
- Boolean expressions and short-circuit expressions (`andalso` and `orelse`)
- The following functions: `hd/1`, `is_atom/1`, `is_binary/1`, `is_bitstring/1`, `is_boolean/1`, `is_float/1`, `is_function/1`, `is_function/2`, `is_integer/1`, `is_list/1`, `is_number/1`, `is_pid/1`, `is_port/1`, `is_record/2`, `is_record/3`, `is_reference/1`, `is_tuple/1`

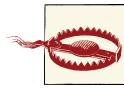
Common Functions

Table A-9. Mathematical functions

Function	Use
<code>math:pi/0</code>	The constant pi
<code>math:sin/1</code>	Sine
<code>math:cos/1</code>	Cosine
<code>math:tan/1</code>	Tangent
<code>math:asin/1</code>	Inverse sine (arcsine)
<code>math:acos/1</code>	Inverse cosine (arcosine)
<code>math:atan/1</code>	Inverse tangent (arctangent)
<code>math:atan2/2</code>	Arctangent that understands quadrants
<code>math:sinh/1</code>	Hyperbolic sine
<code>math:cosh/1</code>	Hyperbolic cosine
<code>math:tanh/1</code>	Hyperbolic tangent
<code>math:asinh/1</code>	Hyperbolic arcsine
<code>math:acosh/1</code>	Hyperbolic arccosine
<code>math:atanh/1</code>	Hyperbolic arctangent
<code>math:exp/1</code>	Exponential function
<code>math:log/1</code>	Natural logarithm (base e)

Function	Use
<code>math:log10/1</code>	Logarithm (base 10)
<code>math:pow/2</code>	First argument to the second argument power
<code>math:sqrt/1</code>	Square root
<code>math:erf/1</code>	Error function
<code>math:erfc/1</code>	Complementary error function

Arguments for all trigonometric functions are expressed in radians. To convert degrees to radians, divide by 180 and multiply by pi.



The `erf/1` and `erfc/1` functions may not be implemented in Windows. The Erlang documentation also warns more broadly that “Not all functions are implemented on all platforms,” but these come directly from the C language libraries.

Table A-10. Approachable higher-order functions for processing lists

function	Returns	Use
<code>lists:foreach/2</code>	ok	Side effects specified in function
<code>lists:map/2</code>	new list	Apply function to list values
<code>lists:filter/2</code>	subset	Creating list where function returns <code>true</code>
<code>lists:all/2</code>	boolean	Returns <code>true</code> if function <code>true</code> for all values, otherwise <code>false</code>
<code>lists:any/2</code>	boolean	Returns <code>true</code> if function <code>true</code> for any values, otherwise <code>false</code>
<code>lists:takewhile/2</code>	subset	Collects the head of the list until the function is <code>true</code>
<code>lists:dropwhile/2</code>	subset	Deletes the head of the list until the function is <code>true</code>
<code>lists:foldl/3</code>	accumulator	Passes function list value and accumulator, forward through list
<code>lists:foldr/3</code>	accumulator	Passes function list value and accumulator, backward through list
<code>lists:partition/3</code>	Tuple of two lists	Split list based on function

Chapter 7 describes these in greater detail.

Strings and Formatting

Table A-11. Simple control sequences for io:format and error_logger functions

Sequence	Produces
<code>~p</code>	Value, pretty-printed
<code>~w</code>	Value, no indentation
<code>~s</code>	Contents of a string
<code>~c</code>	ASCII character corresponding to a number

Sequence	Produces
<code>~t_c</code>	Unicode character corresponding to a number
<code>~i</code>	Ignores that item
<code>~n</code>	Newline (doesn't reference argument list)

Table A-12. Escape sequences for strings

Sequence	Produces
<code>\"</code>	double quote
<code>\'</code>	single quote
<code>\\"</code>	backslash
<code>\b</code>	backspace
<code>\d</code>	delete
<code>\e</code>	escape
<code>\f</code>	form feed
<code>\n</code>	newline
<code>\r</code>	carriage return
<code>\s</code>	space
<code>\t</code>	tab
<code>\v</code>	vertical tab
<code>\XYZ, \YZ, \Z</code>	character with octal representation XYZ, YZ or Z
<code>\XXY</code>	character in hex
<code>\x{X...}</code>	characters in hex, where X... is one or more hexadecimal characters
<code>^a...^z or ^A...^Z</code>	control-A to control-Z

Table A-13. Common string processing functions

Function	Returns
<code>string:len/1</code>	Length of the string. (Traverses string, so slows with big ones.)
<code>length/1</code>	Length of the string. (Traverses string, so slows with big ones.)
<code>string:concat/2</code>	A single string containing the two parts from the arguments.
<code>lists:concat/1</code>	A single string containing all the parts from the arguments.
<code>lists:append/1-2</code>	A single string containing all the parts from the arguments.
<code>lists:nth/2</code>	The character at the specified position.
<code>hd/1</code>	First character of the string.
<code>string:chr/2</code>	The position where the specified character first appears.
<code>string:str/2</code>	The position of a substring in a string.
<code>string:substr/2-3</code>	A segment from a string at a given position of a given length.
<code>string:sub_string/2-3</code>	A segment from a string between two positions.

Function	Returns
<code>string:tokens/2</code>	A list of pieces from a string broken at the specified separators.
<code>string:join/2</code>	A string made from the list of pieces with specified separators added.
<code>string:words/1-2</code>	The number of words in the string.
<code>string:chars/2-3</code>	A string that repeats a given character a given number of times.
<code>string:copies/2</code>	A string that repeats a given string a given number of times.
<code>string:strip/1-3</code>	A string with leading and/or trailing whitespace (or specified characters) removed.
<code>string:left/2-3</code>	A string of a specified length, padded with spaces on the right if needed.
<code>string:right/2-3</code>	A string of a specified length, padded with spaces on the left if needed.
<code>string:centre/2-3</code>	A string of a specified length, padded with spaces on the left <i>and</i> right if needed.
<code>lists:reverse/1-2</code>	A string in backwards order.
<code>string:to_float/1</code>	The float contents of the string, plus leftovers, or an error tuple.
<code>string:to_integer/1</code>	The integer contents of the string, plus leftovers, or an error tuple.
<code>string:to_lower/1</code>	A version of the string with all uppercase (Latin-1) characters converted to lowercase.
<code>string:to_upper/1</code>	A version of the string with all lowercase (Latin-1) characters converted to uppercase.
<code>integer_to_list/1-2</code>	A string version of an integer, optionally in a specified base.
<code>float_to_list/1</code>	A string version of a float.
<code>erlang:fun_to_list/1</code>	A string version of a fun.
<code>list_to_atom/1</code>	An atom version of a string.

Note: I'm working on creating a single wrapper module that assembles Erlang's tools for working with strings into one place. For more, visit <https://github.com/simonstl/erlang-simple-string>.

Data Types for Documentation and Analysis

Table A-14. Basic Data Types for -spec and EDoc

<code>atom()</code>	<code>binary()</code>	<code>float()</code>	<code>fun()</code>	<code>integer()</code>	<code>list()</code>	<code>tuple()</code>
<code>union()</code>	<code>node()</code>	<code>number()</code>	<code>string()</code>	<code>char()</code>	<code>byte()</code>	<code>[] (nil)</code>
<code>any()</code>	<code>none()</code>	<code>pid()</code>	<code>port()</code>	<code>reference()</code>		

For more, see http://www.erlang.org/doc/reference_manual/typespec.html .

APPENDIX B

OTP Templates

These are the full templates for `gen_server`, `supervisor`, and `application` from the Emacs mode for Erlang. Some pieces are more useful than others, but seeing the full set of expected responses can be useful.



Remember, the `noreply` atom doesn't mean "there will never be a reply" but rather that "this response isn't a reply."

Example B-1. A gen_server template from the Erlang mode for Emacs

```
%%%-----  
%%% @author $author  
%%% @copyright (C) $year, $company  
%%% @doc  
%%%  
%%% @end  
%%% Created : $fulldate  
%%%-----  
-module($basename).  
  
-behaviour(gen_server).  
  
%% API  
-export([start_link/0]).  
  
%% gen_server callbacks  
-export([init/1,  
        handle_call/3,  
        handle_cast/2,  
        handle_info/2,  
        terminate/2,
```

```

        code_change/3]).

-define(SERVER, ?MODULE).

-record(state, {}).

%%%=====
%%% API
%%%=====

%%-
%% @doc
%% Starts the server
%%
%% @spec start_link() -> {ok, Pid} | ignore | {error, Error}
%% @end
%%-
start_link() ->
    gen_server:start_link({local, ?SERVER}, ?MODULE, [], []).

%%%=====
%%% gen_server callbacks
%%%=====

%%-
%% @private
%% @doc
%% Initializes the server
%%
%% @spec init(Args) -> {ok, State} |
%% {ok, State, Timeout} |
%% {ignore, Reason} |
%% {stop, Reason}
%% @end
%%-
init([]) ->
    {ok, #state{}}.

%%-
%% @private
%% @doc
%% Handling call messages
%%
%% @spec handle_call(Request, From, State) ->
%% {reply, Reply, State} |
%% {reply, Reply, State, Timeout} |
%% {noreply, State} |
%% {noreply, State, Timeout} |
%% {stop, Reason, Reply, State} |
%% {stop, Reason, State}
%% @end
%%-

```

```

handle_call(_Request, _From, State) ->
    Reply = ok,
    {reply, Reply, State}.

%-
%% @private
%% @doc
%% Handling cast messages
%%
%% @spec handle_cast(Msg, State) -> {noreply, State} |
%% {noreply, State, Timeout} |
%% {stop, Reason, State}
%% @end
%-
handle_cast(_Msg, State) ->
    {noreply, State}.

%-
%% @private
%% @doc
%% Handling all non call/cast messages
%%
%% @spec handle_info(Info, State) -> {noreply, State} |
%% {noreply, State, Timeout} |
%% {stop, Reason, State}
%% @end
%-
handle_info(_Info, State) ->
    {noreply, State}.

%-
%% @private
%% @doc
%% This function is called by a gen_server when it is about to
%% terminate. It should be the opposite of Module:init/1 and do any
%% necessary cleaning up. When it returns, the gen_server terminates
%% with Reason. The return value is ignored.
%%
%% @spec terminate(Reason, State) -> void()
%% @end
%-
terminate(_Reason, _State) ->
    ok.

%-
%% @private
%% @doc
%% Convert process state when code is changed
%%
%% @spec code_change(OldVsn, State, Extra) -> {ok, NewState}
%% @end
%-

```

```

code_change(_OldVsn, State, _Extra) ->
{ok, State}.

%%%=====
%%% Internal functions
%%%=====

Example B-2. A supervisor template from the Erlang mode for Emacs

%%%-----
%%% @author $author
%%% @copyright (C) $year, $company
%%% @doc
%%% 
%%% @end
%%% Created : $fulldate
%%%-
-module($basename).

-behaviour(supervisor).

%% API
-export([start_link/0]).

%% Supervisor callbacks
-export([init/1]).

-define(SERVER, ?MODULE).

%%%=====
%%% API functions
%%%=====

%%-
%% @doc
%% Starts the supervisor
%%
%% @spec start_link() -> {ok, Pid} | ignore | {error, Error}
%% @end
%%-
start_link() ->
    supervisor:start_link({local, ?SERVER}, ?MODULE, []).

%%%=====
%%% Supervisor callbacks
%%%=====

%%-
%% @private
%% @doc
%% Whenever a supervisor is started using supervisor:start_link/[2,3],
%% this function is called by the new process to find out about
%% restart strategy, maximum restart frequency and child

```

```

%% specifications.

%% @spec init(Args) -> {ok, {SupFlags, [ChildSpec]}} |
%% ignore |
%% {error, Reason}
%% @end
%%-
init([]) ->
    RestartStrategy = one_for_one,
    MaxRestarts = 1000,
    MaxSecondsBetweenRestarts = 3600,

    SupFlags = {RestartStrategy, MaxRestarts, MaxSecondsBetweenRestarts},

    Restart = permanent,
    Shutdown = 2000,
    Type = worker,

    AChild = {'AName', {'AModule', start_link, []},
              Restart, Shutdown, Type, ['AModule']},

    {ok, {SupFlags, [AChild]}}.

%%-
%% Internal functions
%%-

```

Example B-3. An application module template from the Erlang mode for Emacs

```

%%%-
%%% @author $author
%%% @copyright (C) $year, $company
%%% @doc
%%% @end
%%% Created : $fulldate
%%%-
-module($basename).

-behaviour(application).

%% Application callbacks
-export([start/2, stop/1]).

%%-
%% Application callbacks
%%-
%%%-
%% @private
%% @doc
%% This function is called whenever an application is started using
%% application:start/[1,2], and should start the processes of the

```

```

%% application. If the application is structured according to the OTP
%% design principles as a supervision tree, this means starting the
%% top supervisor of the tree.
%%
%% @spec start(StartType, StartArgs) -> {ok, Pid} |
%% {ok, Pid, State} |
%% {error, Reason}
%% StartType = normal | {takeover, Node} | {failover, Node}
%% StartArgs = term()
%% @end
%%-----
start(_StartType, _StartArgs) ->
    case 'TopSupervisor':start_link() of
        {ok, Pid} ->
            {ok, Pid};
        Error ->
            Error
    end.

%%-----
%% @private
%% @doc
%% This function is called whenever an application has stopped. It
%% is intended to be the opposite of Module:start/2 and should do
%% any necessary cleaning up. The return value is ignored.
%%
%% @spec stop(State) -> void()
%% @end
%%-----
stop(_State) ->
    ok.

%%%=====
%%% Internal functions
%%%=====

```

About the Author

Simon St. Laurent is a web developer, network administrator, computer book author, and XML troublemaker living in Ithaca, NY. His books include *XML: A Primer*, *XML Elements of Style*, and *Building XML Applications, Cookies, and Sharing Bandwidth*. He is a contributing editor to XMLhack.com and an occasional contributor to XML.com.

Colophon

The cover image is from *Wood's Animate Creation*. The cover font is Adobe ITC Garamond. The text font is Adobe Minion Pro; the heading font is Adobe Myriad Condensed; and the code font is Dalton Maag's Ubuntu Mono.